

oberon_trading_bot_main

April 13, 2025

0.1 Pip Installs

You should really install these.

```
[1]: %%capture
!pip install beautifulsoup4 yfinance torch alpaca-trade-api alpaca-py
```

```
[2]: %%capture
!pip install "stable-baselines3[extra]" sb3-contrib gym pandas numpy matplotlib
↪ "shimmy>=2.0"
```

```
[3]: %%capture
!pip install python-dotenv
```

0.2 SEED (Set Before Trading Env)

Set the initial seed and environment.

```
[12]: import os
os.environ["CUBLAS_WORKSPACE_CONFIG"] = ":4096:8"

import numpy as np
import random
import torch

# Fix all seeds
SEED = 83819
np.random.seed(SEED)
random.seed(SEED)
torch.manual_seed(SEED)
torch.cuda.manual_seed_all(SEED)
torch.use_deterministic_algorithms(True)
```

0.3 Main Logic (Initialize)

Where the technical indicators reside.

```
[13]: import pandas as pd
import numpy as np
```

```

import yfinance as yf
import matplotlib.pyplot as plt
import matplotlib.dates as mdates
from matplotlib.patches import Rectangle
import requests
from bs4 import BeautifulSoup
import re
import io
from PIL import Image
import math
from datetime import datetime, timedelta

import torch
import torch.nn as nn
import torch.optim as optim
import torch.nn.functional as F

DEBUG = True
def debug_print(msg):
    if DEBUG:
        print(msg)

```

```

[14]: # FOMC Dates Scraper
def get_fomc_dates(start_date, end_date):
    url = "https://www.federalreserve.gov/monetarypolicy/fomccalendars.htm"
    try:
        response = requests.get(url, timeout=10)
        if response.status_code != 200:
            debug_print(f"Error: Received status code {response.status_code}")
            return []
        soup = BeautifulSoup(response.text, "html.parser")
        date_objs = []
        for text in soup.stripped_strings:
            matches = re.findall(r'([A-Za-z]+ \d{1,2}, \d{4})', text)
            for date_str in matches:
                try:
                    dt = datetime.strptime(date_str, "%B %d, %Y")
                    if dt not in date_objs:
                        date_objs.append(dt)
                except Exception:
                    continue
        date_objs = sorted(date_objs)
        start_dt = pd.to_datetime(start_date)
        end_dt = pd.to_datetime(end_date)
        filtered_dates = [dt for dt in date_objs if start_dt <= dt <= end_dt]
        return filtered_dates
    except Exception as e:

```

```

debug_print(f"Error scraping FOMC dates: {e}")
return []

```

```

[15]: # Indicator Functions
def exp_average(series, period):
    return series.ewm(span=period, adjust=False).mean()

def wilder_average(series, length):
    return series.ewm(alpha=1/length, adjust=False).mean()

def weighted_moving_average(series, window):
    weights = np.arange(1, window+1)
    return series.rolling(window).apply(lambda prices: np.dot(prices, weights)/
    ↪ weights.sum(), raw=True)

def t3(source, length=21, vf=0.7):
    ema1 = exp_average(source, length)
    ema2 = exp_average(ema1, length)
    gd1 = ema1*(1+vf) - ema2*vf
    ema11 = exp_average(gd1, length)
    ema22 = exp_average(ema11, length)
    gd2 = ema11*(1+vf) - ema22*vf
    ema111 = exp_average(gd2, length)
    ema222 = exp_average(ema111, length)
    gd3 = ema111*(1+vf) - ema222*vf
    return gd3

def vwma(series, window, volume):
    return (series*volume).rolling(window=window, min_periods=window).sum()/
    ↪ volume.rolling(window=window, min_periods=window).sum()

def rsi_function(close, sensitivity, rsiPeriod, rsiBase):
    delta = close.diff()
    gain = delta.clip(lower=0)
    loss = -delta.clip(upper=0)
    avg_gain = gain.rolling(window=rsiPeriod, min_periods=rsiPeriod).mean()
    avg_loss = loss.rolling(window=rsiPeriod, min_periods=rsiPeriod).mean()
    rs = avg_gain / avg_loss
    rsi = 100 - (100/(1+rs))
    rsi = rsi.fillna(50)
    rsi_adj = sensitivity*(rsi-rsiBase)
    return rsi_adj.clip(lower=0, upper=20)

def download_data(ticker, start_date, end_date):
    df = yf.download(ticker, start=pd.to_datetime(start_date), end=pd.
    ↪ to_datetime(end_date))
    if isinstance(df.columns, pd.MultiIndex):

```

```

        df.columns = [col[0].lower() for col in df.columns]
    else:
        df.columns = [str(col).lower() for col in df.columns]
    return df

def compute_bressert(df, n_period=8, r_period=13):
    df['Ln'] = df['low'].rolling(window=n_period, min_periods=1).min()
    df['Hn'] = df['high'].rolling(window=n_period, min_periods=1).max()
    df['Y'] = ((df['close']-df['Ln'])/(df['Hn']-df['Ln']))*100
    df['X'] = exp_average(df['Y'], r_period)
    df['Lxn'] = df['X'].rolling(window=n_period, min_periods=1).min()
    df['Hxn'] = df['X'].rolling(window=n_period, min_periods=1).max()
    df['DSS'] = ((df['X']-df['Lxn'])/(df['Hxn']-df['Lxn']))*100
    df['DSSb'] = exp_average(df['DSS'], r_period)
    df['DSSsignal'] = df['DSSb'].shift(1)
    return df

def compute_zscore(df, length_m=14):
    momentum = df['close'] - df['close'].shift(length_m)
    avgMomentum = momentum.rolling(window=length_m, min_periods=length_m).mean()
    stdDevMomentum = momentum.rolling(window=length_m, min_periods=length_m).
    ↪std().fillna(0)
    zScore = (momentum - avgMomentum)/stdDevMomentum
    return zScore

def compute_zero_lag_macd(source, fastLength=12, slowLength=26, signalLength=9, ↪
    ↪MacdEmaLength=9, useEma=True, useOldAlgo=False):
    if useEma:
        ma1 = source.ewm(span=fastLength, adjust=False).mean()
        ma2 = ma1.ewm(span=fastLength, adjust=False).mean()
    else:
        ma1 = source.rolling(window=fastLength, min_periods=fastLength).mean()
        ma2 = ma1.rolling(window=fastLength, min_periods=fastLength).mean()
    zerolagEMA = (2*ma1) - ma2
    if useEma:
        mas1 = source.ewm(span=slowLength, adjust=False).mean()
        mas2 = mas1.ewm(span=slowLength, adjust=False).mean()
    else:
        mas1 = source.rolling(window=slowLength, min_periods=slowLength).mean()
        mas2 = mas1.rolling(window=slowLength, min_periods=slowLength).mean()
    zerolagslowMA = (2*mas1) - mas2
    ZeroLagMACD = zerolagEMA - zerolagslowMA
    emasig1 = ZeroLagMACD.ewm(span=signalLength, adjust=False).mean()
    emasig2 = emasig1.ewm(span=signalLength, adjust=False).mean()
    if useOldAlgo:
        signal = ZeroLagMACD.rolling(window=signalLength, ↪
    ↪min_periods=signalLength).mean()

```

```

else:
    signal = (2*emasig1) - emasig2
    hist = ZeroLagMACD - signal
    upHist = hist.copy()
    upHist[hist <= 0] = 0
    downHist = hist.copy()
    downHist[hist > 0] = 0
    EMALine = ZeroLagMACD.ewm(span=MacdEmaLength, adjust=False).mean()
    dotUP = ZeroLagMACD.copy()
    dotUP[(ZeroLagMACD.shift(1) >= signal.shift(1)) | (ZeroLagMACD < signal)] = np.nan
    dotDN = ZeroLagMACD.copy()
    dotDN[(ZeroLagMACD.shift(1) <= signal.shift(1)) | (ZeroLagMACD > signal)] = np.nan
    return {
        "ZeroLagMACD": ZeroLagMACD,
        "signal": signal,
        "hist": hist,
        "upHist": upHist,
        "downHist": downHist,
        "EMALine": EMALine,
        "dotUP": dotUP,
        "dotDN": dotDN
    }

def compute_basic_macd(source, fast=12, slow=26, signal=9):
    ema_fast = source.ewm(span=fast, adjust=False).mean()
    ema_slow = source.ewm(span=slow, adjust=False).mean()
    macd_line = ema_fast - ema_slow
    signal_line = macd_line.ewm(span=signal, adjust=False).mean()
    hist_line = macd_line - signal_line
    return {
        "basicMACD": macd_line,
        "signal": signal_line,
        "hist": hist_line
    }

```

[16]: *# Historical Signals Extraction*

```

def extract_signals(df, signalUp_ZLMA, signalDn_ZLMA, bullPt, bearPt,
                    upSig_MCDX, dnSig_MCDX, length_m=14):
    """
    Combines ZLMA, RSI, MCDX, DSS signals in a single table.
    """
    signals = []
    zScore = compute_zscore(df, length_m)

    # ZLMA

```

```

for dt in df.index[signalUp_ZLMA.fillna(False)]:
    signals.append({
        "Date": dt.strftime("%Y-%m-%d"),
        "Signal": "ZLMA Buy",
        "Z-Score": round(zScore.loc[dt],2) if not pd.isna(zScore.loc[dt])
↪else None
    })
for dt in df.index[signalDn_ZLMA.fillna(False)]:
    signals.append({
        "Date": dt.strftime("%Y-%m-%d"),
        "Signal": "ZLMA Sell",
        "Z-Score": round(zScore.loc[dt],2) if not pd.isna(zScore.loc[dt])
↪else None
    })
# RSI
if isinstance(bullPt, pd.Series):
    for dt in bullPt.dropna().index:
        signals.append({
            "Date": dt.strftime("%Y-%m-%d"),
            "Signal": "RSI Buy",
            "Z-Score": round(zScore.loc[dt],2) if not pd.isna(zScore.
↪loc[dt]) else None
        })
if isinstance(bearPt, pd.Series):
    for dt in bearPt.dropna().index:
        signals.append({
            "Date": dt.strftime("%Y-%m-%d"),
            "Signal": "RSI Sell",
            "Z-Score": round(zScore.loc[dt],2) if not pd.isna(zScore.
↪loc[dt]) else None
        })
# MCDX
if isinstance(upSig_MCDX, pd.Series):
    for dt in upSig_MCDX.dropna().index:
        signals.append({
            "Date": dt.strftime("%Y-%m-%d"),
            "Signal": "MCDX Buy",
            "Z-Score": round(zScore.loc[dt],2) if not pd.isna(zScore.
↪loc[dt]) else None
        })
if isinstance(dnSig_MCDX, pd.Series):
    for dt in dnSig_MCDX.dropna().index:
        signals.append({
            "Date": dt.strftime("%Y-%m-%d"),
            "Signal": "MCDX Sell",
            "Z-Score": round(zScore.loc[dt],2) if not pd.isna(zScore.
↪loc[dt]) else None

```

```

    })
    # DSS
    for i in range(1, len(df)):
        if (pd.notna(df['DSSb'].iloc[i]) and pd.notna(df['DSSsignal'].iloc[i]))
        and
            pd.notna(df['DSSb'].iloc[i-1]) and pd.notna(df['DSSsignal'].
            iloc[i-1])):
            if df['DSSb'].iloc[i] > df['DSSsignal'].iloc[i] and df['DSSb'].
            iloc[i-1] <= df['DSSsignal'].iloc[i-1]:
                dt = df.index[i]
                signals.append({
                    "Date": dt.strftime("%Y-%m-%d"),
                    "Signal": "DSS Buy",
                    "Z-Score": round(zScore.loc[dt],2) if not pd.isna(zScore.
            loc[dt]) else None
                })
            elif df['DSSb'].iloc[i] < df['DSSsignal'].iloc[i] and df['DSSb'].
            iloc[i-1] >= df['DSSsignal'].iloc[i-1]:
                dt = df.index[i]
                signals.append({
                    "Date": dt.strftime("%Y-%m-%d"),
                    "Signal": "DSS Sell",
                    "Z-Score": round(zScore.loc[dt],2) if not pd.isna(zScore.
            loc[dt]) else None
                })

    signals_df = pd.DataFrame(signals)
    if not signals_df.empty:
        signals_df["Date"] = pd.to_datetime(signals_df["Date"])
        signals_df = signals_df.sort_values("Date", ascending=False)
    return signals_df

def extract_momentum_signals(df, length_m=14):
    """
    Momentum-based signals. Checks momentum grade changes, direction changes,
    and state changes.
    """
    momentum = df['close'] - df['close'].shift(length_m)
    avgMomentum = momentum.rolling(window=length_m, min_periods=length_m).mean()
    stdDevMomentum = momentum.rolling(window=length_m, min_periods=length_m).
    std().fillna(0)
    zScore = (momentum - avgMomentum) / stdDevMomentum

    def grade(x):
        if x >= 2:
            return "A"

```

```

elif x >= 1:
    return "B"
elif x >= 0:
    return "C"
elif x >= -1:
    return "D"
elif x >= -2:
    return "E"
else:
    return "F"

momentum_grade = zScore.apply(grade)
momentum_direction = momentum.apply(lambda x: "Increasing" if x>0 else
↪ "Decreasing")

momentum_state = []
for i in range(len(momentum)):
    if i == 0:
        momentum_state.append("N/A")
    else:
        if abs(momentum.iloc[i]) < abs(momentum.iloc[i-1]) * 0.0001:
            momentum_state.append("Consolidating")
        elif momentum.iloc[i] * momentum.iloc[i-1] < 0:
            if momentum.iloc[i] > 0:
                momentum_state.append("Turning Up")
            else:
                momentum_state.append("Turning Down")
        elif (momentum.iloc[i] > 0 and (momentum.iloc[i] - momentum.
↪ iloc[i-1]) < 0) or (momentum.iloc[i] < 0 and (momentum.iloc[i] - momentum.
↪ iloc[i-1]) > 0):
            momentum_state.append("Stalling")
        elif momentum.iloc[i] > 0:
            momentum_state.append("Positive Trending")
        else:
            momentum_state.append("Negative Trending")
momentum_state = pd.Series(momentum_state, index=df.index)

signals = []
for i in range(1, len(df)):
    # Grade
    if momentum_grade.iloc[i] != momentum_grade.iloc[i-1]:
        signals.append({
            "Date": df.index[i].strftime("%Y-%m-%d"),
            "Signal": f"Momentum Grade Changed to {momentum_grade.iloc[i]}",
            "Z-Score": round(zScore.iloc[i],2)
        })
    # Direction

```



```

        if momentum_direction.iloc[i] != momentum_direction.iloc[i-1]:
            signals.append({
                "Date": df.index[i].strftime("%Y-%m-%d"),
                "Signal": f"Momentum Direction Changed to {momentum_direction.
↪iloc[i]}",
                "Z-Score": round(zScore.iloc[i],2)
            })
        # State
        if momentum_state.iloc[i] != momentum_state.iloc[i-1]:
            signals.append({
                "Date": df.index[i].strftime("%Y-%m-%d"),
                "Signal": f"Momentum State Changed to {momentum_state.iloc[i]}",
                "Z-Score": round(zScore.iloc[i],2)
            })
    signals_df = pd.DataFrame(signals)
    if not signals_df.empty:
        signals_df["Date"] = pd.to_datetime(signals_df["Date"])
        signals_df = signals_df.sort_values("Date", ascending=False)
    return signals_df

def extract_zero_macd_signals(df, zero_macd_dict, length_m=14):
    macd_line = zero_macd_dict["ZeroLagMACD"]
    macd_mean = macd_line.rolling(window=length_m, min_periods=length_m).mean()
    macd_std = macd_line.rolling(window=length_m, min_periods=length_m).std().
↪replace(0, np.nan)
    macd_zscore = (macd_line - macd_mean)/macd_std

    signals = []
    for i in range(1, len(df)):
        if (pd.notna(zero_macd_dict["ZeroLagMACD"].iloc[i]) and
            pd.notna(zero_macd_dict["signal"].iloc[i]) and
            pd.notna(zero_macd_dict["ZeroLagMACD"].iloc[i-1]) and
            pd.notna(zero_macd_dict["signal"].iloc[i-1])):
            dt = df.index[i]
            if zero_macd_dict["ZeroLagMACD"].iloc[i] > zero_macd_dict["signal"].
↪iloc[i] and zero_macd_dict["ZeroLagMACD"].iloc[i-1] <=
↪zero_macd_dict["signal"].iloc[i-1]:
                signals.append({
                    "Date": dt.strftime("%Y-%m-%d"),
                    "Signal": "ZeroLag MACD Buy",
                    "Z-Score": round(macd_zscore.iloc[i],2) if not pd.
↪isna(macd_zscore.iloc[i]) else None
                })
            elif zero_macd_dict["ZeroLagMACD"].iloc[i] <
↪zero_macd_dict["signal"].iloc[i] and zero_macd_dict["ZeroLagMACD"].iloc[i-1]
↪>= zero_macd_dict["signal"].iloc[i-1]:
                signals.append({

```

```

        "Date": dt.strftime("%Y-%m-%d"),
        "Signal": "ZeroLag MACD Sell",
        "Z-Score": round(macд_zscore.iloc[i],2) if not pd.
↪isna(macд_zscore.iloc[i]) else None
    })
    signals_df = pd.DataFrame(signals)
    if not signals_df.empty:
        signals_df["Date"] = pd.to_datetime(signals_df["Date"])
        signals_df = signals_df.sort_values("Date", ascending=False)
    return signals_df

def extract_basic_macд_signals(df, basic_macд_dict, length_m=14):
    macд_line = basic_macд_dict["basicMACD"]
    signal_line = basic_macд_dict["signal"]
    macд_mean = macд_line.rolling(window=length_m, min_periods=length_m).
↪mean()
    macд_std = macд_line.rolling(window=length_m, min_periods=length_m).
↪std().replace(0, np.nan)
    macд_zscore = (macд_line - macд_mean)/macд_std

    signals=[]
    for i in range(1, len(df)):
        if (pd.notna(macд_line.iloc[i]) and pd.notna(signal_line.iloc[i]) and
            pd.notna(macд_line.iloc[i-1]) and pd.notna(signal_line.iloc[i-1])):
            dt= df.index[i]
            # cross up
            if macд_line.iloc[i]> signal_line.iloc[i] and macд_line.iloc[i-1]<=
↪signal_line.iloc[i-1]:
                signals.append({
                    "Date": dt.strftime("%Y-%m-%d"),
                    "Signal": "Basic MACD Buy",
                    "Z-Score": round(macд_zscore.iloc[i],2) if not pd.
↪isna(macд_zscore.iloc[i]) else None
                })
            # cross down
            elif macд_line.iloc[i]< signal_line.iloc[i] and macд_line.
↪iloc[i-1]>= signal_line.iloc[i-1]:
                signals.append({
                    "Date": dt.strftime("%Y-%m-%d"),
                    "Signal": "Basic MACD Sell",
                    "Z-Score": round(macд_zscore.iloc[i],2) if not pd.
↪isna(macд_zscore.iloc[i]) else None
                })
    signals_df= pd.DataFrame(signals)
    if not signals_df.empty:
        signals_df["Date"]= pd.to_datetime(signals_df["Date"])

```

```

        signals_df= signals_df.sort_values("Date", ascending=False)
    return signals_df

```

[17]: *# RSI and MCDX signals*

```

def calc_rsi_entire_series(df, bullPt, bearPt):
    rsi_series = []
    last_signal = "Sell"
    for i in range(len(df)):
        if isinstance(bullPt, pd.Series) and pd.notna(bullPt.iloc[i]):
            last_signal = "Buy"
        elif isinstance(bearPt, pd.Series) and pd.notna(bearPt.iloc[i]):
            last_signal = "Sell"
        rsi_series.append(last_signal)
    return pd.Series(rsi_series, index=df.index)

def calc_mcdx_entire_series(df, upSig_MCDX, dnSig_MCDX):
    mcdx_series = []
    last_signal = "Sell"
    for i in range(len(df)):
        if pd.notna(upSig_MCDX.iloc[i]):
            last_signal = "Buy"
        elif pd.notna(dnSig_MCDX.iloc[i]):
            last_signal = "Sell"
        mcdx_series.append(last_signal)
    return pd.Series(mcdx_series, index=df.index)

```

[18]: *# Weighted scoreboard*

```

def scoreboard_for_day(
    df, i,
    rsi_series, mcdx_series,
    zero_macd_dict, basic_macd_dict,
    length_m=14,
    weight_zlma=1.0, weight_rsi=1.0, weight_mcdx=1.0, weight_dss=1.0,
    weight_zscore=1.0, weight_mg=1.0, weight_md=1.0, weight_ms=1.0,
    weight_zeromacd=1.0, weight_basicmacd=1.0
):
    # ZLMA => Buy if zlma>ema_value
    zlma_status= "Buy" if df['zlma'].iloc[i]> df['ema_value'].iloc[i] else
    ↪ "Sell"

    # RSI => from daily rsi_series
    rsi_status= rsi_series.iloc[i]
    # MCDX => from daily mcdx_series
    mcdx_status= mcdx_series.iloc[i]
    # DSS => Buy if DSSb>DSSsignal
    dss_status= "Buy" if df['DSSb'].iloc[i]> df['DSSsignal'].iloc[i] else "Sell"

    # Z-Score => positive => Buy

```

```

zVal= compute_zscore(df, length_m).iloc[i]
z_status= "Buy" if zVal>=0 else "Sell"

# Momentum Grade
momentum= df['close']- df['close'].shift(length_m)
avgM= momentum.rolling(window=length_m, min_periods=length_m).mean()
stdM= momentum.rolling(window=length_m, min_periods=length_m).std().
↳fillna(0)
z_m= (momentum-avgM)/stdM
mg_z= z_m.iloc[i] if not pd.isna(z_m.iloc[i]) else 0
if mg_z>=2:
    mg_grade="A"
elif mg_z>=1:
    mg_grade="B"
elif mg_z>=0:
    mg_grade="C"
elif mg_z>=-1:
    mg_grade="D"
elif mg_z>=-2:
    mg_grade="E"
else:
    mg_grade="F"
mg_status= "Buy" if mg_grade in ["A","B","C"] else "Sell"

# Momentum Direction
if i==0:
    md_status="Sell"
else:
    md_status= "Buy" if momentum.iloc[i]> momentum.iloc[i-1] else "Sell"

# Momentum State
if i==0:
    ms_val=0
else:
    val_now= momentum.iloc[i]
    val_prev= momentum.iloc[i-1]
    change= val_now- val_prev
    ms_val=0
    if val_now>0 and val_prev<0:
        ms_val=1
    elif val_now<0 and val_prev>0:
        ms_val=-1
    elif abs(val_now)< abs(avgM.iloc[i])*0.1:
        ms_val=0
    elif (val_now>0 and change<0) or (val_now<0 and change>0):
        ms_val=0
    elif val_now>0:

```

```

        ms_val=1
    else:
        ms_val=-1
    if ms_val>0:
        ms_status="Buy"
    elif ms_val<0:
        ms_status="Sell"
    else:
        ms_status="Neutral"

    zero_macd_status="Sell"
    if zero_macd_dict is not None:
        if zero_macd_dict["ZeroLagMACD"].iloc[i]> zero_macd_dict["signal"].
↪iloc[i]:
            zero_macd_status="Buy"
        else:
            zero_macd_status="Sell"

    basic_macd_status="Sell"
    if basic_macd_dict is not None:
        if basic_macd_dict["basicMACD"].iloc[i]> basic_macd_dict["signal"].
↪iloc[i]:
            basic_macd_status="Buy"
        else:
            basic_macd_status="Sell"

    scoreboard=[]
    scoreboard.append(weight_zlma if zlma_status=="Buy" else -weight_zlma)
    scoreboard.append(weight_rsi if rsi_status=="Buy" else -weight_rsi)
    scoreboard.append(weight_mcdx if mcdx_status=="Buy" else -weight_mcdx)
    scoreboard.append(weight_dss if dss_status=="Buy" else -weight_dss)
    scoreboard.append(weight_zscore if zVal>=0 else -weight_zscore)
    scoreboard.append(weight_mg if mg_status=="Buy" else -weight_mg)
    scoreboard.append(weight_md if md_status=="Buy" else -weight_md)

    ms_val_weighted=0
    if ms_status=="Buy":
        ms_val_weighted= weight_ms
    elif ms_status=="Sell":
        ms_val_weighted= -weight_ms
    scoreboard.append(ms_val_weighted)

    scoreboard.append(weight_zeromacd if zero_macd_status=="Buy" else
↪-weight_zeromacd)
    scoreboard.append(weight_basicmacd if basic_macd_status=="Buy" else
↪-weight_basicmacd)

```

```

total_signals= sum(1 for x in scoreboard if x!=0)
bullish_count= sum(1 for x in scoreboard if x>0)

if total_signals==0:
    return "Neutral"
else:
    if bullish_count>= (total_signals/2.0):
        return "Buy"
    else:
        return "Sell"

```

```

[19]: # Historical Trade Table
def simulate_trades_overall_signal(
    df, rsi_series, mcdx_series,
    zero_macd_dict, basic_macd_dict,
    length_m=14,
    weight_zlma=1.0, weight_rsi=1.0, weight_mcdx=1.0, weight_dss=1.0,
    weight_zscore=1.0, weight_mg=1.0, weight_md=1.0, weight_ms=1.0,
    weight_zeromacd=1.0, weight_basicmacd=1.0
):
    position = 0
    entry_price = 0.0
    trade_log = []
    cumulative_pnl_sum = 0.0
    prev_signal = "Neutral"

    for i in range(len(df)):
        day_signal = scoreboard_for_day(
            df, i, rsi_series, mcdx_series,
            zero_macd_dict, basic_macd_dict,
            length_m,
            weight_zlma, weight_rsi, weight_mcdx, weight_dss,
            weight_zscore, weight_mg, weight_md, weight_ms,
            weight_zeromacd, weight_basicmacd
        )
        price_i = df['close'].iloc[i]
        date_i = df.index[i]

        # Flip from Sell->Buy => close short
        if prev_signal in ["Sell", "Neutral"] and day_signal == "Buy":
            if position == -1:
                exit_price = price_i
                pnl_pct = ((entry_price / exit_price) - 1)*100
                cumulative_pnl_sum += pnl_pct
                trade_log.append({
                    "EntryDate": entry_dt_str,
                    "ExitDate": date_i.strftime("%Y-%m-%d"),

```

```

        "Position": "Short",
        # Round to 4 decimal places:
        "EntryPrice": round(entry_price, 4),
        "ExitPrice": round(exit_price, 4),
        "PnL%": round(pnl_pct, 2),
        "CumulativePnL%": round(cumulative_pnl_sum, 2)
    })
    position = 0
    position = 1
    entry_price = price_i
    entry_dt_str = date_i.strftime("%Y-%m-%d")

# Flip from Buy->Sell => close long
elif prev_signal in ["Buy", "Neutral"] and day_signal == "Sell":
    if position == 1:
        exit_price = price_i
        pnl_pct = ((exit_price / entry_price) - 1)*100
        cumulative_pnl_sum += pnl_pct
        trade_log.append({
            "EntryDate": entry_dt_str,
            "ExitDate": date_i.strftime("%Y-%m-%d"),
            "Position": "Long",
            "EntryPrice": round(entry_price, 4),
            "ExitPrice": round(exit_price, 4),
            "PnL%": round(pnl_pct, 2),
            "CumulativePnL%": round(cumulative_pnl_sum, 2)
        })
        position = 0
    position = -1
    entry_price = price_i
    entry_dt_str = date_i.strftime("%Y-%m-%d")

prev_signal = day_signal

# Close any open position at the end
if position != 0:
    final_price = df['close'].iloc[-1]
    final_date = df.index[-1]
    if position == 1:
        exit_price = final_price
        pnl_pct = ((exit_price / entry_price) - 1)*100
        cumulative_pnl_sum += pnl_pct
        trade_log.append({
            "EntryDate": entry_dt_str,
            "ExitDate": final_date.strftime("%Y-%m-%d"),
            "Position": "Long",
            "EntryPrice": round(entry_price, 4),

```

```

        "ExitPrice": round(exit_price, 4),
        "PnL%": round(pnl_pct, 2),
        "CumulativePnL%": round(cumulative_pnl_sum, 2)
    })
    else:
        exit_price = final_price
        pnl_pct = ((entry_price / exit_price) - 1)*100
        cumulative_pnl_sum += pnl_pct
        trade_log.append({
            "EntryDate": entry_dt_str,
            "ExitDate": final_date.strftime("%Y-%m-%d"),
            "Position": "Short",
            "EntryPrice": round(entry_price, 4),
            "ExitPrice": round(exit_price, 4),
            "PnL%": round(pnl_pct, 2),
            "CumulativePnL%": round(cumulative_pnl_sum, 2)
        })

trade_df = pd.DataFrame(trade_log)
trade_df["ExitDate_dt"] = pd.to_datetime(trade_df["ExitDate"])
trade_df = trade_df.sort_values("ExitDate_dt", ascending=False).
↪reset_index(drop=True)
trade_df.drop(columns=["ExitDate_dt"], inplace=True)
return trade_df

```

```

[20]: # Scoreboard with ZeroLag + Basic MACD
def extract_current_status(
    df, signalUp_ZLMA, signalDn_ZLMA, bullPt, bearPt,
    upSig_MCDX, dnSig_MCDX,
    length_m=14,
    zero_macd_dict=None, basic_macd_dict=None,

    # Weighted scoreboard
    weight_zlma=1.0, weight_rsi=1.0, weight_mcdx=1.0, weight_dss=1.0,
    weight_zscore=1.0, weight_mg=1.0, weight_md=1.0, weight_ms=1.0,
    weight_zeromacd=1.0, weight_basicmacd=1.0
):
    i= len(df)-1
    zlma_status= "Buy" if df['zlma'].iloc[i]> df['ema_value'].iloc[i] else_
↪"Sell"

    # RSI daily
    daily_rsi_series = calc_rsi_entire_series(df, bullPt, bearPt)
    rsi_status = daily_rsi_series.iloc[i]
    # MCDX daily
    daily_mcdx_series = calc_mcdx_entire_series(df, upSig_MCDX, dnSig_MCDX)
    mcdx_status = daily_mcdx_series.iloc[i]

```



```

# DSS => buy if DSSb>DSSsignal
dss_status= "Buy" if df['DSSb'].iloc[i]> df['DSSsignal'].iloc[i] else "Sell"

# Z-Score
zVal= compute_zscore(df, length_m).iloc[i]
z_status= "Buy" if zVal>=0 else "Sell"

# Momentum
momentum= df['close']- df['close'].shift(length_m)
avgM= momentum.rolling(window=length_m, min_periods=length_m).mean()
stdM= momentum.rolling(window=length_m, min_periods=length_m).std().
↳fillna(0)
z_m= (momentum-avgM)/ stdM
mg_z= z_m.iloc[i] if not pd.isna(z_m.iloc[i]) else 0
if mg_z>=2:
    mg_letter="A"
elif mg_z>=1:
    mg_letter="B"
elif mg_z>=0:
    mg_letter="C"
elif mg_z>=-1:
    mg_letter="D"
elif mg_z>=-2:
    mg_letter="E"
else:
    mg_letter="F"
mg_status= "Buy" if mg_letter in ["A","B","C"] else "Sell"

# momentum direction
if i==0:
    md_status="Sell"
else:
    md_status= "Buy" if momentum.iloc[i]> momentum.iloc[i-1] else "Sell"

# momentum state
if i==0:
    ms_val=0
else:
    val_now= momentum.iloc[i]
    val_prev= momentum.iloc[i-1]
    change= val_now- val_prev
    ms_val=0
    if val_now>0 and val_prev<0:
        ms_val=1
    elif val_now<0 and val_prev>0:
        ms_val=-1
    elif abs(val_now)< abs(avgM.iloc[i])*0.1:

```

```

        ms_val=0
    elif (val_now>0 and change<0) or (val_now<0 and change>0):
        ms_val=0
    elif val_now>0:
        ms_val=1
    else:
        ms_val=-1
if ms_val>0:
    ms_status="Buy"
elif ms_val<0:
    ms_status="Sell"
else:
    ms_status="Neutral"

zero_macd_status= "Sell"
if zero_macd_dict is not None:
    if zero_macd_dict["ZeroLagMACD"].iloc[i]> zero_macd_dict["signal"].
↪iloc[i]:
        zero_macd_status= "Buy"
    else:
        zero_macd_status= "Sell"

basic_macd_status="Sell"
if basic_macd_dict is not None:
    if basic_macd_dict["basicMACD"].iloc[i]> basic_macd_dict["signal"].
↪iloc[i]:
        basic_macd_status= "Buy"
    else:
        basic_macd_status= "Sell"

scoreboard=[]
scoreboard.append(weight_zlma if zlma_status=="Buy" else -weight_zlma)
scoreboard.append(weight_rsi if rsi_status=="Buy" else -weight_rsi)
scoreboard.append(weight_mcdx if mcdx_status=="Buy" else -weight_mcdx)
scoreboard.append(weight_dss if dss_status=="Buy" else -weight_dss)
scoreboard.append(weight_zscore if zVal>=0 else -weight_zscore)
scoreboard.append(weight_mg if mg_status=="Buy" else -weight_mg)
scoreboard.append(weight_md if md_status=="Buy" else -weight_md)

ms_val_weighted=0
if ms_status=="Buy":
    ms_val_weighted= weight_ms
elif ms_status=="Sell":
    ms_val_weighted= -weight_ms
scoreboard.append(ms_val_weighted)

```

```

        scoreboard.append(weight_zeromacd if zero_macd_status=="Buy" else_
↪-weight_zeromacd)
        scoreboard.append(weight_basicmacd if basic_macd_status=="Buy" else_
↪-weight_basicmacd)

    total_signals= sum(1 for x in scoreboard if x!=0)
    bullish_count= sum(1 for x in scoreboard if x>0)
    total_score= sum(scoreboard)

    if total_signals==0:
        overall= "Neutral"
    else:
        if bullish_count>= (total_signals/2.0):
            overall= "Buy"
        else:
            overall= "Sell"

    indicators= [
        "ZLMA","RSI","MCDX","DSS","Z-Score",
        "Momentum Grade","Momentum Direction","Momentum State",
        "ZeroLag MACD","Basic MACD","ScoreSum","Overall Trade"
    ]
    signals_=[
        zlma_status,
        rsi_status,
        mcdx_status,
        dss_status,
        round(zVal,2),
        mg_letter,
        ("Increasing" if md_status=="Buy" else "Decreasing"),
        ms_status,
        zero_macd_status,
        basic_macd_status,
        str(round(total_score,2)),
        overall
    ]
    return pd.DataFrame({"Indicator": indicators, "Current Signal": signals_})

```

```

[21]: # Plot
def create_generic_plot(df, ticker, start_date, end_date,
                        ema_value, zlma, signalUp_ZLMA, signalDn_ZLMA,
↪zlma_color, ema_color,
                        rsi_ma_base, rsi_upper_bound, rsi_lower_bound, bullPt,
↪bearPt,
                        b_X, b_DSSb, b_DSSsignal,
                        hbma, threshold, upSig_MCDX, dnSig_MCDX,
                        Dump, DnCandle, PumpCandle, Retest, Banker,

```

```

        iv_series, zero_macd_dict, basic_macd_dict=None,
        momentum_length=14):
fig, axs = plt.subplots(8,1,sharex=True,figsize=(12,20),
                        gridspec_kw={"height_ratios": [2,1,1,1,1,1,1,1]})
fig.suptitle(f"{ticker} - Generic Multi-Panel Chart with Momentum & Dual
MACD", fontsize=14)
x_vals = mdates.date2num(df.index.to_pydatetime())

# Panel 1: Price + ZLMA + RSI + Momentum test
for i in range(len(df)):
    o, c, h, l = df['open'].iloc[i], df['close'].iloc[i], df['high'].
    iloc[i], df['low'].iloc[i]
    color= 'green' if c>= o else 'red'
    axs[0].plot([x_vals[i], x_vals[i]], [l,h], color=color, linewidth=1,
    zorder=1)
    candle_width= 0.6
    axs[0].add_patch(Rectangle((x_vals[i]-candle_width/2, o), candle_width,
    c-o,
                                facecolor=color, edgecolor=color, zorder=2))
    axs[0].plot(df.index, df['EMA_50'], label="EMA 50", color='blue',
    linewidth=1.5, zorder=3)
    axs[0].plot(df.index, df['EMA_100'], label="EMA 100", color='orange',
    linewidth=1.5, zorder=3)
    axs[0].plot(df.index, df['EMA_200'], label="EMA 200", color='purple',
    linewidth=1.5, zorder=3)
    axs[0].plot(df.index, df['EMA_500'], label="EMA 500", color='brown',
    linewidth=1.5, zorder=3)
    axs[0].plot(df.index, ema_value, label="EMA (Trend)", color=ema_color,
    linewidth=2, zorder=4)
    axs[0].plot(df.index, zlma, label="ZLMA", color=zlma_color,
    linewidth=2, zorder=4)
    axs[0].fill_between(df.index, zlma, ema_value, where=(zlma>=ema_value),
    facecolor="darkgreen", alpha=0.3, interpolate=True, zorder=3)
    axs[0].fill_between(df.index, zlma, ema_value, where=(zlma<ema_value),
    facecolor="darkred", alpha=0.3, interpolate=True, zorder=3)
    axs[0].scatter(df.index, zlma.where(signalUp_ZLMA), color="cyan",
    marker="o", s=50, label="ZLMA Buy", zorder=5)
    axs[0].scatter(df.index, zlma.where(signalDn_ZLMA), color="magenta",
    marker="o", s=50, label="ZLMA Sell", zorder=5)

    axs[0].plot(df.index, rsi_ma_base, label="RSI Trail Base",
    color="gray", linestyle="--", linewidth=1)
    axs[0].plot(df.index, rsi_upper_bound, label="RSI Trail Upper",
    color="blue", linewidth=1)
    axs[0].plot(df.index, rsi_lower_bound, label="RSI Trail Lower",
    color="red", linewidth=1)

```

```

    if isinstance(bullPt, pd.Series):
        axs[0].scatter(df.index, bullPt, color="cyan", marker="^", s=50,
↳label="RSI Buy", zorder=6)
    if isinstance(bearPt, pd.Series):
        axs[0].scatter(df.index, bearPt, color="magenta", marker="v", s=50,
↳label="RSI Sell", zorder=6)
    axs[0].fill_between(df.index, rsi_ma_base, rsi_upper_bound,
↳facecolor="darkgreen", alpha=0.2, interpolate=True)
    axs[0].fill_between(df.index, rsi_lower_bound, rsi_ma_base,
↳facecolor="darkred", alpha=0.2, interpolate=True)

    fomc_dates= get_fomc_dates(start_date, end_date)
    for i, dt in enumerate(fomc_dates):
        axs[0].axvline(dt, color="purple", linestyle="--", linewidth=1,
↳label="FOMC" if i==0 else "")

    axs[0].set_ylabel("Price")
    # Legend bottom-right
    axs[0].legend(loc="lower right", ncol=3, fontsize=8)

    # Momentum text
    momentum= df['close']- df['close'].shift(momentum_length)
    avgM= momentum.rolling(window=momentum_length, min_periods=momentum_length).
↳mean()
    stdM= momentum.rolling(window=momentum_length, min_periods=momentum_length).
↳std().fillna(0)
    zScore_m= (momentum- avgM)/ stdM
    if len(zScore_m.dropna())>0:
        last_z= zScore_m.iloc[-1]
    else:
        last_z= np.nan
    if not np.isnan(last_z):
        if last_z>=2:
            gradeStr= "A (Strong Positive Momentum)"
            gradeColor= "green"
        elif last_z>=1:
            gradeStr= "B (Moderate Positive Momentum)"
            gradeColor= "lightgreen"
        elif last_z>=0:
            gradeStr= "C (Weak Positive Momentum)"
            gradeColor= "goldenrod"
        elif last_z>=-1:
            gradeStr= "D (Weak Negative Momentum)"
            gradeColor= "orange"
        elif last_z>=-2:
            gradeStr= "E (Moderate Negative Momentum)"

```

```

        gradeColor= "red"
    else:
        gradeStr= "F (Strong Negative Momentum)"
        gradeColor= "darkred"
    else:
        gradeStr= "N/A"
        gradeColor= "white"
    if len(df)>1:
        directionIncreasing= (momentum.iloc[-1]> momentum.iloc[-2])
        dirStr= "Increasing" if directionIncreasing else "Decreasing"
        dirColor= "green" if directionIncreasing else "red"
        change= momentum.iloc[-1]- momentum.iloc[-2]
        if momentum.iloc[-1]* momentum.iloc[-2]<0:
            stateStr= "Turning Up" if momentum.iloc[-1]>0 else "Turning Down"
            stateColor= "orange"
        elif abs(momentum.iloc[-1])< abs(avgM.iloc[-1])*0.1:
            stateStr= "Consolidating"
            stateColor= "yellow"
        elif (momentum.iloc[-1]>0 and change<0) or (momentum.iloc[-1]<0 and
↪change>0):
            stateStr= "Stalling"
            stateColor= "lightgray"
        elif momentum.iloc[-1]>0:
            stateStr= "Positive Trending"
            stateColor= "green"
        else:
            stateStr= "Negative Trending"
            stateColor= "red"
    else:
        dirStr, dirColor= "N/A","white"
        stateStr, stateColor= "N/A","white"

    axs[0].text(0.01,0.95, f"Momentum Grade: {gradeStr} (Z-Score: {last_z:.
↪2f})",
                transform= axs[0].transAxes, fontsize=10, color= gradeColor,
                bbox=dict(facecolor='white', alpha=0.7, edgecolor='none'))
    axs[0].text(0.01,0.90, f"Momentum Direction: {dirStr}",
                transform= axs[0].transAxes, fontsize=10, color= dirColor,
                bbox=dict(facecolor='white', alpha=0.7, edgecolor='none'))
    axs[0].text(0.01,0.85, f"Momentum State: {stateStr}",
                transform= axs[0].transAxes, fontsize=10, color= stateColor,
                bbox=dict(facecolor='white', alpha=0.7, edgecolor='none'))

    # Panel 2: Bressert
    axs[1].plot(df.index, b_X, label="X (EMA of Y)", color="black", linewidth=2)
    marker_colors= ['black']+ ['red' if b_X.iloc[i]< b_X.iloc[i-1] else 'green']
↪for i in range(1,len(b_X))]
```

```

    axs[1].scatter(df.index, b_X, c=marker_colors, s=20)
    axs[1].plot(df.index, b_DSSb, label="DSSb", color="blue",
↳linewidth=2)
    axs[1].plot(df.index, b_DSSsignal, label="DSSsignal", color="magenta",
↳linewidth=2)
    axs[1].axhline(50, color="gray", linewidth=1)
    axs[1].axhline(80, color="red", linewidth=2)
    axs[1].axhline(20, color="green", linewidth=2)
    axs[1].set_ylabel("Bressert")
    axs[1].legend(loc="lower left", fontsize=8)

# Panel 3: MCDX HBMA & Signals
    axs[2].plot(df.index, hbma, label="HBMA", color="black", linewidth=2,
↳zorder=3)
    axs[2].axhline(threshold, color="gray", linestyle="--", label="Threshold",
↳zorder=2)
    axs[2].scatter(df.index, upSig_MCDX, color="green", marker="o", s=50,
↳label="MCDX Buy", zorder=4)
    axs[2].scatter(df.index, dnSig_MCDX, color="red", marker="o", s=50,
↳label="MCDX Sell", zorder=4)
    axs[2].set_ylabel("MCDX HBMA")
    axs[2].legend(loc="lower left", fontsize=8)

# Panel 4: MCDX Bars
    axs[3].bar(df.index, Dump, width=0.8, color="red", alpha=0.7,
↳label="Dump", zorder=1)
    axs[3].bar(df.index, DnCandle, width=0.8, color="darkgray", alpha=0.7,
↳label="Down Candle", zorder=1)
    axs[3].bar(df.index, PumpCandle, width=0.8, color="green", alpha=0.7,
↳label="Pump Candle", zorder=1)
    axs[3].bar(df.index, Retest, width=0.8, color="darkred", alpha=0.7,
↳label="Retest", zorder=1)
    axs[3].bar(df.index, Banker, width=0.8, color="#84AFC9", alpha=0.7,
↳label="Banker", zorder=1)
    axs[3].set_ylabel("MCDX Bars")
    axs[3].legend(loc="lower left", fontsize=8)

# Panel 5: Zero Lag MACD
    zmacd= zero_macd_dict
    axs[4].fill_between(df.index, zmacd["ZeroLagMACD"], zmacd["signal"],
                        where=(zmacd["ZeroLagMACD"]>= zmacd["signal"]),
                        facecolor="green", alpha=0.3, interpolate=True)
    axs[4].fill_between(df.index, zmacd["ZeroLagMACD"], zmacd["signal"],
                        where=(zmacd["ZeroLagMACD"]< zmacd["signal"]),
                        facecolor="red", alpha=0.3, interpolate=True)

```

```

    axs[4].plot(df.index, zmacd["ZeroLagMACD"], label="ZeroLag MACD",
    ↪color="green", linewidth=1)
    axs[4].plot(df.index, zmacd["signal"], label="ZMACD Signal",
    ↪color="red", linewidth=1)
    axs[4].bar(df.index, zmacd["upHist"]*2, label="ZMACD Hist Up",
    ↪color="gray", width=0.8)
    axs[4].bar(df.index, zmacd["downHist"]*2, label="ZMACD Hist Down",
    ↪color="red", width=0.8)
    axs[4].scatter(df.index, zmacd["dotUP"], color="green", marker="o", s=50,
    ↪label="ZMACD Dot Up")
    axs[4].scatter(df.index, zmacd["dotDN"], color="red", marker="o", s=50,
    ↪label="ZMACD Dot Down")
    axs[4].set_ylabel("ZeroLag MACD")
    axs[4].legend(loc="lower left", fontsize=8)

    # Panel 6: VIX
    axs[5].plot(df.index, iv_series, label="VIX", color="darkorange",
    ↪linewidth=2)
    axs[5].set_ylabel("IV")
    axs[5].legend(loc="lower left", fontsize=8)
    axs[5].xaxis.set_major_formatter(mdates.DateFormatter("%Y-%m-%d"))
    for tick in axs[5].get_xticklabels():
        tick.set_rotation(45)

    # Panel 7: Momentum Z-Score
    zScore_panel= compute_zscore(df, momentum_length)
    axs[6].axhline(0, color="gray", linestyle="--")
    axs[6].fill_between(df.index, zScore_panel, 0, where=(zScore_panel>0),
    ↪facecolor="green", alpha=0.3)
    axs[6].fill_between(df.index, zScore_panel, 0, where=(zScore_panel<0),
    ↪facecolor="red", alpha=0.3)
    axs[6].plot(df.index, zScore_panel, label="Momentum Z-Score",
    ↪color="black", linewidth=1.5)
    axs[6].set_ylabel("Momentum")
    axs[6].legend(loc="lower left", fontsize=8)

    # Panel 8: Basic MACD
    if basic_macd_dict:
        macd_line= basic_macd_dict["basicMACD"]
        sig_line= basic_macd_dict["signal"]
        hist_line= basic_macd_dict["hist"]
        axs[7].fill_between(df.index, macd_line, sig_line,
                           where=(macd_line>= sig_line),
                           facecolor="green", alpha=0.3, interpolate=True)
        axs[7].fill_between(df.index, macd_line, sig_line,
                           where=(macd_line< sig_line),

```



```

        facecolor="red", alpha=0.3, interpolate=True)
    axs[7].plot(df.index, macd_line, label="Basic MACD", color="green",
↳linewidth=1)
    axs[7].plot(df.index, sig_line, label="Basic MACD Signal",
↳color="red", linewidth=1)
    axs[7].bar(df.index, hist_line*2, label="Basic MACD Hist",
↳color="gray", width=0.8)
    axs[7].set_ylabel("Basic MACD")
    axs[7].legend(loc="lower left", fontsize=8)
    else:
        axs[7].text(0.5,0.5, "No Basic MACD Data", ha="center", va="center",
↳transform=axs[7].transAxes)
        axs[7].set_ylabel("Basic MACD")

    plt.tight_layout()
    return fig

def figure_to_pil(fig):
    buf = io.BytesIO()
    fig.savefig(buf, format="png", bbox_inches="tight")
    buf.seek(0)
    return Image.open(buf)

```

```

[22]: # Main chart function and trade log
def generate_plot(
    ticker="SPY",
    start_date="2022-01-01",
    end_date="2023-01-01",

    # Weighted scoreboard
    weight_zlma=1.0,
    weight_rsi=1.0,
    weight_mcdx=1.0,
    weight_dss=1.0,
    weight_zscore=1.0,
    weight_mg=1.0,
    weight_md=1.0,
    weight_ms=1.0,

    # Two MACD weights
    weight_zeromacd=1.0,
    weight_basicmacd=1.0
):
    try:
        df = download_data(ticker, start_date, end_date)
        if df.empty:

```

```

        raise gr.Error(f"No data for {ticker} from {start_date} to {end_date}")
    for col in ["open", "high", "low", "close", "volume"]:
        if col not in df.columns:
            raise gr.Error(f"Missing {col} data for {ticker}")

    # Basic EMAs
    df['EMA_50'] = exp_average(df['close'], 50)
    df['EMA_100'] = exp_average(df['close'], 100)
    df['EMA_200'] = exp_average(df['close'], 200)
    df['EMA_500'] = exp_average(df['close'], 500)

    # ZLMA
    movAvgLength= 15
    ema_value= exp_average(df['close'], movAvgLength)
    df['ema_value']= ema_value
    correction= df['close']+(df['close']-ema_value)
    zlma= exp_average(correction, movAvgLength)
    df['zlma']= zlma

    signalUp_ZLMA= (zlma>ema_value)&(zlma.shift(1)<= ema_value.shift(1))
    signalDn_ZLMA= (zlma<ema_value)&(zlma.shift(1)>= ema_value.shift(1))
    zlma_color= "green" if zlma.iloc[-1]> zlma.iloc[-2] else "red"
    ema_color= "green" if ema_value.iloc[-1]< zlma.iloc[-1] else "red"

    # Bressert
    df= compute_bressert(df,8,13)
    b_X= df['X']
    b_DSSb= df['DSSb']
    b_DSSsignal= df['DSSsignal']

    # MCDX + RSI
    RSIBaseBanker= 50; RSIPeriodBanker= 50
    RSIBaseHotMoney=30; RSIPeriodHotMoney= 40
    SensitivityBanker=1.5; SensitivityHotMoney=0.7
    threshold= 8.5

    rsi_Banker= rsi_function(df['close'], SensitivityBanker,
    RSIPeriodBanker, RSIBaseBanker)
    rsi_HotMoney= rsi_function(df['close'], SensitivityHotMoney,
    RSIPeriodHotMoney, RSIBaseHotMoney)
    hot= rsi_HotMoney
    bank= rsi_Banker
    hotma2= wilder_average(hot,2)
    hotma7= wilder_average(hot,7)
    hotma31= wilder_average(hot,31)
    hotma= exp_average((hotma2*34 + hotma7*33 + hotma31*33)/100,2)

```

```

bankma2= df['close'].rolling(window=2, min_periods=2).mean()
bankma7= exp_average(bank,7)
bankma31=exp_average(bank,31)
bankma= ((bankma2*70 + bankma7*20 + bankma31*10)/100).rolling(window=2,
↪min_periods=2).mean()

banksignal= wilder_average(bankma,4)
df['banksignal']= banksignal

hbAvg=
↪((hot*10)+(hotma*35)+(wilder_average(hotma,2)*15)+(bankma*35)+(banksignal*5))/
↪100

hbma= vwma(hbAvg,2, df['volume'])
df['hbma']= hbma

downtrendsignal= (hotma.shift(1)>= wilder_average(hotma,2).
↪shift(1))&(hotma< wilder_average(hotma,2))
uptrendsignal= (hotma.shift(1)<= wilder_average(hotma,2).
↪shift(1))&(hotma> wilder_average(hotma,2))
upSig_MCDX= hbma.where(uptrendsignal, np.nan)
dnSig_MCDX= hbma.where(downtrendsignal, np.nan)

Dump= bank.where(bank< bank.shift(1)/1.75, np.nan)
dnCond= (bank<bank.shift(1))&(bank<bank.shift(2))&(bank.shift(1)<bank.
↪shift(2))& \
        (bank<bank.shift(3))&(bank<bank.shift(4))&(bank.shift(3)<bank.
↪shift(4))& \
        (bank.shift(6)>8.5)&(bank<10)
DnCandle= bank.where(dnCond, np.nan)
PumpCandle= bank.where(bank> hbma, np.nan)
Retest= bank.where((banksignal> bankma)&(bank>0), np.nan)
Banker= bank

# RSI trail
lookbackPeriod=15
atrLength=27
atrMultiplier=0.1
rsiLowerThreshold=40
rsiUpperThreshold=60
ohlc4= (df['open']+df['high']+df['low']+df['close'])/4
rsi_ma_base= t3(ohlc4, length=lookbackPeriod, vf=0.7)
df['rsi_ma_base']= rsi_ma_base
tr_series= pd.concat([
    df['high']-df['low'],
    abs(df['high']-df['close'].shift(1)),
    abs(df['low']-df['close'].shift(1))
], axis=1).max(axis=1)

```

```

nzTR= tr_series.fillna(df['high']-df['low'])
f_volatility= wilder_average(nzTR,atrLength)* atrMultiplier
rsi_upper_bound= rsi_ma_base+ ((rsiUpperThreshold-50)/10)* f_volatility
rsi_lower_bound= rsi_ma_base- ((50-rsiLowerThreshold)/10)* f_volatility

crossUp= (ohlc4> rsi_upper_bound)&(ohlc4.shift(1)<= rsi_upper_bound.
↪shift(1))
crossDn= (df['close']<rsi_lower_bound)&(df['close'].shift(1)>=
↪rsi_lower_bound.shift(1))
bullPt= rsi_lower_bound.where(crossUp, np.nan)
bearPt= rsi_upper_bound.where(crossDn, np.nan)

# VIX => iv_series
vix_df= yf.download("^VIX", start=pd.to_datetime(start_date), end=pd.
↪to_datetime(end_date))
if vix_df.empty:
    iv_series= pd.Series(np.nan, index=df.index)
else:
    vix_df.index= pd.to_datetime(vix_df.index)
    iv_series= vix_df["Close"].reindex(df.index, method="ffill")

# ZeroLag MACD
zero_macd_dict= compute_zero_lag_macd(df['close'], fastLength=12,
↪slowLength=26, signalLength=9,
MacdEmaLength=9, useEma=True,
↪useOldAlgo=False)

# Basic MACD
# basic_macd_dict= compute_basic_macd(df['close'], fast=12, slow=26,
↪signal=9)
basic_macd_dict= compute_basic_macd(df['close'], fast=30, slow=50,
↪signal=9)

# Create multi-panel figure
fig= create_generic_plot(
    df, ticker, start_date, end_date,
    ema_value, zlma, signalUp_ZLMA, signalDn_ZLMA, zlma_color,
↪ema_color,
    rsi_ma_base, rsi_upper_bound, rsi_lower_bound, bullPt, bearPt,
    b_X, b_DSSb, b_DSSsignal,
    hbma, threshold, upSig_MCDX, dnSig_MCDX,
    Dump, DnCandle, PumpCandle, Retest, Banker,
    iv_series, zero_macd_dict, basic_macd_dict,
    momentum_length=14
)

```

```

    # Build historical signals
    signals_zmacd_df= extract_zero_macd_signals(df, zero_macd_dict,
↪length_m=14)
    signals_bmacd_df= extract_basic_macd_signals(df, basic_macd_dict,
↪length_m=14)
    signals_others_df= extract_signals(df, signalUp_ZLMA, signalDn_ZLMA,
↪bullPt, bearPt,
                                upSig_MCDX, dnSig_MCDX, length_m=14)
    momentum_signals_df= extract_momentum_signals(df, length_m=14)
    historical_signals_df= pd.concat([
        signals_others_df,
        momentum_signals_df,
        signals_zmacd_df,
        signals_bmacd_df
    ], ignore_index=True)
    if not historical_signals_df.empty:
        historical_signals_df["Date"]= pd.
↪to_datetime(historical_signals_df["Date"])
        historical_signals_df= historical_signals_df.sort_values("Date",
↪ascending=False)

    # Current scoreboard
    current_status_df= extract_current_status(
        df, signalUp_ZLMA, signalDn_ZLMA, bullPt, bearPt,
        upSig_MCDX, dnSig_MCDX,
        length_m=14,
        zero_macd_dict= zero_macd_dict,
        basic_macd_dict= basic_macd_dict,
        weight_zlma= weight_zlma,
        weight_rsi= weight_rsi,
        weight_mcdx= weight_mcdx,
        weight_dss= weight_dss,
        weight_zscore= weight_zscore,
        weight_mg= weight_mg,
        weight_md= weight_md,
        weight_ms= weight_ms,
        weight_zeromacd= weight_zeromacd,
        weight_basicmacd= weight_basicmacd
    )

    # Day-by-day scoreboard => trade log
    daily_rsi_series= calc_rsi_entire_series(df, bullPt, bearPt)
    daily_mcdx_series= calc_mcdx_entire_series(df, upSig_MCDX, dnSig_MCDX)
    trade_log_df= simulate_trades_overall_signal(
        df, daily_rsi_series, daily_mcdx_series,
        zero_macd_dict, basic_macd_dict,
        length_m=14,

```

```

        weight_zlma= weight_zlma,
        weight_rsi= weight_rsi,
        weight_mcdx= weight_mcdx,
        weight_dss= weight_dss,
        weight_zscore= weight_zscore,
        weight_mg= weight_mg,
        weight_md= weight_md,
        weight_ms= weight_ms,
        weight_zeromacd= weight_zeromacd,
        weight_basicmacd= weight_basicmacd
    )

    buf= io.BytesIO()
    fig.savefig(buf, format="png", bbox_inches="tight")
    buf.seek(0)
    pil_img= Image.open(buf)
    plt.close(fig)

    return pil_img, current_status_df, historical_signals_df, trade_log_df
except Exception as e:
    debug_print(f"Error: {e}")
    raise gr.Error(f"An error occurred: {e}")

```

```

[23]: def save_historical_data(ticker="SPY", start_date="2022-01-01",
    ↪end_date="2023-01-01",
        data_filename="full_data.csv",
    ↪signals_filename="signals_data.csv"):
    try:
        df = download_data(ticker, start_date, end_date)
        if df.empty:
            raise gr.Error(f"No data for {ticker} from {start_date} to
    ↪{end_date}")

        # --- Basic EMAs and ZLMA ---
        df['EMA_50'] = exp_average(df['close'], 50)
        df['EMA_100'] = exp_average(df['close'], 100)
        df['EMA_200'] = exp_average(df['close'], 200)
        df['EMA_500'] = exp_average(df['close'], 500)

        movAvgLength = 15
        df['ema_value'] = exp_average(df['close'], movAvgLength)
        correction = df['close'] + (df['close'] - df['ema_value'])
        df['zlma'] = exp_average(correction, movAvgLength)

        signalUp_ZLMA = (df['zlma'] > df['ema_value']) & (df['zlma'].shift(1)
    ↪<= df['ema_value'].shift(1))

```

```

    signalDn_ZLMA = (df['zlma'] < df['ema_value']) & (df['zlma'].shift(1)
↪>= df['ema_value'].shift(1))

    # --- Bressert DSS ---
    df = compute_bressert(df, 8, 13)

    # --- MCDX ---
    RSIBaseBanker = 50; RSIPeriodBanker = 50
    RSIBaseHotMoney = 30; RSIPeriodHotMoney = 40
    SensitivityBanker = 1.5; SensitivityHotMoney = 0.7
    rsi_Banker = rsi_function(df['close'], SensitivityBanker,
↪RSIPeriodBanker, RSIBaseBanker)
    rsi_HotMoney = rsi_function(df['close'], SensitivityHotMoney,
↪RSIPeriodHotMoney, RSIBaseHotMoney)

    hot = rsi_HotMoney
    bank = rsi_Banker
    hotma2 = wilder_average(hot, 2)
    hotma7 = wilder_average(hot, 7)
    hotma31 = wilder_average(hot, 31)
    hotma = exp_average((hotma2 * 34 + hotma7 * 33 + hotma31 * 33) / 100, 2)

    bankma2 = df['close'].rolling(window=2).mean()
    bankma7 = exp_average(bank, 7)
    bankma31 = exp_average(bank, 31)
    bankma = ((bankma2 * 70 + bankma7 * 20 + bankma31 * 10) / 100).
↪rolling(window=2).mean()
    banksignal = wilder_average(bankma, 4)
    df['banksignal'] = banksignal

    hbAvg = ((hot * 10) + (hotma * 35) + (wilder_average(hotma, 2) * 15) +
↪(bankma * 35) + (banksignal * 5)) / 100
    df['hbma'] = vwma(hbAvg, 2, df['volume'])

    hotma_slow = wilder_average(hotma, 2)
    uptrendsignal = (hotma.shift(1) <= hotma_slow.shift(1)) & (hotma >
↪hotma_slow)
    downtrendsignal = (hotma.shift(1) >= hotma_slow.shift(1)) & (hotma <
↪hotma_slow)
    upSig_MCDX = df['hbma'].where(uptrendsignal, np.nan)
    dnSig_MCDX = df['hbma'].where(downtrendsignal, np.nan)

    # --- RSI Trail ---
    ohlc4 = (df['open'] + df['high'] + df['low'] + df['close']) / 4
    df['rsi_ma_base'] = t3(ohlc4, length=15, vf=0.7)

```

```

tr_series = pd.concat([
    df['high'] - df['low'],
    abs(df['high'] - df['close'].shift(1)),
    abs(df['low'] - df['close'].shift(1))
], axis=1).max(axis=1)
f_volatility = wilder_average(tr_series.fillna(df['high'] - df['low']),
↪27) * 0.1
rsi_upper_bound = df['rsi_ma_base'] + ((60 - 50) / 10) * f_volatility
rsi_lower_bound = df['rsi_ma_base'] - ((50 - 40) / 10) * f_volatility

crossUp = (ohlc4 > rsi_upper_bound) & (ohlc4.shift(1) <=
↪rsi_upper_bound.shift(1))
crossDn = (df['close'] < rsi_lower_bound) & (df['close'].shift(1) >=
↪rsi_lower_bound.shift(1))
bullPt = rsi_lower_bound.where(crossUp, np.nan)
bearPt = rsi_upper_bound.where(crossDn, np.nan)

# --- ZeroLag MACD & Basic MACD ---
zero_macd_dict = compute_zero_lag_macd(df['close'])
basic_macd_dict = compute_basic_macd(df['close'], fast=30, slow=50,
↪signal=9)

df['ZeroLagMACD'] = zero_macd_dict["ZeroLagMACD"]
df['ZeroLagMACD_signal'] = zero_macd_dict["signal"]
df['basicMACD'] = basic_macd_dict["basicMACD"]
df['basicMACD_signal'] = basic_macd_dict["signal"]

# --- Z-Score ---
df['ZScore'] = compute_zscore(df)

# --- Signal Extraction ---
signals_df = extract_signals(df, signalUp_ZLMA, signalDn_ZLMA, bullPt,
↪bearPt, upSig_MCDX, dnSig_MCDX)
signals_df = pd.concat([
    signals_df,
    extract_zero_macd_signals(df, zero_macd_dict),
    extract_basic_macd_signals(df, basic_macd_dict)
])
signals_df['Date'] = pd.to_datetime(signals_df['Date'])
signals_df = signals_df.sort_values('Date').reset_index(drop=True)
signals_grouped = signals_df.groupby("Date")["Signal"].apply(list)

# --- Persistent Buy/Sell Columns ---
def persistent_state(indicator):
    state = []
    last = 0
    for date in df.index:

```



```

        signal_list = signals_grouped.get(date, [])
        if any(f"{indicator} Buy" in s for s in signal_list):
            last = 1
        elif any(f"{indicator} Sell" in s for s in signal_list):
            last = 0
        state.append(last)
    return pd.Series(state, index=df.index)

indicators = ["ZLMA", "RSI", "MCDX", "DSS", "ZeroLag MACD", "Basic_
↪MACD"]
for ind in indicators:
    df[ind + "_Buy"] = persistent_state(ind)
    df[ind + "_Sell"] = 1 - df[ind + "_Buy"]

# --- Scoreboard Signal (OverallTrade) ---
daily_rsi_series = calc_rsi_entire_series(df, bullPt, bearPt)
daily_mcdx_series = calc_mcdx_entire_series(df, upSig_MCDX, dnSig_MCDX)
df['OverallTrade'] = [
    scoreboard_for_day(df, i, daily_rsi_series, daily_mcdx_series,
                        zero_macd_dict, basic_macd_dict)
    for i in range(len(df))
]

# Save to CSV
df.reset_index().to_csv(data_filename, index=False)
signals_df.to_csv(signals_filename, index=False)
return f"Saved {data_filename} and {signals_filename} successfully."

except Exception as e:
    debug_print(f"Error: {e}")

```

0.4 Entry Point

Calls indicator functions

```

[24]: from datetime import datetime, timedelta

default_end_date = datetime.now().strftime("%Y-%m-%d")
default_start_date = (datetime.now() - timedelta(days=365)).strftime("%Y-%m-%d")

[25]: # Imports
import pandas as pd
from datetime import datetime
import matplotlib.pyplot as plt

# Define inputs
ticker = "SPY"

```

```

start_date = default_start_date
end_date = default_end_date
weights = {
    "weight_zlma": 1.0,
    "weight_rsi": 1.0,
    "weight_mcdx": 0.0,
    "weight_dss": 0.0,
    "weight_zscore": 1.0,
    "weight_mg": 0.0,
    "weight_md": 0.0,
    "weight_ms": 0.0,
    "weight_zeromacd": 1.0,
    "weight_basicmacd": 1.0
}

_, current_status_df, historical_signals_df, trade_log_df = generate_plot(
    ticker=ticker,
    start_date=start_date,
    end_date=end_date,
    **weights
)

df = download_data(ticker, start_date, end_date)
df['EMA_50'] = exp_average(df['close'], 50)
df['EMA_100'] = exp_average(df['close'], 100)
df['EMA_200'] = exp_average(df['close'], 200)
df['EMA_500'] = exp_average(df['close'], 500)

movAvgLength = 15
ema_value = exp_average(df['close'], movAvgLength)
df['ema_value'] = ema_value
correction = df['close'] + (df['close'] - ema_value)
zlma = exp_average(correction, movAvgLength)
df['zlma'] = zlma

signalUp_ZLMA = (zlma > ema_value) & (zlma.shift(1) <= ema_value.shift(1))
signalDn_ZLMA = (zlma < ema_value) & (zlma.shift(1) >= ema_value.shift(1))
zlma_color = "green" if zlma.iloc[-1] > zlma.iloc[-2] else "red"
ema_color = "green" if ema_value.iloc[-1] < zlma.iloc[-1] else "red"

df = compute_bressert(df, 8, 13)
b_X = df['X']
b_DSSb = df['DSSb']
b_DSSsignal = df['DSSsignal']

RSIBaseBanker = 50
RSIPeriodBanker = 50

```

```

RSIBaseHotMoney = 30
RSIPeriodHotMoney = 40
SensitivityBanker = 1.5
SensitivityHotMoney = 0.7
threshold = 8.5

rsi_Banker = rsi_function(df['close'], SensitivityBanker, RSIPeriodBanker,
    ↳RSIBaseBanker)
rsi_HotMoney = rsi_function(df['close'], SensitivityHotMoney,
    ↳RSIPeriodHotMoney, RSIBaseHotMoney)
hot = rsi_HotMoney
bank = rsi_Banker
hotma2 = wilder_average(hot, 2)
hotma7 = wilder_average(hot, 7)
hotma31 = wilder_average(hot, 31)
hotma = exp_average((hotma2 * 34 + hotma7 * 33 + hotma31 * 33) / 100, 2)
bankma2 = df['close'].rolling(window=2).mean()
bankma7 = exp_average(bank, 7)
bankma31 = exp_average(bank, 31)
bankma = ((bankma2 * 70 + bankma7 * 20 + bankma31 * 10) / 100).
    ↳rolling(window=2).mean()
banksignal = wilder_average(bankma, 4)
df['banksignal'] = banksignal
hbAvg = ((hot * 10) + (hotma * 35) + (wilder_average(hotma, 2) * 15) + (bankma
    ↳* 35) + (banksignal * 5)) / 100
hbma = vwma(hbAvg, 2, df['volume'])
df['hbma'] = hbma

hotma_slow = wilder_average(hotma, 2)
uptrendsignal = (hotma.shift(1) <= hotma_slow.shift(1)) & (hotma > hotma_slow)
downtrendsignal = (hotma.shift(1) >= hotma_slow.shift(1)) & (hotma < hotma_slow)
upSig_MCDX = hbma.where(uptrendsignal, np.nan)
dnSig_MCDX = hbma.where(downtrendsignal, np.nan)

Dump = bank.where(bank < bank.shift(1) / 1.75, np.nan)
dnCond = (
    (bank < bank.shift(1)) & (bank < bank.shift(2)) & (bank.shift(1) < bank.
    ↳shift(2)) &
    (bank < bank.shift(3)) & (bank < bank.shift(4)) & (bank.shift(3) < bank.
    ↳shift(4)) &
    (bank.shift(6) > 8.5) & (bank < 10)
)
DnCandle = bank.where(dnCond, np.nan)
PumpCandle = bank.where(bank > hbma, np.nan)
Retest = bank.where((banksignal > bankma) & (bank > 0), np.nan)
Banker = bank

```

```

lookbackPeriod = 15
atrLength = 27
atrMultiplier = 0.1
rsiLowerThreshold = 40
rsiUpperThreshold = 60
ohlc4 = (df['open'] + df['high'] + df['low'] + df['close']) / 4
rsi_ma_base = t3(ohlc4, length=lookbackPeriod, vf=0.7)
df['rsi_ma_base'] = rsi_ma_base
tr_series = pd.concat([
    df['high'] - df['low'],
    abs(df['high'] - df['close'].shift(1)),
    abs(df['low'] - df['close'].shift(1))
], axis=1).max(axis=1)
f_volatility = wilder_average(tr_series.fillna(df['high'] - df['low']),
    ↪atrLength) * atrMultiplier
rsi_upper_bound = rsi_ma_base + ((rsiUpperThreshold - 50) / 10) * f_volatility
rsi_lower_bound = rsi_ma_base - ((50 - rsiLowerThreshold) / 10) * f_volatility
crossUp = (ohlc4 > rsi_upper_bound) & (ohlc4.shift(1) <= rsi_upper_bound.
    ↪shift(1))
crossDn = (df['close'] < rsi_lower_bound) & (df['close'].shift(1) >=
    ↪rsi_lower_bound.shift(1))
bullPt = rsi_lower_bound.where(crossUp, np.nan)
bearPt = rsi_upper_bound.where(crossDn, np.nan)

# VIX
vix_df = yf.download("^VIX", start=pd.to_datetime(start_date), end=pd.
    ↪to_datetime(end_date))
iv_series = vix_df["Close"].reindex(df.index, method="ffill") if not vix_df.
    ↪empty else pd.Series(np.nan, index=df.index)

# MACDs
zero_macd_dict = compute_zero_lag_macd(df['close'])
basic_macd_dict = compute_basic_macd(df['close'], fast=30, slow=50, signal=9)

# Call plot generator
fig = create_generic_plot(
    df, ticker, start_date, end_date,
    ema_value, zlma, signalUp_ZLMA, signalDn_ZLMA, zlma_color, ema_color,
    rsi_ma_base, rsi_upper_bound, rsi_lower_bound, bullPt, bearPt,
    b_X, b_DSSb, b_DSSsignal,
    hbma, threshold, upSig_MCDX, dnSig_MCDX,
    Dump, DnCandle, PumpCandle, Retest, Banker,
    iv_series, zero_macd_dict, basic_macd_dict,
    momentum_length=14
)

```

```
# Save figure as PNG
fig.savefig("chart_output.png", dpi=300)
plt.close(fig)
print("Chart saved to chart_output.png")
```

YF.download() has changed argument auto_adjust default to True

```
[*****100%*****] 1 of 1 completed
[*****100%*****] 1 of 1 completed
[*****100%*****] 1 of 1 completed
[*****100%*****] 1 of 1 completed
```

Chart saved to chart_output.png

```
[26]: # Output current status
print("=== Current Indicator Status ===")
print(current_status_df.to_string(index=False))

# Display top historical signals
print("\n=== Most Recent Historical Signals ===")
print(historical_signals_df.head(10).to_string(index=False))

# Display trade log summary
print("\n=== Most Recent Trades ===")
print(trade_log_df.head(10).to_string(index=False))
```

=== Current Indicator Status ===

Indicator	Current Signal
ZLMA	Sell
RSI	Sell
MCDX	Buy
DSS	Buy
Z-Score	-0.78
Momentum Grade	D
Momentum Direction	Decreasing
Momentum State	Sell
ZeroLag MACD	Buy
Basic MACD	Sell
ScoreSum	-3.0
Overall Trade	Sell

=== Most Recent Historical Signals ===

Date	Signal	Z-Score
2025-04-10	MCDX Buy	-0.78
2025-04-10	Momentum State Changed to Negative Trending	-0.78
2025-04-10	Momentum Grade Changed to D	-0.78
2025-04-09	Momentum Grade Changed to C	0.07
2025-04-09	DSS Buy	0.07
2025-04-09	Momentum State Changed to Stalling	0.07

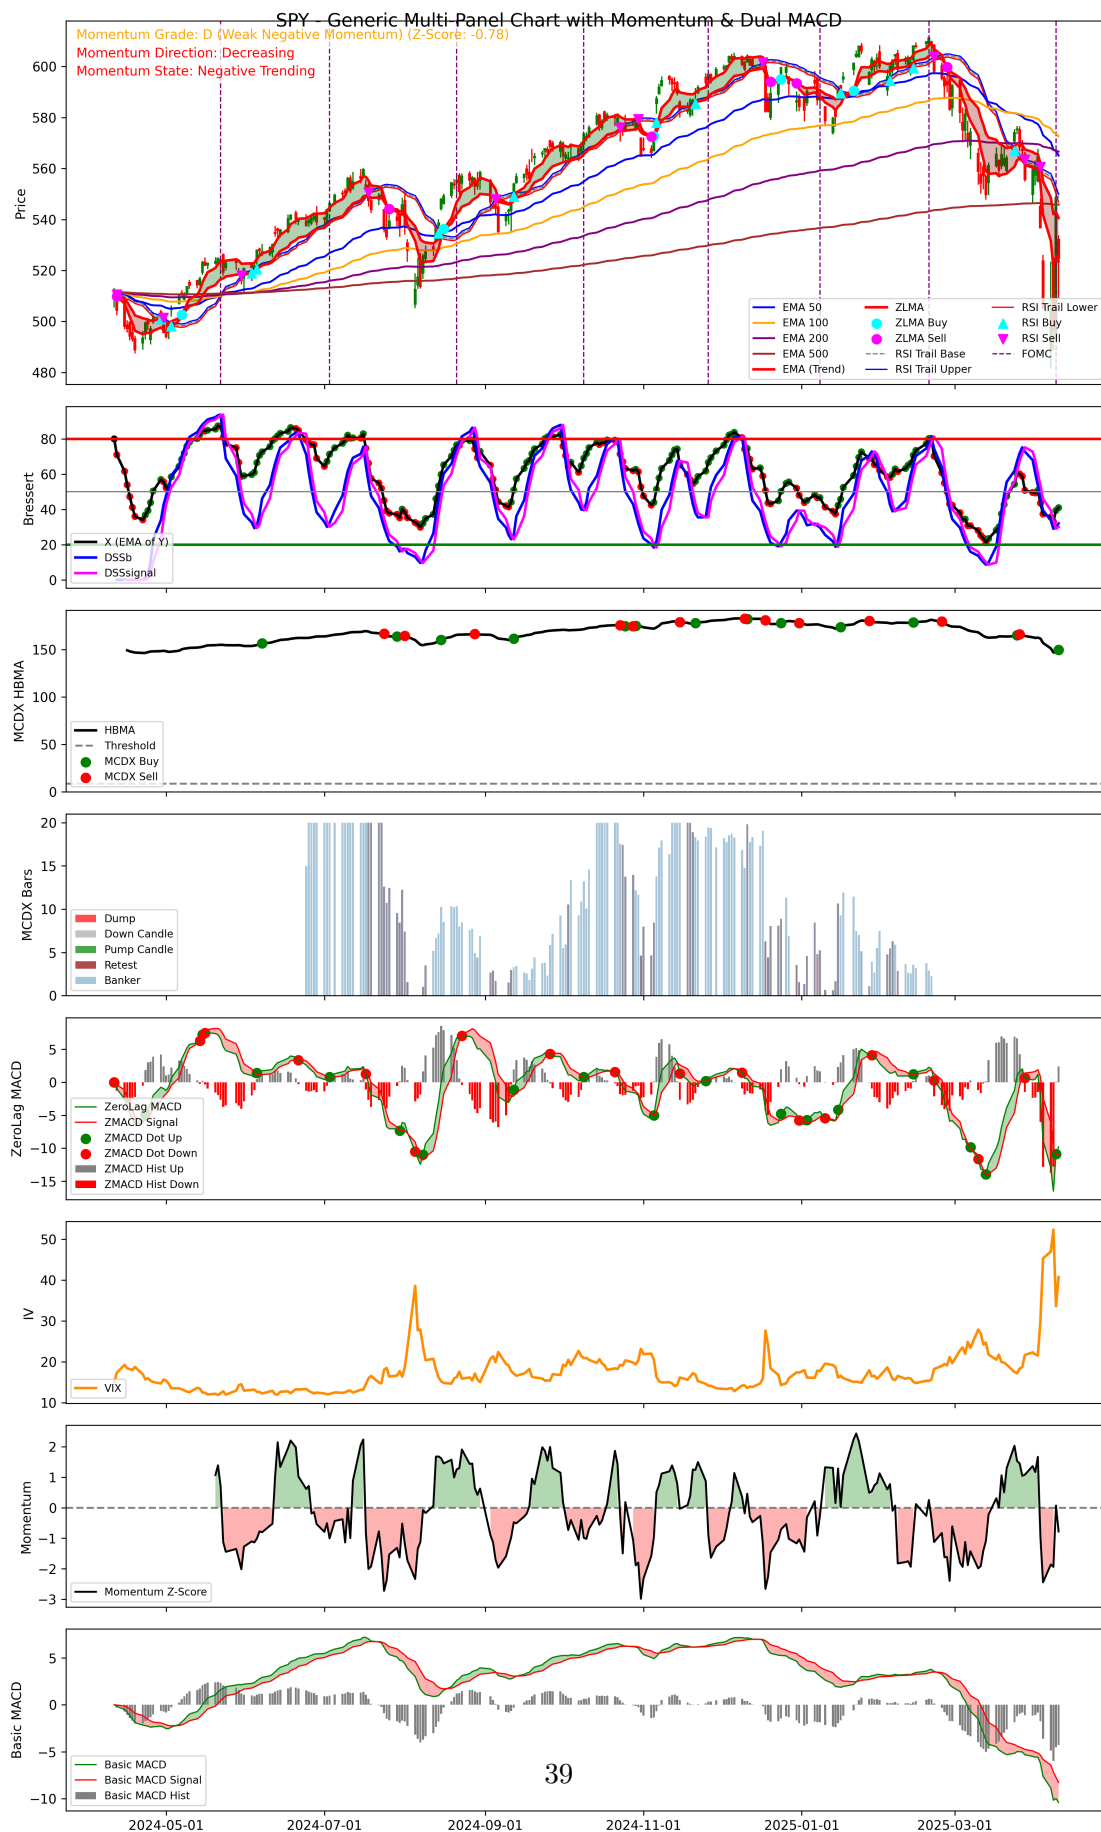
2025-04-09	ZeroLag MACD Buy	-1.19
2025-04-08	Momentum State Changed to Negative Trending	-1.94
2025-04-07	Momentum State Changed to Stalling	-1.86
2025-04-07	Momentum Grade Changed to E	-1.86

=== Most Recent Trades ===

EntryDate	ExitDate	Position	EntryPrice	ExitPrice	PnL%	CumulativePnL%
2025-03-28	2025-04-10	Short	555.6600	524.5800	5.92	20.98
2025-03-24	2025-03-28	Long	574.0800	555.6600	-3.21	15.05
2025-02-21	2025-03-24	Short	598.1407	574.0800	4.19	18.26
2025-01-16	2025-02-21	Long	589.8656	598.1407	1.40	14.07
2024-12-17	2025-01-16	Short	600.4567	589.8656	1.80	12.67
2024-11-06	2024-12-17	Long	587.2908	600.4567	2.24	10.87
2024-10-28	2024-11-06	Short	577.1456	587.2908	-1.73	8.63
2024-10-25	2024-10-28	Long	575.3668	577.1456	0.31	10.36
2024-10-23	2024-10-25	Short	574.3235	575.3668	-0.18	10.05
2024-09-12	2024-10-23	Long	553.8446	574.3235	3.70	10.23

```
[27]: from PIL import Image
      from IPython.display import Image as IPyImage, display

      img = "chart_output.png"
      display(IPyImage(filename=img, width=900))
```



0.5 Save Stuff (For Seed Searching)

Change the ticker you want to search, also required for initialization.

```
[321]: from datetime import datetime, timedelta

# Add +1 day to end date
default_end_date = (datetime.now() + timedelta(days=1)).strftime("%Y-%m-%d")
default_start_date = (datetime.now() - timedelta(days=665)).strftime("%Y-%m-%d")
```

```
[322]: # Import required before calling
from datetime import datetime

# Parameters
ticker = "TSLA"
start_date = default_start_date
end_date = default_end_date
data_filename = "full_data.csv"
signals_filename = "signals_data.csv"

# Call the original function exactly as it exists in your script
save_historical_data(
    ticker=ticker,
    start_date=start_date,
    end_date=end_date,
    data_filename=data_filename,
    signals_filename=signals_filename
)
```

```
[*****100%*****] 1 of 1 completed
```

```
[322]: 'Saved full_data.csv and signals_data.csv successfully.'
```

```
[323]: import pandas as pd

# Preview full data
df_full = pd.read_csv("full_data.csv", parse_dates=["Date"])
print(df_full.head())

# Preview signal data
df_signals = pd.read_csv("signals_data.csv", parse_dates=["Date"])
print(df_signals.head())
```

	Date	close	high	low	open	volume	\
0	2023-06-16	260.540009	263.600006	257.209991	258.920013	167563700	
1	2023-06-20	274.450012	274.750000	261.119995	261.500000	165611200	

2	2023-06-21	259.459991	276.989990	257.779999	275.130005	211797100
3	2023-06-22	264.609985	265.000000	248.250000	250.770004	166875900
4	2023-06-23	256.600006	262.450012	252.800003	259.290009	176584100

	EMA_50	EMA_100	EMA_200	EMA_500	...	RSI_Sell	MCDX_Buy	\
0	260.540009	260.540009	260.540009	260.540009	...	1	0	
1	261.085499	260.815454	260.678417	260.595538	...	0	0	
2	261.021753	260.788613	260.666293	260.591004	...	1	0	
3	261.162468	260.864284	260.705534	260.607048	...	1	0	
4	260.983548	260.779843	260.664683	260.591052	...	1	0	

	MCDX_Sell	DSS_Buy	DSS_Sell	ZeroLag	MACD_Buy	ZeroLag	MACD_Sell	\
0	1	0	1		0		1	
1	1	0	1		1		0	
2	1	0	1		1		0	
3	1	0	1		1		0	
4	1	0	1		0		1	

	Basic	MACD_Buy	Basic	MACD_Sell	OverallTrade
0		0		1	Sell
1		1		0	Buy
2		1		0	Sell
3		1		0	Sell
4		1		0	Sell

[5 rows x 42 columns]

	Date	Signal	Z-Score
0	2023-06-20	Basic MACD Buy	NaN
1	2023-06-20	RSI Buy	NaN
2	2023-06-20	ZLMA Buy	NaN
3	2023-06-20	ZeroLag MACD Buy	NaN
4	2023-06-21	RSI Sell	NaN

1

1.1 Master Trading Env (Start Here)

Halt, run this before continuing.

```
[ ]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import gymnasium as gym
from gymnasium import spaces
from sb3_contrib import RecurrentPPO
from sb3_contrib.ppo_recurrent.policies import MlpLstmPolicy
from stable_baselines3.common.vec_env import DummyVecEnv
import torch.nn as nn
```

```

from functools import partial

SEED = 102022

# --- Load & Clean Data ---
data_path = "full_data.csv"
data = pd.read_csv(data_path, parse_dates=["Date"])

required_cols = [
    'open', 'high', 'low', 'close',
    'zlma', 'ema_value',
    'DSSb', 'DSSsignal',
    'rsi_ma_base',
    'ZeroLagMACD', 'ZeroLagMACD_signal',
    'basicMACD', 'basicMACD_signal',
    'ZScore',
    'ZLMA_Buy', 'ZLMA_Sell',
    'RSI_Buy', 'RSI_Sell',
    'MCDX_Buy', 'MCDX_Sell',
    'DSS_Buy', 'DSS_Sell',
    'ZeroLag MACD_Buy', 'ZeroLag MACD_Sell',
    'Basic MACD_Buy', 'Basic MACD_Sell'
]

missing = [col for col in required_cols if col not in data.columns]
assert not missing, f"Missing required columns: {missing}"

data = data.dropna(subset=required_cols).reset_index(drop=True)

SWITCH_COST = 1.0
TRANSACTION_COST = 0.001

ACTIVATION = torch.nn.Tanh
# ACTIVATION = partial(nn.LeakyReLU, negative_slope=0.01)

policy_kwargs = dict(
    activation_fn=ACTIVATION
)

# --- Main Trading Environment ---
class TradingEnvRL(gym.Env):
    metadata = {'render_modes': ['human']}

    def __init__(self, data, initial_balance=10000, hold_cost=0.02,
                 volatility_window=14, exploration_steps=500,
                 switch_cost=1.0, reentry_threshold=0.01,

```

```

        dynamic_threshold=False, loss_penalty=0.75, drawdown_penalty=5.
↪0, large_loss_threshold=-2.0):

    super().__init__()
    self.loss_penalty = loss_penalty
    self.drawdown_penalty = drawdown_penalty
    self.large_loss_threshold = large_loss_threshold
    self.data = data.reset_index(drop=True).copy()
    self.n_steps = len(self.data)
    self.initial_balance = initial_balance
    self.hold_cost = hold_cost
    self.volatility_window = volatility_window
    self.exploration_steps = exploration_steps
    self.switch_cost = switch_cost
    self.reentry_threshold = reentry_threshold
    self.dynamic_threshold = dynamic_threshold

    self.feature_cols = [
        'open', 'high', 'low', 'close',
        'basicMACD', 'basicMACD_signal',
        'Basic MACD_Buy', 'Basic MACD_Sell'
    ]

    obs_dim = len(self.feature_cols) + 1
    self.observation_space = spaces.Box(low=-np.inf, high=np.inf,
↪shape=(obs_dim,), dtype=np.float32)
    self.action_space = spaces.Discrete(2) # 0=Long, 1=Short

    self._compute_volatility_limit()

    def _compute_volatility_limit(self):
        returns = self.data['close'].pct_change()
        self.data['volatility'] = returns.rolling(self.volatility_window).std()
        self.data['adaptive_hold'] = (10 / (self.data['volatility'] * 100)).
↪clip(lower=3, upper=20).fillna(10).astype(int)

    def reset(self, seed=None, options=None):
        if seed is not None:
            np.random.seed(seed)
            random.seed(seed)
            torch.manual_seed(seed)

        self.current_step = 0
        self.position = 0
        self.entry_price = 0.0
        self.entry_date = None
        self.hold_counter = 0

```

```

self.switch_count = 0
self.balance = self.initial_balance
self.cumulative_pnl = 0.0
self.trade_log = []
self.action_counts = {0: 0, 1: 0}
self.reward_tracker = {0: [], 1: []}
self.consecutive_losses = 0
self.equity_curve = [self.initial_balance]
return self._get_obs(), {}

def _get_obs(self):
    row = self.data.iloc[self.current_step]
    features = row[self.feature_cols].values.astype(np.float32)
    pos_feature = np.array([self.position], dtype=np.float32)
    return np.concatenate([features, pos_feature])

def _force_close(self):
    row = self.data.iloc[self.current_step]
    current_price = float(row['close'])
    current_date = row['Date'].strftime("%Y-%m-%d")

    if self.position == 0:
        return 0.0

    # --- Core Return Logic ---
    trade_pct = ((current_price / self.entry_price - 1) * 100) if self.
↪position == 1 else ((self.entry_price / current_price - 1) * 100)
    pos_str = 'Long' if self.position == 1 else 'Short'
    gross_return = trade_pct / 100
    transaction_cost = TRANSACTION_COST * current_price

    old_balance = self.balance
    self.balance -= transaction_cost
    self.balance *= (1 + gross_return)
    net_profit = self.balance - old_balance

    reward = net_profit

    # --- Track Trade History ---
    self.cumulative_pnl += trade_pct
    compounded_pnl = (self.trade_log[-1]['CompoundedFactor'] * (1 +
↪gross_return)) if self.trade_log else (1 + gross_return)
    compounded_pnl_pct = (compounded_pnl - 1) * 100

    self.trade_log.append({
        'EntryDate': self.entry_date,
        'ExitDate': current_date,

```

```

        'Position': pos_str,
        'EntryPrice': round(self.entry_price, 4),
        'ExitPrice': round(current_price, 4),
        'PnL%': round(trade_pct, 2),
        'CumulativePnL%': round(self.cumulative_pnl, 2),
        'CompoundedFactor': compounded_pnl,
        'CompoundedPnL%': round(compounded_pnl_pct, 2)
    })

    # --- Penalty for Large Loss ---
    if trade_pct < -2.0:
        reward += trade_pct * 2 # Stronger penalty for large loss

    # --- Track & Penalize Consecutive Losses ---
    if trade_pct < 0:
        self.consecutive_losses += 1
        reward -= self.consecutive_losses * 0.75 # Growing penalty
    else:
        self.consecutive_losses = 0

    # --- Optional: Drawdown Penalty ---
    self.equity_curve.append(self.balance)
    max_balance = max(self.equity_curve)
    if max_balance > 0:
        drawdown = (max_balance - self.balance) / max_balance
        reward -= 5.0 * drawdown # Penalize deeper drawdown

    # --- Reset State ---
    self.position = 0
    self.entry_price = 0.0
    self.entry_date = None
    self.hold_counter = 0

    return reward

def step(self, action):
    if isinstance(action, np.ndarray):
        action = int(action.item())

    row = self.data.iloc[self.current_step]
    current_price = float(row['close'])
    current_date = row['Date'].strftime("%Y-%m-%d")

    if self.current_step < self.exploration_steps:
        action = self.action_space.sample()

    self.action_counts[action] += 1

```

```

reward = 0.0
target_position = 1 if action == 0 else -1

# === Holding Same Position ===
if self.position == target_position:
    self.hold_counter += 1
    price_change = (current_price - self.entry_price) / self.entry_price
    step_return = price_change if self.position == 1 else -price_change

    step_reward = np.sign(step_return) * np.sqrt(abs(step_return)) * 10
    step_reward = np.clip(step_reward, -50, 50)
    step_reward -= self.hold_cost
    reward += step_reward
    self.reward_tracker[action].append(step_reward)

# === New Position from Flat ===
elif self.position == 0:
    self.position = target_position
    self.entry_price = current_price
    self.entry_date = current_date
    self.hold_counter = 1
    reward -= TRANSACTION_COST * current_price

# === Switch Position ===
else:
    hold_penalty = max(0, 3 - self.hold_counter) * 5.0
    switch_penalty = self.switch_cost + hold_penalty
    reward += self._force_close()
    reward -= switch_penalty

# === Dynamic Re-entry Threshold ===
if self.current_step > 0:
    prev_close = self.data.iloc[self.current_step - 1]['close']
    recent_return = abs((current_price - prev_close) / prev_close)

    # Volatility-aware reentry threshold
    recent_volatility = row['volatility']
    reentry_threshold = 0.5 * recent_volatility
else:
    recent_return = 0
    reentry_threshold = 0.01 # Fallback

if recent_return > reentry_threshold:
    self.position = target_position
    self.entry_price = current_price
    self.entry_date = current_date
    self.hold_counter = 1

```

```

        reward -= TRANSACTION_COST * current_price

    self.current_step += 1
    terminated = self.current_step >= self.n_steps - 1

    if terminated and self.position != 0:
        reward += self._force_close()

    obs = self._get_obs() if not terminated else np.zeros(self.
↪observation_space.shape, dtype=np.float32)
    return obs, reward, terminated, False, {}

    def render(self):
        print(f"Step: {self.current_step}, Position: {self.position}, Balance: ↪
↪{self.balance:.2f}")

    def save_trade_log(self, filename="trade_log.csv"):
        df = pd.DataFrame(self.trade_log)
        if "CompoundedFactor" in df.columns:
            df = df.drop(columns=["CompoundedFactor"])
        df.to_csv(filename, index=False)

# --- Training ---
env = TradingEnvRL(data, initial_balance=10000)
vec_env = DummyVecEnv([lambda: env])
vec_env.seed(SEED)

model = RecurrentPPO(
    policy=MlpLstmPolicy,
    env=vec_env,
    verbose=1,
    n_steps=64,
    batch_size=32,
    learning_rate=0.01,
    gamma=0.99,
    ent_coef=0.01,
    seed=SEED,
    policy_kwargs=policy_kwargs
)

model.learn(total_timesteps=3000)
print(f"\nTraining Complete")

```

```

[325]: # --- Evaluation ---
obs, _ = env.reset()
state = None
done = False

```

```

total_reward = 0
portfolio_values = []

while not done:
    action, state = model.predict(obs, state=state, deterministic=True)
    obs, reward, done, _, _ = env.step(action)
    total_reward += reward
    current_index = min(env.current_step, len(env.data) - 1)
    current_price = env.data.loc[current_index, 'close']
    unrealized = (current_price - env.entry_price) if env.position == 1 else ↪
    (env.entry_price - current_price) if env.position == -1 else 0.0
    mtm_equity = env.balance + unrealized
    portfolio_values.append(mtm_equity)

env.save_trade_log("trade_log_recurrent.csv")
print(f"\nTotal Reward: {total_reward:.2f}")
print(f"Final Balance: {env.balance:.2f}")
print("Trade Log (Recurrent PPO):")
print(pd.DataFrame(env.trade_log))
print("Action counts:", env.action_counts)
print("Average reward per action:")
for k, v in env.reward_tracker.items():
    mean_r = np.mean(v) if v else 0
    print(f"Action {k} ({'Long' if k==0 else 'Short'}): {mean_r:.4f}")

```

Total Reward: -2763.79

Final Balance: 9550.48

Trade Log (Recurrent PPO):

	EntryDate	ExitDate	Position	EntryPrice	ExitPrice	PnL%	\
0	2023-07-27	2023-07-31	Long	255.71	267.43	4.58	
1	2023-08-01	2023-08-03	Short	261.07	259.32	0.67	
2	2023-08-04	2023-08-07	Short	253.86	251.45	0.96	
3	2023-08-08	2023-08-09	Long	249.70	242.19	-3.01	
4	2023-08-10	2023-08-11	Short	245.34	242.65	1.11	
..	
172	2025-03-26	2025-03-27	Long	272.06	273.13	0.39	
173	2025-03-28	2025-04-03	Short	263.55	267.28	-1.40	
174	2025-04-03	2025-04-04	Long	267.28	239.43	-10.42	
175	2025-04-04	2025-04-08	Short	239.43	221.86	7.92	
176	2025-04-08	2025-04-10	Long	221.86	252.40	13.77	

	CumulativePnL%	CompoundedFactor	CompoundedPnL%
0	4.58	1.045833	4.58
1	5.26	1.052891	5.29
2	6.22	1.062982	6.30
3	3.21	1.031012	3.10
4	4.32	1.042442	4.24

..
172	11.82	0.884307	-11.57
173	10.43	0.871966	-12.80
174	0.01	0.781109	-21.89
175	7.93	0.842968	-15.70
176	21.69	0.959007	-4.10

```
[177 rows x 9 columns]
Action counts: {0: 198, 1: 230}
Average reward per action:
Action 0 (Long): 0.0375
Action 1 (Short): 0.0622
```

```
[326]: # --- Plot Performance with Trade Markers ---
buy_hold_line = [env.initial_balance * (p / data['close'].iloc[0]) for p in
    ↪data['close'].iloc[:len(portfolio_values)]]

date_to_step = {row['Date'].strftime('%Y-%m-%d'): i for i, row in data.iloc[:
    ↪len(portfolio_values)].iterrows()}
entry_points = []
exit_points = []


for trade in env.trade_log:
    entry_step = date_to_step.get(trade['EntryDate'])
    exit_step = date_to_step.get(trade['ExitDate'])
    if entry_step is not None and entry_step < len(portfolio_values):
        entry_points.append((entry_step, portfolio_values[entry_step]))
    if exit_step is not None and exit_step < len(portfolio_values):
        exit_points.append((exit_step, portfolio_values[exit_step]))

plt.figure(figsize=(12, 6))
plt.plot(portfolio_values, label="Recurrent PPO Portfolio", linewidth=2)
plt.plot(buy_hold_line, label="Buy & Hold Benchmark", linestyle="--",
    ↪color="gray")

if entry_points:
    entry_steps, entry_vals = zip(*entry_points)
    plt.scatter(entry_steps, entry_vals, color='green', marker='^', s=80,
    ↪label="Trade Entry")
if exit_points:
    exit_steps, exit_vals = zip(*exit_points)
    plt.scatter(exit_steps, exit_vals, color='red', marker='v', s=80,
    ↪label="Trade Exit")

plt.title("Portfolio Value: Recurrent PPO vs. Buy & Hold")
plt.xlabel("Time Step")
plt.ylabel("Portfolio Value")
```

```
plt.legend()
plt.grid(True)
plt.tight_layout()
plt.savefig("recurrent_ppo_performance.png")
plt.show()
```



1.2 Start Your Seed Search Below Here

1.3 Fast Fixed Seed Search Above & Below (Find Seeds Here, Deterministic Start on Center Seed)

Fixed seed outputs centered around the main seed

```
[ ]: import numpy as np
import pandas as pd
import torch
import random
from stable_baselines3.common.utils import set_random_seed
from sb3_contrib import RecurrentPPO
from sb3_contrib.ppo_recurrent.policies import MlpLstmPolicy
from stable_baselines3.common.vec_env import DummyVecEnv

# --- Sweep Parameters ---
CENTER_SEED = 49274
N_TRIALS = 11
TOTAL_TIMESTEPS = 3000
SEED_RANGE = 50000 # how far above/below to search

# --- Prepare Seed List ---
```

```

np.random.seed(CENTER_SEED)

possible_below = np.arange(max(0, CENTER_SEED - SEED_RANGE), CENTER_SEED)
possible_above = np.arange(CENTER_SEED + 1, CENTER_SEED + SEED_RANGE)

n_below = (N_TRIALS - 1) // 2
n_above = N_TRIALS - 1 - n_below

below_seeds = np.random.choice(possible_below, size=n_below, replace=False)
above_seeds = np.random.choice(possible_above, size=n_above, replace=False)

all_seeds = [CENTER_SEED] + list(below_seeds) + list(above_seeds)
np.random.shuffle(all_seeds)

# --- Result Tracking ---
results = []
best_model = None
best_score = -np.inf
best_seed = None

# --- Sweep Loop ---
for seed in all_seeds:
    print(f"\n--- Training with seed {seed} ---")

    seed = int(seed)
    set_random_seed(seed)
    np.random.seed(seed)
    random.seed(seed)
    torch.manual_seed(seed)

    # Re-init environment and model
    env = TradingEnvRL(data.copy(), initial_balance=10000)
    vec_env = DummyVecEnv([lambda: env])
    vec_env.seed(seed)

    model = RecurrentPPO(
        policy=MlpLstmPolicy,
        env=vec_env,
        seed=seed,
        verbose=0,
        n_steps=64,
        batch_size=32,
        learning_rate=3e-4,
        gamma=0.99,
        ent_coef=0.005
    )

```

```

model.learn(total_timesteps=TOTAL_TIMESTEPS)

# Evaluation
obs, _ = env.reset()
state = None
done = False
total_reward = 0

while not done:
    action, state = model.predict(obs, state=state, deterministic=True)
    obs, reward, done, _, _ = env.step(action)
    total_reward += reward

final_balance = env.balance
print(f"Seed {seed}: Reward = {total_reward:.2f}, Final Balance =_{
↪{final_balance:.2f}")

results.append((seed, total_reward, final_balance))

if total_reward > best_score:
    best_score = total_reward
    best_model = model
    best_seed = seed
    model.save("best_recurrent_model.zip")
    print("Best model updated and saved.")

```

```

[ ]: # --- Save Results ---
results_df = pd.DataFrame(results, columns=["Seed", "Reward", "Final Balance"])
results_df.to_csv("seed_sweep_results.csv", index=False)

best_final_balance = results_df.loc[results_df["Seed"] == best_seed, "Final_
↪Balance"].values[0]

print("\n--- Seed Sweep Complete ---")
print(f"Best Seed: {best_seed}")
print(f"Best Total Reward: {best_score:.2f}")
print(f"Best Final Balance: {best_final_balance:.2f}")

```

```

--- Seed Sweep Complete ---
Best Seed: 17394
Best Total Reward: -11765.74
Best Final Balance: 10379.95

```

1.4 Fixed Seed Search Random (Start Exactly on Center Seed)

Useful once you have a good seed to find better ones

```

[ ]: import numpy as np
import pandas as pd
import torch
import random
from stable_baselines3.common.utils import set_random_seed
from sb3_contrib import RecurrentPPO
from sb3_contrib.ppo_recurrent.policies import MlpLstmPolicy
from stable_baselines3.common.vec_env import DummyVecEnv

# --- Sweep Parameters ---
CENTER_SEED = 112948
N_TRIALS = 11
TOTAL_TIMESTEPS = 3000
SEED_RANGE = 50000

# --- Prepare Seed List ---
possible_below = np.arange(max(0, CENTER_SEED - SEED_RANGE), CENTER_SEED)
possible_above = np.arange(CENTER_SEED + 1, CENTER_SEED + SEED_RANGE)

# Number of additional seeds
n_below = (N_TRIALS - 1) // 2
n_above = N_TRIALS - 1 - n_below

# Sample from below and above ranges
np.random.seed(None)
below_seeds = np.random.choice(possible_below, size=n_below, replace=False)
above_seeds = np.random.choice(possible_above, size=n_above, replace=False)

# Combine all seeds, center first
remaining_seeds = list(below_seeds) + list(above_seeds)
np.random.shuffle(remaining_seeds)
all_seeds = [CENTER_SEED] + remaining_seeds

# --- Result Tracking ---
results = []
best_model = None
best_score = -np.inf
best_seed = None

# --- Sweep Loop ---
for seed in all_seeds:
    print(f"\n--- Training with seed {seed} ---")

    seed = int(seed)
    set_random_seed(seed)
    np.random.seed(seed)
    random.seed(seed)

```

```

torch.manual_seed(seed)

env = TradingEnvRL(data.copy(), initial_balance=10000)
vec_env = DummyVecEnv([lambda: env])
vec_env.seed(seed)

model = RecurrentPPO(
    policy=MlpLstmPolicy,
    env=vec_env,
    seed=seed,
    verbose=1,
    n_steps=64,
    batch_size=32,
    learning_rate=3e-4,
    gamma=0.99,
    ent_coef=0.005
)

model.learn(total_timesteps=TOTAL_TIMESTEPS)

# Evaluation
obs, _ = env.reset()
state = None
done = False
total_reward = 0

while not done:
    action, state = model.predict(obs, state=state, deterministic=True)
    obs, reward, done, _, _ = env.step(action)
    total_reward += reward

final_balance = env.balance
print(f"Seed {seed}: Reward = {total_reward:.2f}, Final Balance = {
↪{final_balance:.2f}")

results.append((seed, total_reward, final_balance))

if total_reward > best_score:
    best_score = total_reward
    best_model = model
    best_seed = seed
    model.save("best_recurrent_model.zip")
    print("Best model updated and saved.")

```

```

[ ]: # --- Save Results ---
results_df = pd.DataFrame(results, columns=["Seed", "Reward", "Final Balance"])
results_df.to_csv("seed_sweep_results.csv", index=False)

```

```

best_final_balance = results_df.loc[results_df["Seed"] == best_seed, "Final_
↪Balance"].values[0]

print("\n--- Seed Sweep Complete ---")
print(f"Best Seed: {best_seed}")
print(f"Best Total Reward: {best_score:.2f}")
print(f"Best Final Balance: {best_final_balance:.2f}")

```

1.5 True Random Seed Search (Actually Random)

Nothing works? Find a new seed

```

[327]: import numpy as np
import pandas as pd
import torch
import random
from stable_baselines3.common.utils import set_random_seed
from sb3_contrib import RecurrentPPO
from sb3_contrib.ppo_recurrent.policies import MlpLstmPolicy
from stable_baselines3.common.vec_env import DummyVecEnv

# --- Sweep Parameters ---
CENTER_SEED = 83819
N_TRIALS = 11
TOTAL_TIMESTEPS = 3000
SEED_RANGE = 50000

# --- Prepare Random Seeds Around Center ---
low = max(0, CENTER_SEED - SEED_RANGE)
high = CENTER_SEED + SEED_RANGE

np.random.seed(None)
random_seeds = np.random.choice(np.arange(low, high), size=N_TRIALS,
↪replace=False)

# --- Result Tracking ---
results = []
best_model = None
best_score = -np.inf
best_seed = None

# --- Sweep Loop ---
for seed in random_seeds:
    print(f"\n--- Training with seed {seed} ---")

    seed = int(seed)

```

```

set_random_seed(seed)
np.random.seed(seed)
random.seed(seed)
torch.manual_seed(seed)

env = TradingEnvRL(data.copy(), initial_balance=10000)
vec_env = DummyVecEnv([lambda: env])
vec_env.seed(seed)

model = RecurrentPPO(
    policy=MlpLstmPolicy,
    env=vec_env,
    seed=seed,
    verbose=0,
    n_steps=64,
    batch_size=32,
    learning_rate=0.01,
    gamma=0.99,
    ent_coef=0.01
)

model.learn(total_timesteps=TOTAL_TIMESTEPS)

# Evaluation
obs, _ = env.reset()
state = None
done = False
total_reward = 0

while not done:
    action, state = model.predict(obs, state=state, deterministic=True)
    obs, reward, done, _, _ = env.step(action)
    total_reward += reward

final_balance = env.balance
print(f"Seed {seed}: Reward = {total_reward:.2f}, Final Balance = {
↪final_balance:.2f}")

# Compute average reward per action
avg_rewards = {}
for action, rewards in env.reward_tracker.items():
    label = "Long" if action == 0 else "Short"
    avg = np.mean(rewards) if rewards else 0.0
    avg_rewards[label] = avg
    print(f"Average reward for action {label}: {avg:.4f}")

results.append((seed, total_reward, final_balance))

```



```

    if total_reward > best_score:
        best_score = total_reward
        best_model = model
        best_seed = seed
        model.save("best_recurrent_model.zip")
        print("Best model updated and saved.")

# --- Save Results ---
results_df = pd.DataFrame(results, columns=["Seed", "Reward", "Final Balance"])
results_df.to_csv("seed_sweep_results.csv", index=False)

best_final_balance = results_df.loc[results_df["Seed"] == best_seed, "Final_Balance"].values[0]

print("\n--- Seed Sweep Complete ---")
print(f"Best Seed: {best_seed}")
print(f"Best Total Reward: {best_score:.2f}")
print(f"Best Final Balance: {best_final_balance:.2f}")

```

--- Training with seed 84511 ---

Seed 84511: Reward = -69.15, Final Balance = 12038.12
 Average reward for action Long: 0.3914
 Average reward for action Short: -0.5380
 Best model updated and saved.

--- Training with seed 97974 ---

Seed 97974: Reward = -9070.66, Final Balance = 3522.30
 Average reward for action Long: 0.0136
 Average reward for action Short: -0.0997

--- Training with seed 64982 ---

Seed 64982: Reward = 7790.18, Final Balance = 19863.73
 Average reward for action Long: -0.2441
 Average reward for action Short: 0.6258
 Best model updated and saved.

--- Training with seed 62977 ---

Seed 62977: Reward = 2570.52, Final Balance = 14840.80
 Average reward for action Long: 0.4955
 Average reward for action Short: -0.5567

--- Training with seed 111581 ---

Seed 111581: Reward = -6663.71, Final Balance = 5529.87
 Average reward for action Long: -0.0897
 Average reward for action Short: -0.0450

```

--- Training with seed 43633 ---
Seed 43633: Reward = -4373.75, Final Balance = 8066.42
Average reward for action Long: -0.0223
Average reward for action Short: 0.1117

--- Training with seed 110692 ---
Seed 110692: Reward = 3013.97, Final Balance = 15078.01
Average reward for action Long: -0.2785
Average reward for action Short: 0.6881

--- Training with seed 65785 ---
Seed 65785: Reward = 2912.61, Final Balance = 15068.03
Average reward for action Long: 0.2839
Average reward for action Short: 0.2509

--- Training with seed 94074 ---
Seed 94074: Reward = -4182.82, Final Balance = 8117.05
Average reward for action Long: -0.0971
Average reward for action Short: -0.4204

--- Training with seed 38615 ---
Seed 38615: Reward = -3628.04, Final Balance = 8626.54
Average reward for action Long: -0.6513
Average reward for action Short: -0.1777

--- Training with seed 109961 ---
Seed 109961: Reward = 10069.41, Final Balance = 22325.22
Average reward for action Long: 0.4063
Average reward for action Short: 0.2677
Best model updated and saved.

--- Seed Sweep Complete ---
Best Seed: 109961
Best Total Reward: 10069.41
Best Final Balance: 22325.22

```

```

[328]: # --- Save Results ---
results_df = pd.DataFrame(results, columns=["Seed", "Reward", "Final Balance"])
results_df.to_csv("seed_sweep_results.csv", index=False)

best_final_balance = results_df.loc[results_df["Seed"] == best_seed, "Final_
↪Balance"].values[0]

print("\n--- Seed Sweep Complete ---")
print(f"Best Seed: {best_seed}")
print(f"Best Total Reward: {best_score:.2f}")
print(f"Best Final Balance: {best_final_balance:.2f}")

```

```
--- Seed Sweep Complete ---  
Best Seed: 109961  
Best Total Reward: 10069.41  
Best Final Balance: 22325.22
```

1.6 Stock Test Results

1.7 Last Training on: 4/10/2025

For best replication results try setting the end date to this.

1.8 GOOG

```
[500]: from datetime import datetime, timedelta  
  
# Add +1 day to end date  
default_end_date = (datetime.now() + timedelta(days=1)).strftime("%Y-%m-%d")  
default_start_date = (datetime.now() - timedelta(days=665)).strftime("%Y-%m-%d")
```

```
[501]: # Parameters  
ticker = "GOOG"  
start_date = default_start_date  
end_date = default_end_date  
data_filename = "full_data.csv"  
signals_filename = "signals_data.csv"  
  
save_historical_data(  
    ticker=ticker,  
    start_date=start_date,  
    end_date=end_date,  
    data_filename=data_filename,  
    signals_filename=signals_filename  
)
```

```
[*****100%*****] 1 of 1 completed
```

```
[501]: 'Saved full_data.csv and signals_data.csv successfully.'
```

```
[502]: import pandas as pd  
  
# Preview full data  
df_full = pd.read_csv("full_data.csv", parse_dates=["Date"])  
print(df_full.head())  
  
# Preview signal data  
df_signals = pd.read_csv("signals_data.csv", parse_dates=["Date"])  
print(df_signals.head())
```

	Date	close	high	low	open	volume	\
0	2023-06-16	123.475159	126.102713	123.206435	126.102713	56686800	
1	2023-06-20	123.266144	124.584902	122.250956	122.952634	22698000	
2	2023-06-21	120.688370	122.828236	120.290254	122.654058	22612000	
3	2023-06-22	123.286057	123.350745	119.036182	120.091190	20781900	
4	2023-06-23	122.440063	122.858089	121.285536	121.464688	29542900	

	EMA_50	EMA_100	EMA_200	EMA_500	...	RSI_Sell	MCDX_Buy	\
0	123.475159	123.475159	123.475159	123.475159	...	1	0	
1	123.466962	123.471020	123.473079	123.474324	...	1	0	
2	123.357998	123.415918	123.445370	123.463203	...	1	0	
3	123.355176	123.413346	123.443785	123.462496	...	1	0	
4	123.319290	123.394073	123.433798	123.458414	...	1	0	

	MCDX_Sell	DSS_Buy	DSS_Sell	ZeroLag	MACD_Buy	ZeroLag	MACD_Sell	\
0	1	0	1		0		1	
1	1	0	1		0		1	
2	1	0	1		0		1	
3	1	1	0		0		1	
4	1	1	0		0		1	

	Basic	MACD_Buy	Basic	MACD_Sell	OverallTrade
0		0		1	Sell
1		0		1	Sell
2		0		1	Sell
3		0		1	Sell
4		0		1	Sell

[5 rows x 42 columns]

	Date	Signal	Z-Score
0	2023-06-20	Basic MACD Sell	NaN
1	2023-06-20	ZLMA Sell	NaN
2	2023-06-20	ZeroLag MACD Sell	NaN
3	2023-06-22	DSS Buy	NaN
4	2023-06-26	DSS Sell	NaN

```
[503]: from stable_baselines3.common.callbacks import BaseCallback
```

```
class RewardTrackingCallback(BaseCallback):
    def __init__(self, verbose=0):
        super().__init__(verbose)
        self.episode_rewards = []
        self.timesteps = []

    def _on_step(self) -> bool:
        if self.locals.get("dones") is not None and any(self.locals["dones"]):
            ep_rew = self.locals["rewards"]
```

```

        self.episode_rewards.append(sum(ep_rew))
        self.timesteps.append(self.num_timesteps)
    return True

```

```

[ ]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import gymnasium as gym
from gymnasium import spaces
from sb3_contrib import RecurrentPPO
from sb3_contrib.ppo_recurrent.policies import MlpLstmPolicy
from stable_baselines3.common.vec_env import DummyVecEnv
import torch
import random

# --- Fixed Seed ---
SEED = 88573
np.random.seed(SEED)
random.seed(SEED)
torch.manual_seed(SEED)

# --- Load & Clean Data ---
data_path = "full_data.csv"
data = pd.read_csv(data_path, parse_dates=["Date"])

required_cols = [
    'open', 'high', 'low', 'close',
    'zlma', 'ema_value',
    'DSSb', 'DSSsignal',
    'rsi_ma_base',
    'ZeroLagMACD', 'ZeroLagMACD_signal',
    'basicMACD', 'basicMACD_signal',
    'ZScore',
    'ZLMA_Buy', 'ZLMA_Sell',
    'RSI_Buy', 'RSI_Sell',
    'MCDX_Buy', 'MCDX_Sell',
    'DSS_Buy', 'DSS_Sell',
    'ZeroLag MACD_Buy', 'ZeroLag MACD_Sell',
    'Basic MACD_Buy', 'Basic MACD_Sell'
]

missing = [col for col in required_cols if col not in data.columns]
assert not missing, f"Missing required columns: {missing}"

data = data.dropna(subset=required_cols).reset_index(drop=True)

SWITCH_COST = 1.0

```

```

TRANSACTION_COST = 0.001

ACTIVATION = torch.nn.Tanh
# ACTIVATION = partial(nn.LeakyReLU, negative_slope=0.01)

policy_kwargs = dict(
    activation_fn=ACTIVATION
)

# --- Main Trading Environment ---
class TradingEnvRL(gym.Env):
    metadata = {'render_modes': ['human']}

    def __init__(self, data, initial_balance=10000, hold_cost=0.02,
                 volatility_window=14, exploration_steps=500,
                 switch_cost=1.0, reentry_threshold=0.01,
                 dynamic_threshold=False, loss_penalty=0.75, drawdown_penalty=5.
                 ↪0, large_loss_threshold=-2.0):

        super().__init__()
        self.loss_penalty = loss_penalty
        self.drawdown_penalty = drawdown_penalty
        self.large_loss_threshold = large_loss_threshold
        self.data = data.reset_index(drop=True).copy()
        self.n_steps = len(self.data)
        self.initial_balance = initial_balance
        self.hold_cost = hold_cost
        self.volatility_window = volatility_window
        self.exploration_steps = exploration_steps
        self.switch_cost = switch_cost
        self.reentry_threshold = reentry_threshold
        self.dynamic_threshold = dynamic_threshold
        self.step_rewards = []

        self.feature_cols = [
            'open', 'high', 'low', 'close',
            'basicMACD', 'basicMACD_signal',
            'Basic MACD_Buy', 'Basic MACD_Sell'
        ]

        obs_dim = len(self.feature_cols) + 1
        self.observation_space = spaces.Box(low=-np.inf, high=np.inf,
        ↪shape=(obs_dim,), dtype=np.float32)
        self.action_space = spaces.Discrete(2) # 0=Long, 1=Short

        self._compute_volatility_limit()

```

```

def _compute_volatility_limit(self):
    returns = self.data['close'].pct_change()
    self.data['volatility'] = returns.rolling(self.volatility_window).std()
    self.data['adaptive_hold'] = (10 / (self.data['volatility'] * 100)).
↳ clip(lower=3, upper=20).fillna(10).astype(int)

def reset(self, seed=None, options=None):
    if seed is not None:
        np.random.seed(seed)
        random.seed(seed)
        torch.manual_seed(seed)

    self.current_step = 0
    self.position = 0
    self.entry_price = 0.0
    self.entry_date = None
    self.hold_counter = 0
    self.switch_count = 0
    self.balance = self.initial_balance
    self.cumulative_pnl = 0.0
    self.trade_log = []
    self.action_counts = {0: 0, 1: 0}
    self.reward_tracker = {0: [], 1: []}
    self.consecutive_losses = 0
    self.equity_curve = [self.initial_balance]
    self.step_rewards = []
    return self._get_obs(), {}

def _get_obs(self):
    row = self.data.iloc[self.current_step]
    features = row[self.feature_cols].values.astype(np.float32)
    pos_feature = np.array([self.position], dtype=np.float32)
    return np.concatenate([features, pos_feature])

def _force_close(self):
    row = self.data.iloc[self.current_step]
    current_price = float(row['close'])
    current_date = row['Date'].strftime("%Y-%m-%d")

    if self.position == 0:
        return 0.0

    # --- Core Return Logic ---
    trade_pct = ((current_price / self.entry_price - 1) * 100) if self.
↳ position == 1 else ((self.entry_price / current_price - 1) * 100)
    pos_str = 'Long' if self.position == 1 else 'Short'
    gross_return = trade_pct / 100

```

```

transaction_cost = TRANSACTION_COST * current_price

old_balance = self.balance
self.balance -= transaction_cost
self.balance *= (1 + gross_return)
net_profit = self.balance - old_balance

reward = net_profit

# --- Track Trade History ---
self.cumulative_pnl += trade_pct
compounded_pnl = (self.trade_log[-1]['CompoundedFactor'] * (1 +
↪gross_return)) if self.trade_log else (1 + gross_return)
compounded_pnl_pct = (compounded_pnl - 1) * 100

self.trade_log.append({
    'EntryDate': self.entry_date,
    'ExitDate': current_date,
    'Position': pos_str,
    'EntryPrice': round(self.entry_price, 4),
    'ExitPrice': round(current_price, 4),
    'PnL%': round(trade_pct, 2),
    'CumulativePnL%': round(self.cumulative_pnl, 2),
    'CompoundedFactor': compounded_pnl,
    'CompoundedPnL%': round(compounded_pnl_pct, 2)
})

# --- Penalty for Large Loss ---
if trade_pct < -2.0:
    reward += trade_pct * 2 # Stronger penalty for large loss

# --- Track & Penalize Consecutive Losses ---
if trade_pct < 0:
    self.consecutive_losses += 1
    reward -= self.consecutive_losses * 0.75 # Growing penalty
else:
    self.consecutive_losses = 0

# --- Optional: Drawdown Penalty ---
self.equity_curve.append(self.balance)
max_balance = max(self.equity_curve)
if max_balance > 0:
    drawdown = (max_balance - self.balance) / max_balance
    reward -= 5.0 * drawdown # Penalize deeper drawdown

# --- Reset State ---
self.position = 0

```



```

self.entry_price = 0.0
self.entry_date = None
self.hold_counter = 0

return reward

def step(self, action):
    if isinstance(action, np.ndarray):
        action = int(action.item())

    row = self.data.iloc[self.current_step]
    current_price = float(row['close'])
    current_date = row['Date'].strftime("%Y-%m-%d")

    if self.current_step < self.exploration_steps:
        action = self.action_space.sample()

    self.action_counts[action] += 1
    reward = 0.0
    target_position = 1 if action == 0 else -1

    # === Holding Same Position ===
    if self.position == target_position:
        self.hold_counter += 1
        price_change = (current_price - self.entry_price) / self.entry_price
        step_return = price_change if self.position == 1 else -price_change

        step_reward = np.sign(step_return) * np.sqrt(abs(step_return)) * 10
        step_reward = np.clip(step_reward, -50, 50)
        step_reward -= self.hold_cost
        reward += step_reward
        self.reward_tracker[action].append(step_reward)
        self.step_rewards.append(reward)

    # === New Position from Flat ===
    elif self.position == 0:
        self.position = target_position
        self.entry_price = current_price
        self.entry_date = current_date
        self.hold_counter = 1
        reward -= TRANSACTION_COST * current_price

    # === Switch Position ===
    else:
        hold_penalty = max(0, 3 - self.hold_counter) * 5.0
        switch_penalty = self.switch_cost + hold_penalty
        reward += self._force_close()

```

```

reward -= switch_penalty

# === Dynamic Re-entry Threshold ===
if self.current_step > 0:
    prev_close = self.data.iloc[self.current_step - 1]['close']
    recent_return = abs((current_price - prev_close) / prev_close)

    # Volatility-aware reentry threshold
    recent_volatility = row['volatility']
    reentry_threshold = 0.5 * recent_volatility
else:
    recent_return = 0
    reentry_threshold = 0.01 # Fallback

if recent_return > reentry_threshold:
    self.position = target_position
    self.entry_price = current_price
    self.entry_date = current_date
    self.hold_counter = 1
    reward -= TRANSACTION_COST * current_price

self.current_step += 1
terminated = self.current_step >= self.n_steps - 1

if terminated and self.position != 0:
    reward += self._force_close()

obs = self._get_obs() if not terminated else np.zeros(self.
↪observation_space.shape, dtype=np.float32)
return obs, reward, terminated, False, {}

def render(self):
    print(f"Step: {self.current_step}, Position: {self.position}, Balance: {self.balance:.2f}")
↪{self.balance:.2f}")

def save_trade_log(self, filename="trade_log.csv"):
    df = pd.DataFrame(self.trade_log)
    if "CompoundedFactor" in df.columns:
        df = df.drop(columns=["CompoundedFactor"])
    df.to_csv(filename, index=False)

# --- Training ---
env = TradingEnvRL(data, initial_balance=10000)
vec_env = DummyVecEnv([lambda: env])
vec_env.seed(SEED)

model = RecurrentPPO(

```

```

    policy=MlpLstmPolicy,
    env=vec_env,
    verbose=1,
    n_steps=64,
    batch_size=32,
    learning_rate=0.01,
    gamma=0.99,
    ent_coef=0.01,
    seed=SEED,
    policy_kwargs=policy_kwargs
)

callback = RewardTrackingCallback()

model.learn(total_timesteps=3000, callback=callback)
print(f"\nTraining Complete")

```

```

[505]: # --- Evaluation ---
obs, _ = env.reset()
state = None
done = False
total_reward = 0
portfolio_values = []
final_action = None

while not done:
    action, state = model.predict(obs, state=state, deterministic=True)
    final_action = action
    obs, reward, done, _, _ = env.step(action)
    total_reward += reward

    current_index = min(env.current_step, len(env.data) - 1)
    current_price = env.data.loc[current_index, 'close']
    unrealized = (
        (current_price - env.entry_price) if env.position == 1 else
        (env.entry_price - current_price) if env.position == -1 else
        0.0
    )
    mtm_equity = env.balance + unrealized
    portfolio_values.append(mtm_equity)

# --- Save trade log ---
env.save_trade_log("trade_log_recurrent.csv")

# --- Final Model Signal ---
signal_str = "BUY" if final_action == 0 else "SELL"
latest_date = env.data['Date'].iloc[env.current_step - 1].strftime("%Y-%m-%d")

```

```

print(f"\nLatest model signal at {latest_date}: {signal_str}")

# --- Summary Output ---
print(f"Total Reward: {total_reward:.2f}")
print(f"Final Balance: {env.balance:.2f}")
print("Trade Log (Recurrent PP0):")
print(pd.DataFrame(env.trade_log))
print("Action counts:", env.action_counts)
print("Average reward per action:")
for k, v in env.reward_tracker.items():
    mean_r = np.mean(v) if v else 0
    print(f"Action {k} ({'Long' if k==0 else 'Short'}): {mean_r:.4f}")

```

Latest model signal at 2025-04-09: BUY

Total Reward: 3402.04

Final Balance: 15100.18

Trade Log (Recurrent PP0):

	EntryDate	ExitDate	Position	EntryPrice	ExitPrice	PnL%	\
0	2023-07-27	2023-08-02	Short	129.2578	128.0336	0.96	
1	2023-08-03	2023-08-04	Long	128.1630	127.9340	-0.18	
2	2023-08-07	2023-08-09	Long	131.3180	129.5365	-1.36	
3	2023-08-10	2023-08-11	Long	129.5962	129.5564	-0.03	
4	2023-08-14	2023-08-15	Long	131.2085	129.6559	-1.18	
..	
178	2025-03-27	2025-03-28	Short	164.0800	156.0600	5.14	
179	2025-03-28	2025-03-31	Long	156.0600	156.2300	0.11	
180	2025-04-01	2025-04-03	Short	158.8800	152.6300	4.09	
181	2025-04-03	2025-04-07	Long	152.6300	149.2400	-2.22	
182	2025-04-08	2025-04-10	Short	146.5800	155.3700	-5.66	

	CumulativePnL%	CompoundedFactor	CompoundedPnL%
0	0.96	1.009562	0.96
1	0.78	1.007758	0.78
2	-0.58	0.994086	-0.59
3	-0.61	0.993781	-0.62
4	-1.79	0.982021	-1.80
..
178	49.72	1.574491	57.45
179	49.83	1.576206	57.62
180	53.93	1.640750	64.07
181	51.71	1.604308	60.43
182	46.05	1.513545	51.35

[183 rows x 9 columns]

Action counts: {0: 210, 1: 218}

Average reward per action:

Action 0 (Long): 0.0169

Action 1 (Short): -0.1150

```
[506]: # --- Get Latest Signal from Model ---
last_index = len(env.data) - 1
env.current_step = last_index
obs = env._get_obs()

obs_input = obs[np.newaxis, :]
action, _ = model.predict(obs_input, deterministic=True)

if action == 0:
    current_signal = "BUY"
else:
    current_signal = "SELL"

print(f"\nLatest model signal at {env.data.iloc[last_index]['Date'].date()}:
↳{current_signal}")
```

Latest model signal at 2025-04-10: BUY

```
[507]: # --- Plot Performance with Trade Markers ---
buy_hold_line = [env.initial_balance * (p / data['close'].iloc[0]) for p in
↳data['close'].iloc[:len(portfolio_values)]]

date_to_step = {row['Date'].strftime('%Y-%m-%d'): i for i, row in data.iloc[:
↳len(portfolio_values)].iterrows()}
entry_points = []
exit_points = []

for trade in env.trade_log:
    entry_step = date_to_step.get(trade['EntryDate'])
    exit_step = date_to_step.get(trade['ExitDate'])
    if entry_step is not None and entry_step < len(portfolio_values):
        entry_points.append((entry_step, portfolio_values[entry_step]))
    if exit_step is not None and exit_step < len(portfolio_values):
        exit_points.append((exit_step, portfolio_values[exit_step]))

plt.figure(figsize=(12, 6))
plt.plot(portfolio_values, label="Recurrent PPO Portfolio", linewidth=2)
plt.plot(buy_hold_line, label="Buy & Hold Benchmark", linestyle="--",
↳color="gray")


if entry_points:
    entry_steps, entry_vals = zip(*entry_points)
    plt.scatter(entry_steps, entry_vals, color='green', marker='^', s=80,
↳label="Trade Entry")
if exit_points:
```

```

        exit_steps, exit_vals = zip(*exit_points)
        plt.scatter(exit_steps, exit_vals, color='red', marker='v', s=80,
        ↪label="Trade Exit")

plt.title("Portfolio Value: Recurrent PPO vs. Buy & Hold")
plt.xlabel("Time Step")
plt.ylabel("Portfolio Value")
plt.legend()
plt.grid(True)
plt.tight_layout()
plt.savefig("recurrent_ppo_performance.png")
plt.show()

```




```

[508]: import matplotlib.pyplot as plt

plt.figure(figsize=(10, 5))
plt.plot(callback.timesteps, callback.episode_rewards, marker='o')
plt.title("Episode Reward Over Time")
plt.xlabel("Timesteps")
plt.ylabel("Episode Reward")
plt.grid(True)
plt.tight_layout()
plt.show()

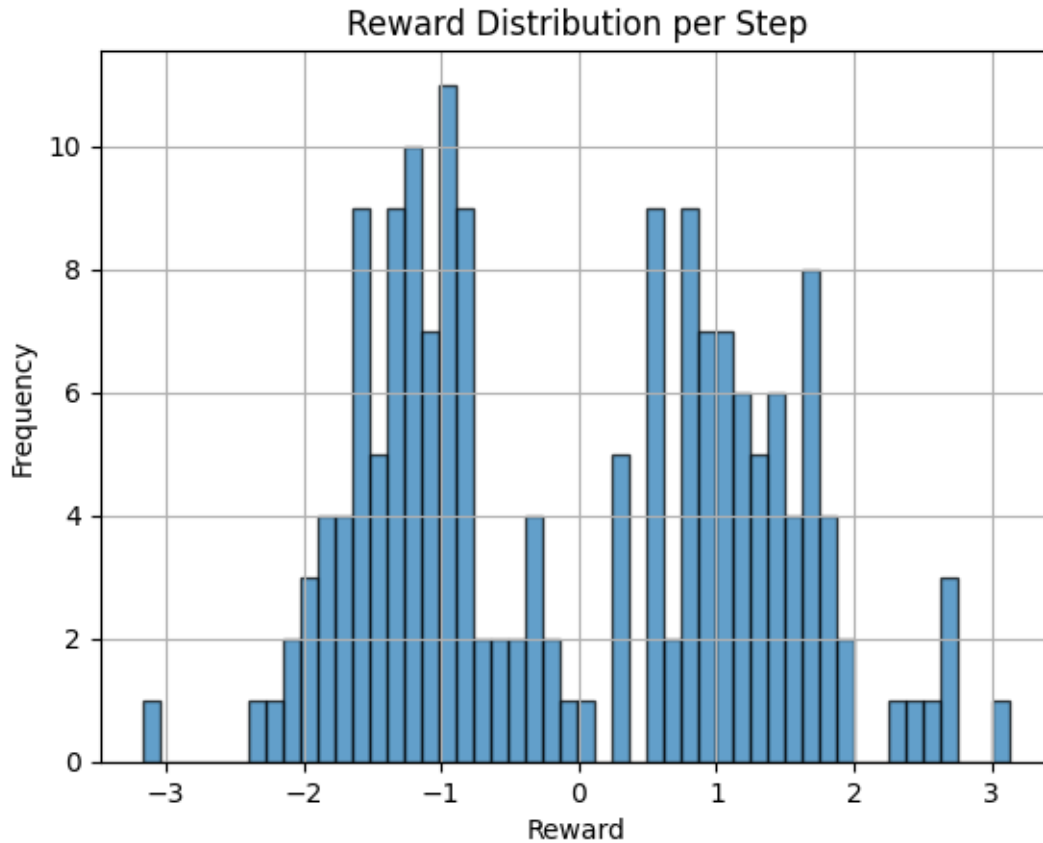
```



```
[509]: import matplotlib.pyplot as plt
import numpy as np
import pandas as pd

rewards = np.array(env.step_rewards)

# Histogram
plt.hist(rewards, bins=50, alpha=0.7, edgecolor='black')
plt.title("Reward Distribution per Step")
plt.xlabel("Reward")
plt.ylabel("Frequency")
plt.grid(True)
plt.show()
```



```
[510]: import numpy as np
import matplotlib.pyplot as plt
import pandas as pd

rewards = np.array(env.step_rewards)

# --- Raw volatility metrics ---
reward_std = np.std(rewards)
reward_mean = np.mean(rewards)
reward_variance = np.var(rewards)
reward_range = np.max(rewards) - np.min(rewards)

# --- Rolling volatility ---
window = 20
rolling_std = pd.Series(rewards).rolling(window=window).std()

# --- Plot ---
plt.figure(figsize=(10, 4))
plt.plot(rolling_std, label=f"Rolling Std (window={window})")
```




```

plt.axhline(reward_std, color='red', linestyle='--', label='Global Std')
plt.title("Reward Volatility (Rolling Std Dev of Step Rewards)")
plt.xlabel("Step")
plt.ylabel("Rolling Std Dev")
plt.legend()
plt.grid(True)
plt.tight_layout()
plt.show()

# --- Output stats ---
print(f"Mean Step Reward: {reward_mean:.4f}")
print(f"Std Dev of Step Rewards: {reward_std:.4f}")
print(f"Variance: {reward_variance:.4f}")
print(f"Range: {reward_range:.4f}")
print(f"Max Rolling Std ({window}): {rolling_std.max():.4f}")
print(f"Min Rolling Std ({window}): {rolling_std.min():.4f}")

```



```

Mean Step Reward: -0.0502
Std Dev of Step Rewards: 1.3521
Variance: 1.8281
Range: 6.2998
Max Rolling Std (20): 1.7136
Min Rolling Std (20): 0.8100

```

```

[511]: import pandas as pd
import matplotlib.pyplot as plt

def compute_switch_rate(trade_log: pd.DataFrame, plot: bool = True):
    if 'Position' not in trade_log.columns or trade_log.empty:
        raise ValueError("Trade log must contain a 'Position' column and be
        non-empty.")

```

```

positions = trade_log['Position'].tolist()
if len(positions) < 2:
    print("Insufficient trades to compute switching behavior.")
    return 0.0

switches = sum(1 for i in range(1, len(positions)) if positions[i] !=
↳positions[i - 1])
total_transitions = len(positions) - 1
switch_rate = switches / total_transitions

print(f"Total Trades: {len(positions)}")
print(f"Total Transitions: {total_transitions}")
print(f"Switches: {switches}")
print(f"Switch Rate: {switch_rate:.4f}")

if plot:
    plt.bar(['Same', 'Switch'], [total_transitions - switches, switches],
↳color=['gray', 'red'])
    plt.title("Trade Direction Transitions")
    plt.ylabel("Count")
    plt.grid(True, axis='y', linestyle='--', alpha=0.6)
    plt.show()

    return switch_rate


df_trades = pd.DataFrame(env.trade_log)
switch_rate = compute_switch_rate(df_trades)

```

```

Total Trades: 183
Total Transitions: 182
Switches: 142
Switch Rate: 0.7802

```



```
[512]: import matplotlib.pyplot as plt

total_actions = sum(env.action_counts.values())

actions = list(env.action_counts.keys())
counts = [env.action_counts[a] for a in actions]
labels = ['Long (0)', 'Short (1)']

proportions = [count / total_actions for count in counts]

plt.bar(labels, proportions, color='gray', edgecolor='black')
plt.title("Action Distribution (Policy Bias)")
plt.ylabel("Proportion of Actions")
plt.ylim(0, 1)
plt.grid(True, axis='y')
plt.show()

for label, count, prop in zip(labels, counts, proportions):
    print(f"{label}: Count = {count}, Proportion = {prop:.2%}")
```



Long (0): Count = 210, Proportion = 49.07%

Short (1): Count = 218, Proportion = 50.93%

— Training with seed 88573 — Seed 88573: Reward = 3402.05, Final Balance = 15100.19 Average reward for action Long: 0.0169 Average reward for action Short: -0.1150 Best model updated and saved.

— Training with seed 65343 — Seed 65343: Reward = -4979.90, Final Balance = 6706.99 Average reward for action Long: 0.2866 Average reward for action Short: -0.0512

— Training with seed 44347 — Seed 44347: Reward = 3249.87, Final Balance = 14809.75 Average reward for action Long: 0.4132 Average reward for action Short: 0.3986

— Training with seed 75599 — Seed 75599: Reward = 2606.58, Final Balance = 14008.65 Average reward for action Long: 0.4323 Average reward for action Short: 0.2494

— Training with seed 81024 — Seed 81024: Reward = -6856.40, Final Balance = 5296.91 Average reward for action Long: -0.0409 Average reward for action Short: -0.4090

— Training with seed 78092 — Seed 78092: Reward = 3352.69, Final Balance = 15044.88 Average reward for action Long: 0.7192 Average reward for action Short: 0.1192

— Training with seed 72471 — Seed 72471: Reward = 732.29, Final Balance = 12317.76 Average reward for action Long: 0.1188 Average reward for action Short: 0.1239

— Training with seed 112103 — Seed 112103: Reward = -6388.21, Final Balance = 5675.00 Average reward for action Long: -0.2011 Average reward for action Short: -0.3009

— Training with seed 57977 — Seed 57977: Reward = -4607.66, Final Balance = 7337.97 Average reward for action Long: -0.0022 Average reward for action Short: 0.1151

— Training with seed 38635 — Seed 38635: Reward = -3719.01, Final Balance = 8201.49 Average reward for action Long: 0.4568 Average reward for action Short: -0.0798

— Training with seed 90866 — Seed 90866: Reward = -2050.35, Final Balance = 9696.36 Average reward for action Long: 0.1171 Average reward for action Short: 0.0206

— Seed Sweep Complete — Best Seed: 88573 Best Total Reward: 3402.05 Best Final Balance: 15100.19

```
[513]: model.save("GOOG_best_model")
```

1.9 AMZN

```
[514]: from datetime import datetime, timedelta

# Add +1 day to end date
default_end_date = (datetime.now() + timedelta(days=1)).strftime("%Y-%m-%d")
default_start_date = (datetime.now() - timedelta(days=665)).strftime("%Y-%m-%d")
```

```
[515]: # Parameters
ticker = "AMZN"
start_date = default_start_date
end_date = default_end_date
data_filename = "full_data.csv"
signals_filename = "signals_data.csv"

save_historical_data(
    ticker=ticker,
    start_date=start_date,
    end_date=end_date,
    data_filename=data_filename,
    signals_filename=signals_filename
)
```

```
[*****100%*****] 1 of 1 completed
```

```
[515]: 'Saved full_data.csv and signals_data.csv successfully.'
```

```
[516]: import pandas as pd

# Preview full data
df_full = pd.read_csv("full_data.csv", parse_dates=["Date"])
print(df_full.head())
```

```
# Preview signal data
```

```
df_signals = pd.read_csv("signals_data.csv", parse_dates=["Date"])
print(df_signals.head())
```

	Date	close	high	low	open	volume	\
0	2023-06-16	125.489998	127.900002	125.300003	127.709999	84188100	
1	2023-06-20	125.779999	127.250000	124.500000	124.970001	56930100	
2	2023-06-21	124.830002	126.730003	123.849998	125.639999	52137700	
3	2023-06-22	130.149994	130.330002	125.139999	125.309998	90354600	
4	2023-06-23	129.330002	130.839996	128.279999	129.110001	71855200	

	EMA_50	EMA_100	EMA_200	EMA_500	...	RSI_Sell	MCDX_Buy	\
0	125.489998	125.489998	125.489998	125.489998	...	1	0	
1	125.501370	125.495740	125.492883	125.491156	...	1	0	
2	125.475042	125.482558	125.486288	125.488516	...	1	0	
3	125.658374	125.574982	125.532693	125.507125	...	0	0	
4	125.802359	125.649339	125.570477	125.522386	...	0	0	

	MCDX_Sell	DSS_Buy	DSS_Sell	ZeroLag	MACD_Buy	ZeroLag	MACD_Sell	\
0	1	0	1		0		1	
1	1	0	1		1		0	
2	1	0	1		0		1	
3	1	0	1		1		0	
4	1	0	1		1		0	

	Basic	MACD_Buy	Basic	MACD_Sell	OverallTrade
0		0		1	Sell
1		1		0	Sell
2		0		1	Sell
3		1		0	Buy
4		1		0	Buy

```
[5 rows x 42 columns]
```

	Date	Signal	Z-Score
0	2023-06-20	Basic MACD Buy	NaN
1	2023-06-20	ZLMA Buy	NaN
2	2023-06-20	ZeroLag MACD Buy	NaN
3	2023-06-21	ZLMA Sell	NaN
4	2023-06-21	ZeroLag MACD Sell	NaN

```
[517]: from stable_baselines3.common.callbacks import BaseCallback
```

```
class RewardTrackingCallback(BaseCallback):
    def __init__(self, verbose=0):
        super().__init__(verbose)
        self.episode_rewards = []
        self.timesteps = []
```

```

def _on_step(self) -> bool:
    if self.locals.get("dones") is not None and any(self.locals["dones"]):
        ep_rew = self.locals["rewards"]
        self.episode_rewards.append(sum(ep_rew))
        self.timesteps.append(self.num_timesteps)
    return True

```

```

[ ]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import gymnasium as gym
from gymnasium import spaces
from sb3_contrib import RecurrentPPO
from sb3_contrib.ppo_recurrent.policies import MlpLstmPolicy
from stable_baselines3.common.vec_env import DummyVecEnv
import torch
import random

# --- Fixed Seed ---
SEED = 78768
np.random.seed(SEED)
random.seed(SEED)
torch.manual_seed(SEED)

# --- Load & Clean Data ---
data_path = "full_data.csv"
data = pd.read_csv(data_path, parse_dates=["Date"])

required_cols = [
    'open', 'high', 'low', 'close',
    'zlma', 'ema_value',
    'DSSb', 'DSSsignal',
    'rsi_ma_base',
    'ZeroLagMACD', 'ZeroLagMACD_signal',
    'basicMACD', 'basicMACD_signal',
    'ZScore',
    'ZLMA_Buy', 'ZLMA_Sell',
    'RSI_Buy', 'RSI_Sell',
    'MCDX_Buy', 'MCDX_Sell',
    'DSS_Buy', 'DSS_Sell',
    'ZeroLag MACD_Buy', 'ZeroLag MACD_Sell',
    'Basic MACD_Buy', 'Basic MACD_Sell'
]

missing = [col for col in required_cols if col not in data.columns]
assert not missing, f"Missing required columns: {missing}"

```

```

data = data.dropna(subset=required_cols).reset_index(drop=True)

SWITCH_COST = 1.0
TRANSACTION_COST = 0.001

ACTIVATION = torch.nn.Tanh
# ACTIVATION = partial(nn.LeakyReLU, negative_slope=0.01)

policy_kwargs = dict(
    activation_fn=ACTIVATION
)

# --- Main Trading Environment ---
class TradingEnvRL(gym.Env):
    metadata = {'render_modes': ['human']}

    def __init__(self, data, initial_balance=10000, hold_cost=0.02,
                 volatility_window=14, exploration_steps=500,
                 switch_cost=1.0, reentry_threshold=0.01,
                 dynamic_threshold=False, loss_penalty=0.75, drawdown_penalty=5.
↪0, large_loss_threshold=-2.0):

        super().__init__()
        self.loss_penalty = loss_penalty
        self.drawdown_penalty = drawdown_penalty
        self.large_loss_threshold = large_loss_threshold
        self.data = data.reset_index(drop=True).copy()
        self.n_steps = len(self.data)
        self.initial_balance = initial_balance
        self.hold_cost = hold_cost
        self.volatility_window = volatility_window
        self.exploration_steps = exploration_steps
        self.switch_cost = switch_cost
        self.reentry_threshold = reentry_threshold
        self.dynamic_threshold = dynamic_threshold
        self.step_rewards = []

        self.feature_cols = [
            'open', 'high', 'low', 'close',
            'basicMACD', 'basicMACD_signal',
            'Basic MACD_Buy', 'Basic MACD_Sell'
        ]

        obs_dim = len(self.feature_cols) + 1
        self.observation_space = spaces.Box(low=-np.inf, high=np.inf,
↪shape=(obs_dim,), dtype=np.float32)

```



```

self.action_space = spaces.Discrete(2) # 0=Long, 1=Short

self._compute_volatility_limit()

def _compute_volatility_limit(self):
    returns = self.data['close'].pct_change()
    self.data['volatility'] = returns.rolling(self.volatility_window).std()
    self.data['adaptive_hold'] = (10 / (self.data['volatility'] * 100)).
    ↪clip(lower=3, upper=20).fillna(10).astype(int)

def reset(self, seed=None, options=None):
    if seed is not None:
        np.random.seed(seed)
        random.seed(seed)
        torch.manual_seed(seed)

    self.current_step = 0
    self.position = 0
    self.entry_price = 0.0
    self.entry_date = None
    self.hold_counter = 0
    self.switch_count = 0
    self.balance = self.initial_balance
    self.cumulative_pnl = 0.0
    self.trade_log = []
    self.action_counts = {0: 0, 1: 0}
    self.reward_tracker = {0: [], 1: []}
    self.consecutive_losses = 0
    self.equity_curve = [self.initial_balance]
    self.step_rewards = []
    return self._get_obs(), {}

def _get_obs(self):
    row = self.data.iloc[self.current_step]
    features = row[self.feature_cols].values.astype(np.float32)
    pos_feature = np.array([self.position], dtype=np.float32)
    return np.concatenate([features, pos_feature])

def _force_close(self):
    row = self.data.iloc[self.current_step]
    current_price = float(row['close'])
    current_date = row['Date'].strftime("%Y-%m-%d")

    if self.position == 0:
        return 0.0

# --- Core Return Logic ---

```

```

        trade_pct = ((current_price / self.entry_price - 1) * 100) if self.
↪position == 1 else ((self.entry_price / current_price - 1) * 100)
        pos_str = 'Long' if self.position == 1 else 'Short'
        gross_return = trade_pct / 100
        transaction_cost = TRANSACTION_COST * current_price

        old_balance = self.balance
        self.balance -= transaction_cost
        self.balance *= (1 + gross_return)
        net_profit = self.balance - old_balance

        reward = net_profit

        # --- Track Trade History ---
        self.cumulative_pnl += trade_pct
        compounded_pnl = (self.trade_log[-1]['CompoundedFactor'] * (1 +
↪gross_return)) if self.trade_log else (1 + gross_return)
        compounded_pnl_pct = (compounded_pnl - 1) * 100

        self.trade_log.append({
            'EntryDate': self.entry_date,
            'ExitDate': current_date,
            'Position': pos_str,
            'EntryPrice': round(self.entry_price, 4),
            'ExitPrice': round(current_price, 4),
            'PnL%': round(trade_pct, 2),
            'CumulativePnL%': round(self.cumulative_pnl, 2),
            'CompoundedFactor': compounded_pnl,
            'CompoundedPnL%': round(compounded_pnl_pct, 2)
        })

        # --- Penalty for Large Loss ---
        if trade_pct < -2.0:
            reward += trade_pct * 2 # Stronger penalty for large loss

        # --- Track & Penalize Consecutive Losses ---
        if trade_pct < 0:
            self.consecutive_losses += 1
            reward -= self.consecutive_losses * 0.75 # Growing penalty
        else:
            self.consecutive_losses = 0

        # --- Optional: Drawdown Penalty ---
        self.equity_curve.append(self.balance)
        max_balance = max(self.equity_curve)
        if max_balance > 0:
            drawdown = (max_balance - self.balance) / max_balance

```

```

        reward -= 5.0 * drawdown # Penalize deeper drawdown

# --- Reset State ---
self.position = 0
self.entry_price = 0.0
self.entry_date = None
self.hold_counter = 0

return reward

def step(self, action):
    if isinstance(action, np.ndarray):
        action = int(action.item())

    row = self.data.iloc[self.current_step]
    current_price = float(row['close'])
    current_date = row['Date'].strftime("%Y-%m-%d")

    if self.current_step < self.exploration_steps:
        action = self.action_space.sample()

    self.action_counts[action] += 1
    reward = 0.0
    target_position = 1 if action == 0 else -1

    # === Holding Same Position ===
    if self.position == target_position:
        self.hold_counter += 1
        price_change = (current_price - self.entry_price) / self.entry_price
        step_return = price_change if self.position == 1 else -price_change

        step_reward = np.sign(step_return) * np.sqrt(abs(step_return)) * 10
        step_reward = np.clip(step_reward, -50, 50)
        step_reward -= self.hold_cost
        reward += step_reward
        self.reward_tracker[action].append(step_reward)
        self.step_rewards.append(reward)

    # === New Position from Flat ===
    elif self.position == 0:
        self.position = target_position
        self.entry_price = current_price
        self.entry_date = current_date
        self.hold_counter = 1
        reward -= TRANSACTION_COST * current_price

    # === Switch Position ===

```

```

else:
    hold_penalty = max(0, 3 - self.hold_counter) * 5.0
    switch_penalty = self.switch_cost + hold_penalty
    reward += self._force_close()
    reward -= switch_penalty

    # === Dynamic Re-entry Threshold ===
    if self.current_step > 0:
        prev_close = self.data.iloc[self.current_step - 1]['close']
        recent_return = abs((current_price - prev_close) / prev_close)

        # Volatility-aware reentry threshold
        recent_volatility = row['volatility']
        reentry_threshold = 0.5 * recent_volatility
    else:
        recent_return = 0
        reentry_threshold = 0.01 # Fallback

    if recent_return > reentry_threshold:
        self.position = target_position
        self.entry_price = current_price
        self.entry_date = current_date
        self.hold_counter = 1
        reward -= TRANSACTION_COST * current_price

self.current_step += 1
terminated = self.current_step >= self.n_steps - 1

if terminated and self.position != 0:
    reward += self._force_close()

obs = self._get_obs() if not terminated else np.zeros(self.
↪observation_space.shape, dtype=np.float32)
return obs, reward, terminated, False, {}

def render(self):
    print(f"Step: {self.current_step}, Position: {self.position}, Balance: ↵
↪{self.balance:.2f}")

def save_trade_log(self, filename="trade_log.csv"):
    df = pd.DataFrame(self.trade_log)
    if "CompoundedFactor" in df.columns:
        df = df.drop(columns=["CompoundedFactor"])
    df.to_csv(filename, index=False)

# --- Training ---
env = TradingEnvRL(data, initial_balance=10000)

```

```

vec_env = DummyVecEnv([lambda: env])
vec_env.seed(SEED)

model = RecurrentPP0(
    policy=MlpLstmPolicy,
    env=vec_env,
    verbose=1,
    n_steps=64,
    batch_size=32,
    learning_rate=0.01,
    gamma=0.99,
    ent_coef=0.01,
    seed=SEED,
    policy_kwargs=policy_kwargs
)

callback = RewardTrackingCallback()

model.learn(total_timesteps=3000, callback=callback)
print(f"\nTraining Complete")

```

```

[519]: # --- Evaluation ---
obs, _ = env.reset()
state = None
done = False
total_reward = 0
portfolio_values = []
final_action = None

while not done:
    action, state = model.predict(obs, state=state, deterministic=True)
    final_action = action
    obs, reward, done, _, _ = env.step(action)
    total_reward += reward

    current_index = min(env.current_step, len(env.data) - 1)
    current_price = env.data.loc[current_index, 'close']
    unrealized = (
        (current_price - env.entry_price) if env.position == 1 else
        (env.entry_price - current_price) if env.position == -1 else
        0.0
    )
    mtm_equity = env.balance + unrealized
    portfolio_values.append(mtm_equity)

# --- Save trade log ---
env.save_trade_log("trade_log_recurrent.csv")

```

```

# --- Final Model Signal ---
signal_str = "BUY" if final_action == 0 else "SELL"
latest_date = env.data['Date'].iloc[env.current_step - 1].strftime("%Y-%m-%d")
print(f"\nLatest model signal at {latest_date}: {signal_str}")

# --- Summary Output ---
print(f"Total Reward: {total_reward:.2f}")
print(f"Final Balance: {env.balance:.2f}")
print("Trade Log (Recurrent PPO):")
print(pd.DataFrame(env.trade_log))
print("Action counts:", env.action_counts)
print("Average reward per action:")
for k, v in env.reward_tracker.items():
    mean_r = np.mean(v) if v else 0
    print(f"Action {k} ({'Long' if k==0 else 'Short'}): {mean_r:.4f}")

```

Latest model signal at 2025-04-09: BUY

Total Reward: 3329.88

Final Balance: 15003.76

Trade Log (Recurrent PPO):

	EntryDate	ExitDate	Position	EntryPrice	ExitPrice	PnL%	\
0	2023-07-27	2023-07-28	Long	128.25	132.21	3.09	
1	2023-07-31	2023-08-03	Short	133.68	128.91	3.70	
2	2023-08-04	2023-08-09	Short	139.57	137.85	1.25	
3	2023-08-10	2023-08-11	Short	138.56	138.41	0.11	
4	2023-08-14	2023-08-15	Short	140.57	137.67	2.11	
..	
169	2025-03-28	2025-03-31	Short	192.72	190.26	1.29	
170	2025-03-31	2025-04-01	Long	190.26	192.17	1.00	
171	2025-04-02	2025-04-08	Short	196.01	170.66	14.85	
172	2025-04-08	2025-04-09	Long	170.66	191.10	11.98	
173	2025-04-09	2025-04-10	Short	191.10	181.22	5.45	

	CumulativePnL%	CompoundedFactor	CompoundedPnL%
0	3.09	1.030877	3.09
1	6.79	1.069022	6.90
2	8.04	1.082361	8.24
3	8.14	1.083534	8.35
4	10.25	1.106358	10.64
..
169	14.36	1.098457	9.85
170	15.36	1.109484	10.95
171	30.22	1.274288	27.43
172	42.19	1.426910	42.69
173	47.64	1.504704	50.47

```
[174 rows x 9 columns]
Action counts: {0: 197, 1: 231}
Average reward per action:
Action 0 (Long): 0.0559
Action 1 (Short): -0.0480
```

```
[520]: # --- Get Latest Signal from Model ---
last_index = len(env.data) - 1
env.current_step = last_index
obs = env._get_obs()

# Add batch dimension and run prediction
obs_input = obs[np.newaxis, :]
action, _ = model.predict(obs_input, deterministic=True)

# Convert action to trading signal
if action == 0:
    current_signal = "BUY"
else:
    current_signal = "SELL"

print(f"\nLatest model signal at {env.data.iloc[last_index]['Date'].date()}:
↳ {current_signal}")
```

Latest model signal at 2025-04-10: BUY

```
[521]: # --- Plot Performance with Trade Markers ---
buy_hold_line = [env.initial_balance * (p / data['close'].iloc[0]) for p in
↳ data['close'].iloc[:len(portfolio_values)]]

date_to_step = {row['Date'].strftime('%Y-%m-%d'): i for i, row in data.iloc[:
↳ len(portfolio_values)].iterrows()}
entry_points = []
exit_points = []

for trade in env.trade_log:
    entry_step = date_to_step.get(trade['EntryDate'])
    exit_step = date_to_step.get(trade['ExitDate'])
    if entry_step is not None and entry_step < len(portfolio_values):
        entry_points.append((entry_step, portfolio_values[entry_step]))
    if exit_step is not None and exit_step < len(portfolio_values):
        exit_points.append((exit_step, portfolio_values[exit_step]))


plt.figure(figsize=(12, 6))
plt.plot(portfolio_values, label="Recurrent PPO Portfolio", linewidth=2)
plt.plot(buy_hold_line, label="Buy & Hold Benchmark", linestyle="--",
↳ color="gray")
```

```

if entry_points:
    entry_steps, entry_vals = zip(*entry_points)
    plt.scatter(entry_steps, entry_vals, color='green', marker='^', s=80,
    ↪label="Trade Entry")
if exit_points:
    exit_steps, exit_vals = zip(*exit_points)
    plt.scatter(exit_steps, exit_vals, color='red', marker='v', s=80,
    ↪label="Trade Exit")

plt.title("Portfolio Value: Recurrent PPO vs. Buy & Hold")
plt.xlabel("Time Step")
plt.ylabel("Portfolio Value")
plt.legend()
plt.grid(True)
plt.tight_layout()
plt.savefig("recurrent_ppo_performance.png")
plt.show()

```



```

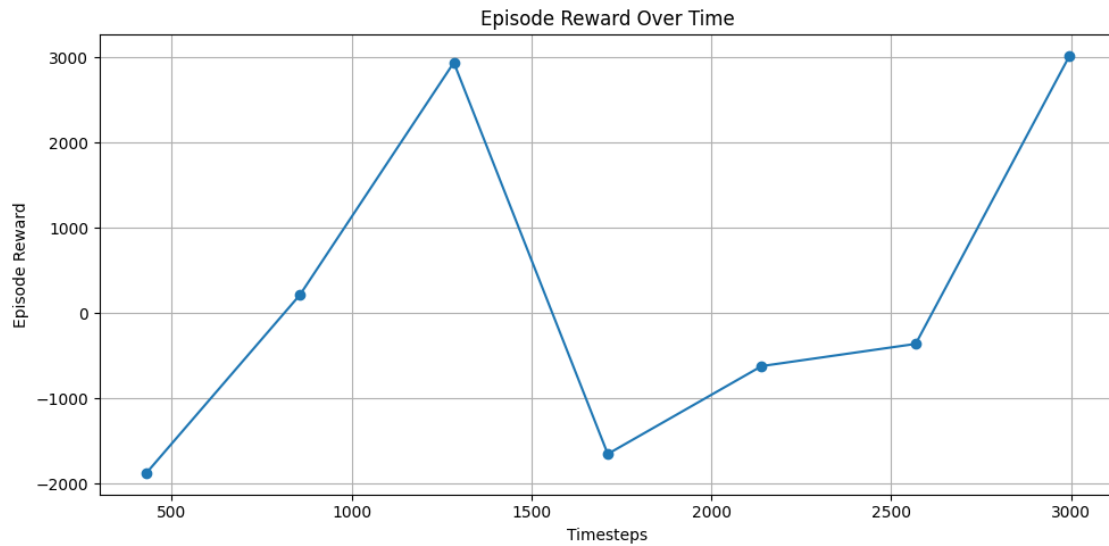
[522]: import matplotlib.pyplot as plt

plt.figure(figsize=(10, 5))
plt.plot(callback.timesteps, callback.episode_rewards, marker='o')
plt.title("Episode Reward Over Time")
plt.xlabel("Timesteps")
plt.ylabel("Episode Reward")
plt.grid(True)
plt.tight_layout()

```




```
plt.show()
```



```
[523]: import matplotlib.pyplot as plt
import numpy as np
import pandas as pd

rewards = np.array(env.step_rewards)

# Histogram
plt.hist(rewards, bins=50, alpha=0.7, edgecolor='black')
plt.title("Reward Distribution per Step")
plt.xlabel("Reward")
plt.ylabel("Frequency")
plt.grid(True)
plt.show()
```



```
[524]: import numpy as np
import matplotlib.pyplot as plt
import pandas as pd

rewards = np.array(env.step_rewards)

# --- Raw volatility metrics ---
reward_std = np.std(rewards)
reward_mean = np.mean(rewards)
reward_variance = np.var(rewards)
reward_range = np.max(rewards) - np.min(rewards)

# --- Rolling volatility ---
window = 20
rolling_std = pd.Series(rewards).rolling(window=window).std()


# --- Plot ---
plt.figure(figsize=(10, 4))
plt.plot(rolling_std, label=f"Rolling Std (window={window})")
```

```

plt.axhline(reward_std, color='red', linestyle='--', label='Global Std')
plt.title("Reward Volatility (Rolling Std Dev of Step Rewards)")
plt.xlabel("Step")
plt.ylabel("Rolling Std Dev")
plt.legend()
plt.grid(True)
plt.tight_layout()
plt.show()

# --- Output stats ---
print(f"Mean Step Reward: {reward_mean:.4f}")
print(f"Std Dev of Step Rewards: {reward_std:.4f}")
print(f"Variance: {reward_variance:.4f}")
print(f"Range: {reward_range:.4f}")
print(f"Max Rolling Std ({window}): {rolling_std.max():.4f}")
print(f"Min Rolling Std ({window}): {rolling_std.min():.4f}")

```



```

Mean Step Reward: -0.0036
Std Dev of Step Rewards: 1.3703
Variance: 1.8777
Range: 6.2073
Max Rolling Std (20): 2.1944
Min Rolling Std (20): 0.8812

```

```

[525]: import pandas as pd
import matplotlib.pyplot as plt

def compute_switch_rate(trade_log: pd.DataFrame, plot: bool = True):
    if 'Position' not in trade_log.columns or trade_log.empty:
        raise ValueError("Trade log must contain a 'Position' column and be
        non-empty.")

```

```

positions = trade_log['Position'].tolist()
if len(positions) < 2:
    print("Insufficient trades to compute switching behavior.")
    return 0.0

switches = sum(1 for i in range(1, len(positions)) if positions[i] !=
↳positions[i - 1])
total_transitions = len(positions) - 1
switch_rate = switches / total_transitions

print(f"Total Trades: {len(positions)}")
print(f"Total Transitions: {total_transitions}")
print(f"Switches: {switches}")
print(f"Switch Rate: {switch_rate:.4f}")

if plot:
    plt.bar(['Same', 'Switch'], [total_transitions - switches, switches],
↳color=['gray', 'red'])
    plt.title("Trade Direction Transitions")
    plt.ylabel("Count")
    plt.grid(True, axis='y', linestyle='--', alpha=0.6)
    plt.show()

return switch_rate

df_trades = pd.DataFrame(env.trade_log)
switch_rate = compute_switch_rate(df_trades)

```

```

Total Trades: 174
Total Transitions: 173
Switches: 129
Switch Rate: 0.7457

```



```
[526]: import matplotlib.pyplot as plt

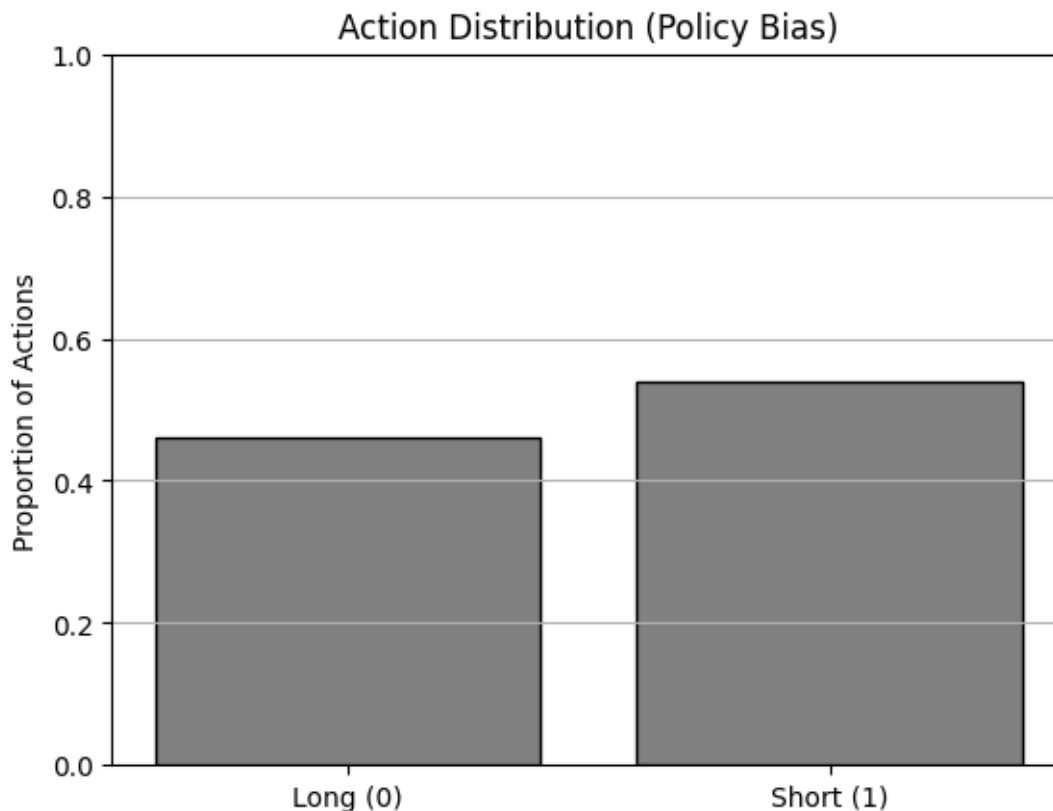
total_actions = sum(env.action_counts.values())

actions = list(env.action_counts.keys())
counts = [env.action_counts[a] for a in actions]
labels = ['Long (0)', 'Short (1)']

proportions = [count / total_actions for count in counts]

plt.bar(labels, proportions, color='gray', edgecolor='black')
plt.title("Action Distribution (Policy Bias)")
plt.ylabel("Proportion of Actions")
plt.ylim(0, 1)
plt.grid(True, axis='y')
plt.show()

for label, count, prop in zip(labels, counts, proportions):
    print(f"{label}: Count = {count}, Proportion = {prop:.2%}")
```



Long (0): Count = 197, Proportion = 46.03%

Short (1): Count = 231, Proportion = 53.97%

— Training with seed 58630 — Seed 58630: Reward = 1182.75, Final Balance = 12750.29 Average reward for action Long: 0.3288 Average reward for action Short: -0.4190 Best model updated and saved.

— Training with seed 91454 — Seed 91454: Reward = -3917.27, Final Balance = 8096.02 Average reward for action Long: 0.0848 Average reward for action Short: -0.3334

— Training with seed 81001 — Seed 81001: Reward = -4347.41, Final Balance = 7382.52 Average reward for action Long: -0.5212 Average reward for action Short: -0.3672

— Training with seed 40163 — Seed 40163: Reward = 1815.62, Final Balance = 13615.02 Average reward for action Long: 0.2187 Average reward for action Short: -0.2444 Best model updated and saved.

— Training with seed 78768 — Seed 78768: Reward = 5456.93, Final Balance = 17079.26 Average reward for action Long: 0.1406 Average reward for action Short: -0.0176 Best model updated and saved.

— Training with seed 64909 — Seed 64909: Reward = 2967.96, Final Balance = 14889.38 Average reward for action Long: 0.1946 Average reward for action Short: 0.0206

— Training with seed 120070 — Seed 120070: Reward = -4524.22, Final Balance = 7404.34 Average

reward for action Long: -0.1343 Average reward for action Short: -0.2350

— Training with seed 97257 — Seed 97257: Reward = -2579.63, Final Balance = 9205.13 Average reward for action Long: 0.0203 Average reward for action Short: -0.3204

— Training with seed 123783 — Seed 123783: Reward = 1518.37, Final Balance = 13395.80 Average reward for action Long: 0.0246 Average reward for action Short: -0.0679

— Training with seed 120497 — Seed 120497: Reward = -2572.49, Final Balance = 9149.12 Average reward for action Long: 0.1572 Average reward for action Short: -0.1886

— Training with seed 107358 — Seed 107358: Reward = -2230.98, Final Balance = 9478.17 Average reward for action Long: -0.1151 Average reward for action Short: -0.1609

— Seed Sweep Complete — Best Seed: 78768 Best Total Reward: 5456.93 Best Final Balance: 17079.26

```
[527]: model.save("AMZN_best_model")
```

1.10 AAPL

```
[528]: from datetime import datetime, timedelta

# Add +1 day to end date
default_end_date = (datetime.now() + timedelta(days=1)).strftime("%Y-%m-%d")
default_start_date = (datetime.now() - timedelta(days=665)).strftime("%Y-%m-%d")
```

```
[529]: # Parameters
ticker = "AAPL"
start_date = default_start_date
end_date = default_end_date
data_filename = "full_data.csv"
signals_filename = "signals_data.csv"

save_historical_data(
    ticker=ticker,
    start_date=start_date,
    end_date=end_date,
    data_filename=data_filename,
    signals_filename=signals_filename
)
```

```
[*****100%*****] 1 of 1 completed
```

```
[529]: 'Saved full_data.csv and signals_data.csv successfully.'
```

```
[530]: import pandas as pd

# Preview full data
df_full = pd.read_csv("full_data.csv", parse_dates=["Date"])
print(df_full.head())
```

```
# Preview signal data
df_signals = pd.read_csv("signals_data.csv", parse_dates=["Date"])
print(df_signals.head())
```

	Date	close	high	low	open	volume \
0	2023-06-16	183.326981	185.379156	182.682586	185.121386	101235600
1	2023-06-20	183.416199	184.496820	182.821377	182.821377	49799100
2	2023-06-21	182.375275	183.812781	181.017066	183.307164	49515700
3	2023-06-22	185.389069	185.438641	182.087753	182.157158	51245300
4	2023-06-23	185.071808	185.944232	183.416196	183.951553	53079300

	EMA_50	EMA_100	EMA_200	EMA_500	...	RSI_Sell	MCDX_Buy \
0	183.326981	183.326981	183.326981	183.326981	...	1	0
1	183.330479	183.328747	183.327868	183.327337	...	1	0
2	183.293020	183.309867	183.318390	183.323536	...	1	0
3	183.375218	183.351039	183.338994	183.331782	...	1	0
4	183.441751	183.385114	183.356235	183.338728	...	0	0

	MCDX_Sell	DSS_Buy	DSS_Sell	ZeroLag	MACD_Buy	ZeroLag	MACD_Sell \
0	1	0	1		0		1
1	1	0	1		1		0
2	1	0	1		0		1
3	1	0	1		1		0
4	1	0	1		1		0

	Basic	MACD_Buy	Basic	MACD_Sell	OverallTrade
0		0		1	Sell
1		1		0	Sell
2		0		1	Sell
3		1		0	Sell
4		1		0	Buy

[5 rows x 42 columns]

	Date	Signal	Z-Score
0	2023-06-20	Basic MACD Buy	NaN
1	2023-06-20	ZLMA Buy	NaN
2	2023-06-20	ZeroLag MACD Buy	NaN
3	2023-06-21	ZLMA Sell	NaN
4	2023-06-21	ZeroLag MACD Sell	NaN

```
[531]: from stable_baselines3.common.callbacks import BaseCallback
```

```
class RewardTrackingCallback(BaseCallback):
    def __init__(self, verbose=0):
        super().__init__(verbose)
        self.episode_rewards = []
```



```

        self.timesteps = []

    def _on_step(self) -> bool:
        if self.locals.get("dones") is not None and any(self.locals["dones"]):
            ep_rew = self.locals["rewards"]
            self.episode_rewards.append(sum(ep_rew))
            self.timesteps.append(self.num_timesteps)
        return True

```

```

[532]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import gymnasium as gym
from gymnasium import spaces
from sb3_contrib import RecurrentPPO
from sb3_contrib.ppo_recurrent.policies import MlpLstmPolicy
from stable_baselines3.common.vec_env import DummyVecEnv
import torch
import random

# --- Fixed Seed ---
SEED = 128030
np.random.seed(SEED)
random.seed(SEED)
torch.manual_seed(SEED)

# --- Load & Clean Data ---
data_path = "full_data.csv"
data = pd.read_csv(data_path, parse_dates=["Date"])

required_cols = [
    'open', 'high', 'low', 'close',
    'zlma', 'ema_value',
    'DSSb', 'DSSsignal',
    'rsi_ma_base',
    'ZeroLagMACD', 'ZeroLagMACD_signal',
    'basicMACD', 'basicMACD_signal',
    'ZScore',
    'ZLMA_Buy', 'ZLMA_Sell',
    'RSI_Buy', 'RSI_Sell',
    'MCDX_Buy', 'MCDX_Sell',
    'DSS_Buy', 'DSS_Sell',
    'ZeroLag MACD_Buy', 'ZeroLag MACD_Sell',
    'Basic MACD_Buy', 'Basic MACD_Sell'
]

missing = [col for col in required_cols if col not in data.columns]

```

```

assert not missing, f"Missing required columns: {missing}"

data = data.dropna(subset=required_cols).reset_index(drop=True)

SWITCH_COST = 1.0
TRANSACTION_COST = 0.001

ACTIVATION = torch.nn.Tanh
# ACTIVATION = partial(nn.LeakyReLU, negative_slope=0.01)

policy_kwargs = dict(
    activation_fn=ACTIVATION
)

# --- Main Trading Environment ---
class TradingEnvRL(gym.Env):
    metadata = {'render_modes': ['human']}

    def __init__(self, data, initial_balance=10000, hold_cost=0.02,
                  volatility_window=14, exploration_steps=500,
                  switch_cost=1.0, reentry_threshold=0.01,
                  dynamic_threshold=False, loss_penalty=0.75, drawdown_penalty=5.
↪0, large_loss_threshold=-2.0):

        super().__init__()
        self.loss_penalty = loss_penalty
        self.drawdown_penalty = drawdown_penalty
        self.large_loss_threshold = large_loss_threshold
        self.data = data.reset_index(drop=True).copy()
        self.n_steps = len(self.data)
        self.initial_balance = initial_balance
        self.hold_cost = hold_cost
        self.volatility_window = volatility_window
        self.exploration_steps = exploration_steps
        self.switch_cost = switch_cost
        self.reentry_threshold = reentry_threshold
        self.dynamic_threshold = dynamic_threshold
        self.step_rewards = []

        self.feature_cols = [
            'open', 'high', 'low', 'close',
            'basicMACD', 'basicMACD_signal',
            'Basic MACD_Buy', 'Basic MACD_Sell'
        ]

        obs_dim = len(self.feature_cols) + 1

```

```

        self.observation_space = spaces.Box(low=-np.inf, high=np.inf,
↪shape=(obs_dim,), dtype=np.float32)
        self.action_space = spaces.Discrete(2) # 0=Long, 1=Short

        self._compute_volatility_limit()

    def _compute_volatility_limit(self):
        returns = self.data['close'].pct_change()
        self.data['volatility'] = returns.rolling(self.volatility_window).std()
        self.data['adaptive_hold'] = (10 / (self.data['volatility'] * 100)).
↪clip(lower=3, upper=20).fillna(10).astype(int)

    def reset(self, seed=None, options=None):
        if seed is not None:
            np.random.seed(seed)
            random.seed(seed)
            torch.manual_seed(seed)

        self.current_step = 0
        self.position = 0
        self.entry_price = 0.0
        self.entry_date = None
        self.hold_counter = 0
        self.switch_count = 0
        self.balance = self.initial_balance
        self.cumulative_pnl = 0.0
        self.trade_log = []
        self.action_counts = {0: 0, 1: 0}
        self.reward_tracker = {0: [], 1: []}
        self.consecutive_losses = 0
        self.equity_curve = [self.initial_balance]
        self.step_rewards = []
        return self._get_obs(), {}

    def _get_obs(self):
        row = self.data.iloc[self.current_step]
        features = row[self.feature_cols].values.astype(np.float32)
        pos_feature = np.array([self.position], dtype=np.float32)
        return np.concatenate([features, pos_feature])

    def _force_close(self):
        row = self.data.iloc[self.current_step]
        current_price = float(row['close'])
        current_date = row['Date'].strftime("%Y-%m-%d")

        if self.position == 0:
            return 0.0

```

```

# --- Core Return Logic ---
trade_pct = ((current_price / self.entry_price - 1) * 100) if self.
↪position == 1 else ((self.entry_price / current_price - 1) * 100)
pos_str = 'Long' if self.position == 1 else 'Short'
gross_return = trade_pct / 100
transaction_cost = TRANSACTION_COST * current_price

old_balance = self.balance
self.balance -= transaction_cost
self.balance *= (1 + gross_return)
net_profit = self.balance - old_balance

reward = net_profit

# --- Track Trade History ---
self.cumulative_pnl += trade_pct
compounded_pnl = (self.trade_log[-1]['CompoundedFactor'] * (1 + ↪
↪gross_return)) if self.trade_log else (1 + gross_return)
compounded_pnl_pct = (compounded_pnl - 1) * 100

self.trade_log.append({
    'EntryDate': self.entry_date,
    'ExitDate': current_date,
    'Position': pos_str,
    'EntryPrice': round(self.entry_price, 4),
    'ExitPrice': round(current_price, 4),
    'PnL%': round(trade_pct, 2),
    'CumulativePnL%': round(self.cumulative_pnl, 2),
    'CompoundedFactor': compounded_pnl,
    'CompoundedPnL%': round(compounded_pnl_pct, 2)
})

# --- Penalty for Large Loss ---
if trade_pct < -2.0:
    reward += trade_pct * 2 # Stronger penalty for large loss

# --- Track & Penalize Consecutive Losses ---
if trade_pct < 0:
    self.consecutive_losses += 1
    reward -= self.consecutive_losses * 0.75 # Growing penalty
else:
    self.consecutive_losses = 0

# --- Optional: Drawdown Penalty ---
self.equity_curve.append(self.balance)
max_balance = max(self.equity_curve)

```

```

    if max_balance > 0:
        drawdown = (max_balance - self.balance) / max_balance
        reward -= 5.0 * drawdown # Penalize deeper drawdown

    # --- Reset State ---
    self.position = 0
    self.entry_price = 0.0
    self.entry_date = None
    self.hold_counter = 0

    return reward

def step(self, action):
    if isinstance(action, np.ndarray):
        action = int(action.item())

    row = self.data.iloc[self.current_step]
    current_price = float(row['close'])
    current_date = row['Date'].strftime("%Y-%m-%d")

    if self.current_step < self.exploration_steps:
        action = self.action_space.sample()

    self.action_counts[action] += 1
    reward = 0.0
    target_position = 1 if action == 0 else -1

    # === Holding Same Position ===
    if self.position == target_position:
        self.hold_counter += 1
        price_change = (current_price - self.entry_price) / self.entry_price
        step_return = price_change if self.position == 1 else -price_change

        step_reward = np.sign(step_return) * np.sqrt(abs(step_return)) * 10
        step_reward = np.clip(step_reward, -50, 50)
        step_reward -= self.hold_cost
        reward += step_reward
        self.reward_tracker[action].append(step_reward)
        self.step_rewards.append(reward)

    # === New Position from Flat ===
    elif self.position == 0:
        self.position = target_position
        self.entry_price = current_price
        self.entry_date = current_date
        self.hold_counter = 1
        reward -= TRANSACTION_COST * current_price

```

```

# === Switch Position ===
else:
    hold_penalty = max(0, 3 - self.hold_counter) * 5.0
    switch_penalty = self.switch_cost + hold_penalty
    reward += self._force_close()
    reward -= switch_penalty

# === Dynamic Re-entry Threshold ===
if self.current_step > 0:
    prev_close = self.data.iloc[self.current_step - 1]['close']
    recent_return = abs((current_price - prev_close) / prev_close)

    # Volatility-aware reentry threshold
    recent_volatility = row['volatility']
    reentry_threshold = 0.5 * recent_volatility
else:
    recent_return = 0
    reentry_threshold = 0.01 # Fallback

if recent_return > reentry_threshold:
    self.position = target_position
    self.entry_price = current_price
    self.entry_date = current_date
    self.hold_counter = 1
    reward -= TRANSACTION_COST * current_price

self.current_step += 1
terminated = self.current_step >= self.n_steps - 1

if terminated and self.position != 0:
    reward += self._force_close()

obs = self._get_obs() if not terminated else np.zeros(self.
↪observation_space.shape, dtype=np.float32)
return obs, reward, terminated, False, {}

def render(self):
    print(f"Step: {self.current_step}, Position: {self.position}, Balance:␣
↪{self.balance:.2f}")

def save_trade_log(self, filename="trade_log.csv"):
    df = pd.DataFrame(self.trade_log)
    if "CompoundedFactor" in df.columns:
        df = df.drop(columns=["CompoundedFactor"])
    df.to_csv(filename, index=False)

```

```

# --- Training ---
env = TradingEnvRL(data, initial_balance=10000)
vec_env = DummyVecEnv([lambda: env])
vec_env.seed(SEED)

model = RecurrentPPO(
    policy=MlpLstmPolicy,
    env=vec_env,
    verbose=1,
    n_steps=64,
    batch_size=32,
    learning_rate=0.01,
    gamma=0.99,
    ent_coef=0.01,
    seed=SEED,
    policy_kwargs=policy_kwargs
)

callback = RewardTrackingCallback()

model.learn(total_timesteps=3000, callback=callback)
print(f"\nTraining Complete")

```

Using cuda device

```

-----
| time/          |      |
|   fps          | 212  |
|  iterations    | 1    |
| time_elapsed   | 0    |
| total_timesteps | 64   |
-----

```

```

-----
| time/          |      |
|   fps          | 99   |
|  iterations    | 2    |
| time_elapsed   | 1    |
| total_timesteps | 128  |
| train/         |      |
| approx_kl      | 0.022984039 |
| clip_fraction  | 0.483 |
| clip_range     | 0.2   |
| entropy_loss   | -0.667 |
| explained_variance | 0.000254 |
| learning_rate  | 0.01  |
| loss           | 8.64e+03 |
| n_updates      | 10    |
| policy_gradient_loss | -0.0229 |
| value_loss     | 7.3e+04 |
-----

```

time/		
fps	106	
iterations	3	
time_elapsed	1	
total_timesteps	192	
train/		
approx_kl	0.04052279	
clip_fraction	0.314	
clip_range	0.2	
entropy_loss	-0.689	
explained_variance	2.38e-07	
learning_rate	0.01	
loss	4.04e+04	
n_updates	20	
policy_gradient_loss	0.00177	
value_loss	7.92e+04	

time/		
fps	110	
iterations	4	
time_elapsed	2	
total_timesteps	256	
train/		
approx_kl	0.008715083	
clip_fraction	0.145	
clip_range	0.2	
entropy_loss	-0.645	
explained_variance	2.38e-07	
learning_rate	0.01	
loss	3.53e+04	
n_updates	30	
policy_gradient_loss	-0.0036	
value_loss	1.04e+05	

time/		
fps	113	
iterations	5	
time_elapsed	2	
total_timesteps	320	
train/		
approx_kl	0.007861175	
clip_fraction	0.119	
clip_range	0.2	
entropy_loss	-0.614	

explained_variance	0	
learning_rate	0.01	
loss	7.03e+04	
n_updates	40	
policy_gradient_loss	-0.00661	
value_loss	1.31e+05	

time/		
fps	115	
iterations	6	
time_elapsed	3	
total_timesteps	384	
train/		
approx_kl	0.013576364	
clip_fraction	0.131	
clip_range	0.2	
entropy_loss	-0.634	
explained_variance	0	
learning_rate	0.01	
loss	4.3e+04	
n_updates	50	
policy_gradient_loss	-0.00765	
value_loss	7.43e+04	

time/		
fps	117	
iterations	7	
time_elapsed	3	
total_timesteps	448	
train/		
approx_kl	0.0006895084	
clip_fraction	0	
clip_range	0.2	
entropy_loss	-0.658	
explained_variance	0	
learning_rate	0.01	
loss	1.92e+05	
n_updates	60	
policy_gradient_loss	-0.000158	
value_loss	4.33e+05	

time/		
fps	106	
iterations	8	
time_elapsed	4	

	total_timesteps		512	
	train/			
	approx_kl		0.02264718	
	clip_fraction		0.075	
	clip_range		0.2	
	entropy_loss		-0.683	
	explained_variance		-1.19e-07	
	learning_rate		0.01	
	loss		5.35e+05	
	n_updates		70	
	policy_gradient_loss		-0.0158	
	value_loss		7.32e+05	

	time/			
	fps		108	
	iterations		9	
	time_elapsed		5	
	total_timesteps		576	
	train/			
	approx_kl		0.01759915	
	clip_fraction		0	
	clip_range		0.2	
	entropy_loss		-0.684	
	explained_variance		1.79e-07	
	learning_rate		0.01	
	loss		3.15e+04	
	n_updates		80	
	policy_gradient_loss		-0.00253	
	value_loss		6.89e+04	

	time/			
	fps		110	
	iterations		10	
	time_elapsed		5	
	total_timesteps		640	
	train/			
	approx_kl		0.12803873	
	clip_fraction		0.186	
	clip_range		0.2	
	entropy_loss		-0.637	
	explained_variance		0	
	learning_rate		0.01	
	loss		5.65e+03	
	n_updates		90	
	policy_gradient_loss		0.0148	
	value_loss		3e+04	

time/		
fps	111	
iterations	11	
time_elapsed	6	
total_timesteps	704	
train/		
approx_kl	0.074654855	
clip_fraction	0.88	
clip_range	0.2	
entropy_loss	-0.656	
explained_variance	0	
learning_rate	0.01	
loss	2.01e+05	
n_updates	100	
policy_gradient_loss	0.163	
value_loss	3.34e+05	

time/		
fps	112	
iterations	12	
time_elapsed	6	
total_timesteps	768	
train/		
approx_kl	0.00012970204	
clip_fraction	0.22	
clip_range	0.2	
entropy_loss	-0.655	
explained_variance	0	
learning_rate	0.01	
loss	9.7e+03	
n_updates	110	
policy_gradient_loss	-0.00161	
value_loss	1.07e+05	

time/		
fps	113	
iterations	13	
time_elapsed	7	
total_timesteps	832	
train/		
approx_kl	0.02278279	
clip_fraction	0.15	
clip_range	0.2	
entropy_loss	-0.672	

explained_variance	0	
learning_rate	0.01	
loss	8.04e+03	
n_updates	120	
policy_gradient_loss	-0.0132	
value_loss	7.03e+04	

time/		
fps	114	
iterations	14	
time_elapsed	7	
total_timesteps	896	
train/		
approx_kl	0.021921577	
clip_fraction	0.508	
clip_range	0.2	
entropy_loss	-0.678	
explained_variance	0	
learning_rate	0.01	
loss	2.92e+04	
n_updates	130	
policy_gradient_loss	-0.0168	
value_loss	1.42e+05	

time/		
fps	109	
iterations	15	
time_elapsed	8	
total_timesteps	960	
train/		
approx_kl	0.2543922	
clip_fraction	0.873	
clip_range	0.2	
entropy_loss	-0.387	
explained_variance	0	
learning_rate	0.01	
loss	1.46e+06	
n_updates	140	
policy_gradient_loss	0.0959	
value_loss	1.52e+06	

time/		
fps	110	
iterations	16	
time_elapsed	9	

	total_timesteps	1024	
	train/		
	approx_kl	1.8340633	
	clip_fraction	0.211	
	clip_range	0.2	
	entropy_loss	-0.142	
	explained_variance	0	
	learning_rate	0.01	
	loss	3.29e+04	
	n_updates	150	
	policy_gradient_loss	0.0396	
	value_loss	5.26e+04	

	time/		
	fps	111	
	iterations	17	
	time_elapsed	9	
	total_timesteps	1088	
	train/		
	approx_kl	0.0	
	clip_fraction	0	
	clip_range	0.2	
	entropy_loss	-9.92e-05	
	explained_variance	5.96e-08	
	learning_rate	0.01	
	loss	1.22e+05	
	n_updates	160	
	policy_gradient_loss	3.91e-09	
	value_loss	1.86e+05	

	time/		
	fps	111	
	iterations	18	
	time_elapsed	10	
	total_timesteps	1152	
	train/		
	approx_kl	0.0	
	clip_fraction	0	
	clip_range	0.2	
	entropy_loss	-5.59e-05	
	explained_variance	1.19e-07	
	learning_rate	0.01	
	loss	5.51e+05	
	n_updates	170	
	policy_gradient_loss	-4.28e-09	
	value_loss	5.86e+05	

time/		
fps	112	
iterations	19	
time_elapsed	10	
total_timesteps	1216	
train/		
approx_kl	0.0	
clip_fraction	0	
clip_range	0.2	
entropy_loss	-5.26e-05	
explained_variance	0	
learning_rate	0.01	
loss	1.25e+05	
n_updates	180	
policy_gradient_loss	1.49e-09	
value_loss	1.68e+05	

time/		
fps	112	
iterations	20	
time_elapsed	11	
total_timesteps	1280	
train/		
approx_kl	0.0	
clip_fraction	0	
clip_range	0.2	
entropy_loss	-5.22e-05	
explained_variance	1.79e-07	
learning_rate	0.01	
loss	4.99e+04	
n_updates	190	
policy_gradient_loss	-3.04e-08	
value_loss	9.29e+04	

time/		
fps	112	
iterations	21	
time_elapsed	11	
total_timesteps	1344	
train/		
approx_kl	0.0	
clip_fraction	0	
clip_range	0.2	
entropy_loss	-5.22e-05	

explained_variance	0	
learning_rate	0.01	
loss	1.13e+05	
n_updates	200	
policy_gradient_loss	-2.24e-09	
value_loss	5.92e+05	

time/		
fps	108	
iterations	22	
time_elapsed	12	
total_timesteps	1408	
train/		
approx_kl	0.0	
clip_fraction	0	
clip_range	0.2	
entropy_loss	-5.22e-05	
explained_variance	2.38e-07	
learning_rate	0.01	
loss	2.84e+04	
n_updates	210	
policy_gradient_loss	-1.23e-08	
value_loss	2e+05	

time/		
fps	109	
iterations	23	
time_elapsed	13	
total_timesteps	1472	
train/		
approx_kl	0.0	
clip_fraction	0	
clip_range	0.2	
entropy_loss	-5.22e-05	
explained_variance	0	
learning_rate	0.01	
loss	1.25e+04	
n_updates	220	
policy_gradient_loss	-1.08e-08	
value_loss	3.56e+04	

time/		
fps	110	
iterations	24	
time_elapsed	13	

	total_timesteps		1536	
	train/			
	approx_kl		0.0	
	clip_fraction		0	
	clip_range		0.2	
	entropy_loss		-5.22e-05	
	explained_variance		5.96e-08	
	learning_rate		0.01	
	loss		3.19e+04	
	n_updates		230	
	policy_gradient_loss		2.42e-08	
	value_loss		1.64e+05	

	time/			
	fps		110	
	iterations		25	
	time_elapsed		14	
	total_timesteps		1600	
	train/			
	approx_kl		0.0	
	clip_fraction		0	
	clip_range		0.2	
	entropy_loss		-5.22e-05	
	explained_variance		0	
	learning_rate		0.01	
	loss		2.03e+05	
	n_updates		240	
	policy_gradient_loss		2.24e-09	
	value_loss		2.77e+05	

	time/			
	fps		111	
	iterations		26	
	time_elapsed		14	
	total_timesteps		1664	
	train/			
	approx_kl		0.0	
	clip_fraction		0	
	clip_range		0.2	
	entropy_loss		-5.22e-05	
	explained_variance		-1.19e-07	
	learning_rate		0.01	
	loss		5.03e+04	
	n_updates		250	
	policy_gradient_loss		4.88e-08	
	value_loss		2.15e+05	

time/		
fps	111	
iterations	27	
time_elapsed	15	
total_timesteps	1728	
train/		
approx_kl	0.0	
clip_fraction	0	
clip_range	0.2	
entropy_loss	-5.22e-05	
explained_variance	-1.19e-07	
learning_rate	0.01	
loss	9.72e+04	
n_updates	260	
policy_gradient_loss	1.9e-08	
value_loss	1.19e+05	

time/		
fps	109	
iterations	28	
time_elapsed	16	
total_timesteps	1792	
train/		
approx_kl	0.0	
clip_fraction	0	
clip_range	0.2	
entropy_loss	-5.22e-05	
explained_variance	0	
learning_rate	0.01	
loss	2.39e+05	
n_updates	270	
policy_gradient_loss	4.02e-08	
value_loss	8.13e+05	

time/		
fps	109	
iterations	29	
time_elapsed	16	
total_timesteps	1856	
train/		
approx_kl	0.0	
clip_fraction	0	
clip_range	0.2	
entropy_loss	-5.22e-05	

explained_variance	0	
learning_rate	0.01	
loss	2.07e+04	
n_updates	280	
policy_gradient_loss	1.86e-10	
value_loss	4.42e+04	

time/		
fps	110	
iterations	30	
time_elapsed	17	
total_timesteps	1920	
train/		
approx_kl	0.0	
clip_fraction	0	
clip_range	0.2	
entropy_loss	-5.22e-05	
explained_variance	0	
learning_rate	0.01	
loss	6.87e+04	
n_updates	290	
policy_gradient_loss	5.81e-08	
value_loss	1.48e+05	

time/		
fps	110	
iterations	31	
time_elapsed	17	
total_timesteps	1984	
train/		
approx_kl	0.0	
clip_fraction	0	
clip_range	0.2	
entropy_loss	-5.22e-05	
explained_variance	0	
learning_rate	0.01	
loss	1.32e+05	
n_updates	300	
policy_gradient_loss	1.04e-08	
value_loss	1.52e+05	

time/		
fps	110	
iterations	32	
time_elapsed	18	

	total_timesteps		2048	
	train/			
	approx_kl		0.0	
	clip_fraction		0	
	clip_range		0.2	
	entropy_loss		-5.22e-05	
	explained_variance		0	
	learning_rate		0.01	
	loss		3.04e+05	
	n_updates		310	
	policy_gradient_loss		1.53e-08	
	value_loss		3.18e+05	

	time/			
	fps		111	
	iterations		33	
	time_elapsed		18	
	total_timesteps		2112	
	train/			
	approx_kl		0.0	
	clip_fraction		0	
	clip_range		0.2	
	entropy_loss		-5.22e-05	
	explained_variance		0	
	learning_rate		0.01	
	loss		4.96e+03	
	n_updates		320	
	policy_gradient_loss		2.31e-08	
	value_loss		1.73e+04	

	time/			
	fps		111	
	iterations		34	
	time_elapsed		19	
	total_timesteps		2176	
	train/			
	approx_kl		0.0	
	clip_fraction		0	
	clip_range		0.2	
	entropy_loss		-5.22e-05	
	explained_variance		0	
	learning_rate		0.01	
	loss		1.46e+04	
	n_updates		330	
	policy_gradient_loss		3.35e-09	
	value_loss		7.84e+04	

time/		
fps	109	
iterations	35	
time_elapsed	20	
total_timesteps	2240	
train/		
approx_kl	0.0	
clip_fraction	0	
clip_range	0.2	
entropy_loss	-5.22e-05	
explained_variance	1.19e-07	
learning_rate	0.01	
loss	3.95e+04	
n_updates	340	
policy_gradient_loss	3.73e-09	
value_loss	1.47e+05	

time/		
fps	110	
iterations	36	
time_elapsed	20	
total_timesteps	2304	
train/		
approx_kl	0.0	
clip_fraction	0	
clip_range	0.2	
entropy_loss	-5.22e-05	
explained_variance	0	
learning_rate	0.01	
loss	1.33e+04	
n_updates	350	
policy_gradient_loss	7.45e-10	
value_loss	2.51e+04	

time/		
fps	110	
iterations	37	
time_elapsed	21	
total_timesteps	2368	
train/		
approx_kl	0.0	
clip_fraction	0	
clip_range	0.2	
entropy_loss	-5.22e-05	

	explained_variance	0	
	learning_rate	0.01	
	loss	1.63e+04	
	n_updates	360	
	policy_gradient_loss	-5.96e-09	
	value_loss	2.42e+04	

	time/		
	fps	110	
	iterations	38	
	time_elapsed	21	
	total_timesteps	2432	
	train/		
	approx_kl	0.0	
	clip_fraction	0	
	clip_range	0.2	
	entropy_loss	-5.23e-05	
	explained_variance	5.96e-08	
	learning_rate	0.01	
	loss	4.86e+04	
	n_updates	370	
	policy_gradient_loss	-6.33e-09	
	value_loss	2.36e+05	

	time/		
	fps	111	
	iterations	39	
	time_elapsed	22	
	total_timesteps	2496	
	train/		
	approx_kl	0.0	
	clip_fraction	0	
	clip_range	0.2	
	entropy_loss	-5.23e-05	
	explained_variance	1.19e-07	
	learning_rate	0.01	
	loss	3.38e+04	
	n_updates	380	
	policy_gradient_loss	1.86e-08	
	value_loss	8.66e+04	

	time/		
	fps	111	
	iterations	40	
	time_elapsed	22	

	total_timesteps	2560	
	train/		
	approx_kl	0.0	
	clip_fraction	0	
	clip_range	0.2	
	entropy_loss	-5.23e-05	
	explained_variance	1.19e-07	
	learning_rate	0.01	
	loss	1.26e+04	
	n_updates	390	
	policy_gradient_loss	4.45e-08	
	value_loss	2.79e+04	

	time/		
	fps	111	
	iterations	41	
	time_elapsed	23	
	total_timesteps	2624	
	train/		
	approx_kl	0.0	
	clip_fraction	0	
	clip_range	0.2	
	entropy_loss	-5.23e-05	
	explained_variance	-1.19e-07	
	learning_rate	0.01	
	loss	7.9e+04	
	n_updates	400	
	policy_gradient_loss	6.52e-09	
	value_loss	2.02e+05	

	time/		
	fps	110	
	iterations	42	
	time_elapsed	24	
	total_timesteps	2688	
	train/		
	approx_kl	0.0	
	clip_fraction	0	
	clip_range	0.2	
	entropy_loss	-5.23e-05	
	explained_variance	0	
	learning_rate	0.01	
	loss	8.22e+04	
	n_updates	410	
	policy_gradient_loss	9.69e-09	
	value_loss	1.32e+05	

time/		
fps	110	
iterations	43	
time_elapsed	24	
total_timesteps	2752	
train/		
approx_kl	0.0	
clip_fraction	0	
clip_range	0.2	
entropy_loss	-5.23e-05	
explained_variance	0	
learning_rate	0.01	
loss	5.58e+03	
n_updates	420	
policy_gradient_loss	1.79e-08	
value_loss	2.68e+04	

time/		
fps	110	
iterations	44	
time_elapsed	25	
total_timesteps	2816	
train/		
approx_kl	0.0	
clip_fraction	0	
clip_range	0.2	
entropy_loss	-5.23e-05	
explained_variance	0	
learning_rate	0.01	
loss	2.07e+04	
n_updates	430	
policy_gradient_loss	3.61e-08	
value_loss	1.04e+05	

time/		
fps	111	
iterations	45	
time_elapsed	25	
total_timesteps	2880	
train/		
approx_kl	0.0	
clip_fraction	0	
clip_range	0.2	
entropy_loss	-5.23e-05	

	explained_variance		0	
	learning_rate		0.01	
	loss		1.31e+05	
	n_updates		440	
	policy_gradient_loss		-3.73e-10	
	value_loss		1.51e+05	

	time/			
	fps		111	
	iterations		46	
	time_elapsed		26	
	total_timesteps		2944	
	train/			
	approx_kl		0.0	
	clip_fraction		0	
	clip_range		0.2	
	entropy_loss		-5.23e-05	
	explained_variance		0	
	learning_rate		0.01	
	loss		1.22e+05	
	n_updates		450	
	policy_gradient_loss		1.3e-08	
	value_loss		3.11e+05	

	time/			
	fps		111	
	iterations		47	
	time_elapsed		26	
	total_timesteps		3008	
	train/			
	approx_kl		0.0	
	clip_fraction		0	
	clip_range		0.2	
	entropy_loss		-5.23e-05	
	explained_variance		0	
	learning_rate		0.01	
	loss		1.93e+05	
	n_updates		460	
	policy_gradient_loss		-7.45e-10	
	value_loss		2.17e+05	

Training Complete


```

[533]: # --- Evaluation ---
obs, _ = env.reset()
state = None
done = False
total_reward = 0
portfolio_values = []
final_action = None

while not done:
    action, state = model.predict(obs, state=state, deterministic=True)
    final_action = action
    obs, reward, done, _, _ = env.step(action)
    total_reward += reward

    current_index = min(env.current_step, len(env.data) - 1)
    current_price = env.data.loc[current_index, 'close']
    unrealized = (
        (current_price - env.entry_price) if env.position == 1 else
        (env.entry_price - current_price) if env.position == -1 else
        0.0
    )
    mtm_equity = env.balance + unrealized
    portfolio_values.append(mtm_equity)

# --- Save trade log ---
env.save_trade_log("trade_log_recurrent.csv")

# --- Final Model Signal ---
signal_str = "BUY" if final_action == 0 else "SELL"
latest_date = env.data['Date'].iloc[env.current_step - 1].strftime("%Y-%m-%d")
print(f"\nLatest model signal at {latest_date}: {signal_str}")

# --- Summary Output ---
print(f"Total Reward: {total_reward:.2f}")
print(f"Final Balance: {env.balance:.2f}")
print("Trade Log (Recurrent PPO):")
print(pd.DataFrame(env.trade_log))
print("Action counts:", env.action_counts)
print("Average reward per action:")
for k, v in env.reward_tracker.items():
    mean_r = np.mean(v) if v else 0
    print(f"Action {k} ({'Long' if k==0 else 'Short'}): {mean_r:.4f}")

```

```

Latest model signal at 2025-04-09: BUY
Total Reward: 4716.27
Final Balance: 16302.66
Trade Log (Recurrent PPO):

```

	EntryDate	ExitDate	Position	EntryPrice	ExitPrice	PnL%	\
0	2023-07-27	2023-07-28	Long	191.5555	194.1430	1.35	
1	2023-07-31	2023-08-01	Short	194.7576	193.9249	0.43	
2	2023-08-02	2023-08-08	Long	190.9210	178.2511	-6.64	
3	2023-08-09	2023-08-11	Short	176.6550	176.4964	0.09	
4	2023-08-14	2023-08-15	Short	178.1543	176.1589	1.13	
..	
171	2025-03-27	2025-04-03	Short	223.8500	203.1900	10.17	
172	2025-04-03	2025-04-07	Long	203.1900	181.4600	-10.69	
173	2025-04-07	2025-04-08	Short	181.4600	172.4200	5.24	
174	2025-04-08	2025-04-09	Long	172.4200	198.8500	15.33	
175	2025-04-09	2025-04-10	Short	198.8500	190.4200	4.43	

	CumulativePnL%	CompoundedFactor	CompoundedPnL%
0	1.35	1.013508	1.35
1	1.78	1.017860	1.79
2	-4.86	0.950313	-4.97
3	-4.77	0.951166	-4.88
4	-3.63	0.961941	-3.81
..
171	40.86	1.445311	44.53
172	30.16	1.290743	29.07
173	35.40	1.358417	35.84
174	50.73	1.566647	56.66
175	55.16	1.636003	63.60

```
[176 rows x 9 columns]
Action counts: {0: 222, 1: 206}
Average reward per action:
Action 0 (Long): 0.1187
Action 1 (Short): 0.0004
```

```
[534]: # --- Get Latest Signal from Model ---
last_index = len(env.data) - 1
env.current_step = last_index
obs = env._get_obs()

# Add batch dimension and run prediction
obs_input = obs[np.newaxis, :]
action, _ = model.predict(obs_input, deterministic=True)

# Convert action to trading signal
if action == 0:
    current_signal = "BUY"
else:
    current_signal = "SELL"
```

```
print(f"\nLatest model signal at {env.data.iloc[last_index]['Date'].date()}:  
↳ {current_signal}")
```

Latest model signal at 2025-04-10: BUY

```
[535]: # --- Plot Performance with Trade Markers ---
buy_hold_line = [env.initial_balance * (p / data['close'].iloc[0]) for p in  
↳ data['close'].iloc[:len(portfolio_values)]]


date_to_step = {row['Date'].strftime('%Y-%m-%d'): i for i, row in data.iloc[:  
↳ len(portfolio_values)].iterrows()}
entry_points = []
exit_points = []

for trade in env.trade_log:
    entry_step = date_to_step.get(trade['EntryDate'])
    exit_step = date_to_step.get(trade['ExitDate'])
    if entry_step is not None and entry_step < len(portfolio_values):
        entry_points.append((entry_step, portfolio_values[entry_step]))
    if exit_step is not None and exit_step < len(portfolio_values):
        exit_points.append((exit_step, portfolio_values[exit_step]))

plt.figure(figsize=(12, 6))
plt.plot(portfolio_values, label="Recurrent PPO Portfolio", linewidth=2)
plt.plot(buy_hold_line, label="Buy & Hold Benchmark", linestyle="--",  
↳ color="gray")


if entry_points:
    entry_steps, entry_vals = zip(*entry_points)
    plt.scatter(entry_steps, entry_vals, color='green', marker='^', s=80,  
↳ label="Trade Entry")
if exit_points:
    exit_steps, exit_vals = zip(*exit_points)
    plt.scatter(exit_steps, exit_vals, color='red', marker='v', s=80,  
↳ label="Trade Exit")

plt.title("Portfolio Value: Recurrent PPO vs. Buy & Hold")
plt.xlabel("Time Step")
plt.ylabel("Portfolio Value")
plt.legend()
plt.grid(True)
plt.tight_layout()
plt.savefig("recurrent_ppo_performance.png")
plt.show()
```



```
[536]: import matplotlib.pyplot as plt
```


```
plt.figure(figsize=(10, 5))
plt.plot(callback.timesteps, callback.episode_rewards, marker='o')
plt.title("Episode Reward Over Time")
plt.xlabel("Timesteps")
plt.ylabel("Episode Reward")
plt.grid(True)
plt.tight_layout()
plt.show()
```



```
[537]: import matplotlib.pyplot as plt
import numpy as np
import pandas as pd

rewards = np.array(env.step_rewards)

# Histogram
plt.hist(rewards, bins=50, alpha=0.7, edgecolor='black')
plt.title("Reward Distribution per Step")
plt.xlabel("Reward")
plt.ylabel("Frequency")
plt.grid(True)
plt.show()
```



```
[538]: import numpy as np
import matplotlib.pyplot as plt
import pandas as pd

rewards = np.array(env.step_rewards)
```

```


# --- Raw volatility metrics ---
reward_std = np.std(rewards)
reward_mean = np.mean(rewards)
reward_variance = np.var(rewards)
reward_range = np.max(rewards) - np.min(rewards)

# --- Rolling volatility ---
window = 20
rolling_std = pd.Series(rewards).rolling(window=window).std()

# --- Plot ---
plt.figure(figsize=(10, 4))
plt.plot(rolling_std, label=f"Rolling Std (window={window})")
plt.axhline(reward_std, color='red', linestyle='--', label='Global Std')
plt.title("Reward Volatility (Rolling Std Dev of Step Rewards)")
plt.xlabel("Step")
plt.ylabel("Rolling Std Dev")
plt.legend()
plt.grid(True)
plt.tight_layout()
plt.show()

# --- Output stats ---
print(f"Mean Step Reward: {reward_mean:.4f}")
print(f"Std Dev of Step Rewards: {reward_std:.4f}")
print(f"Variance: {reward_variance:.4f}")
print(f"Range: {reward_range:.4f}")
print(f"Max Rolling Std ({window}): {rolling_std.max():.4f}")
print(f"Min Rolling Std ({window}): {rolling_std.min():.4f}")

```



Mean Step Reward: 0.0649

Std Dev of Step Rewards: 1.2698
Variance: 1.6125
Range: 5.5483
Max Rolling Std (20): 1.7133
Min Rolling Std (20): 0.7813

```
[539]: import pandas as pd
import matplotlib.pyplot as plt

def compute_switch_rate(trade_log: pd.DataFrame, plot: bool = True):
    if 'Position' not in trade_log.columns or trade_log.empty:
        raise ValueError("Trade log must contain a 'Position' column and be
        ↪non-empty.")

    positions = trade_log['Position'].tolist()
    if len(positions) < 2:
        print("Insufficient trades to compute switching behavior.")
        return 0.0

    switches = sum(1 for i in range(1, len(positions)) if positions[i] !=
    ↪positions[i - 1])
    total_transitions = len(positions) - 1
    switch_rate = switches / total_transitions


    print(f"Total Trades: {len(positions)}")
    print(f"Total Transitions: {total_transitions}")
    print(f"Switches: {switches}")
    print(f"Switch Rate: {switch_rate:.4f}")

    if plot:
        plt.bar(['Same', 'Switch'], [total_transitions - switches, switches],
        ↪color=['gray', 'red'])
        plt.title("Trade Direction Transitions")
        plt.ylabel("Count")
        plt.grid(True, axis='y', linestyle='--', alpha=0.6)
        plt.show()

    return switch_rate

df_trades = pd.DataFrame(env.trade_log)
switch_rate = compute_switch_rate(df_trades)
```

Total Trades: 176
Total Transitions: 175
Switches: 137
Switch Rate: 0.7829



```
[540]: import matplotlib.pyplot as plt


total_actions = sum(env.action_counts.values())

actions = list(env.action_counts.keys())
counts = [env.action_counts[a] for a in actions]
labels = ['Long (0)', 'Short (1)']

proportions = [count / total_actions for count in counts]

plt.bar(labels, proportions, color='gray', edgecolor='black')
plt.title("Action Distribution (Policy Bias)")
plt.ylabel("Proportion of Actions")
plt.ylim(0, 1)
plt.grid(True, axis='y')
plt.show()

for label, count, prop in zip(labels, counts, proportions):
    print(f"{label}: Count = {count}, Proportion = {prop:.2%}")
```

Long (0): Count = 222, Proportion = 51.87%

Short (1): Count = 206, Proportion = 48.13%

— Training with seed 71719 — Seed 71719: Reward = -164.56, Final Balance = 11574.93 Average reward for action Long: -0.0565 Average reward for action Short: 0.1578 Best model updated and saved.

— Training with seed 34867 — Seed 34867: Reward = -4979.62, Final Balance = 6732.89 Average reward for action Long: 0.0093 Average reward for action Short: 0.0643

— Training with seed 57617 — Seed 57617: Reward = -2588.05, Final Balance = 9053.52 Average reward for action Long: 0.2192 Average reward for action Short: -0.2105

— Training with seed 122493 — Seed 122493: Reward = 331.47, Final Balance = 11925.18 Average reward for action Long: 0.2236 Average reward for action Short: -0.3920 Best model updated and saved.

— Training with seed 47139 — Seed 47139: Reward = 1809.22, Final Balance = 13726.43 Average reward for action Long: 0.0164 Average reward for action Short: -0.1107 Best model updated and saved.

— Training with seed 128030 — Seed 128030: Reward = 4716.23, Final Balance = 16302.62 Average reward for action Long: 0.1187 Average reward for action Short: 0.0004 Best model updated and saved.

— Training with seed 45301 — Seed 45301: Reward = 219.28, Final Balance = 11712.43 Average reward for action Long: 0.2243 Average reward for action Short: 0.0845

— Training with seed 58376 — Seed 58376: Reward = -4153.13, Final Balance = 7403.73 Average reward for action Long: 0.0082 Average reward for action Short: -0.2599

— Training with seed 119695 — Seed 119695: Reward = 4637.83, Final Balance = 16022.76 Average reward for action Long: 0.4593 Average reward for action Short: 0.1460

— Training with seed 37431 — Seed 37431: Reward = -1709.18, Final Balance = 9801.24 Average reward for action Long: 0.1996 Average reward for action Short: -0.0553

— Training with seed 107833 — Seed 107833: Reward = -1798.95, Final Balance = 9891.31 Average reward for action Long: 0.0503 Average reward for action Short: 0.0821

— Seed Sweep Complete — Best Seed: 128030 Best Total Reward: 4716.23 Best Final Balance: 16302.62

```
[541]: model.save("AAPL_best_model")
```

1.11 META

```
[542]: from datetime import datetime, timedelta

# Add +1 day to end date
default_end_date = (datetime.now() + timedelta(days=1)).strftime("%Y-%m-%d")
default_start_date = (datetime.now() - timedelta(days=665)).strftime("%Y-%m-%d")
```

```
[543]: # Parameters
ticker = "META"
start_date = default_start_date
end_date = default_end_date
data_filename = "full_data.csv"
signals_filename = "signals_data.csv"

save_historical_data(
    ticker=ticker,
    start_date=start_date,
    end_date=end_date,
    data_filename=data_filename,
    signals_filename=signals_filename
)
```

```
[*****100%*****] 1 of 1 completed
```

```
[543]: 'Saved full_data.csv and signals_data.csv successfully.'
```

```
[544]: import pandas as pd

# Preview full data
df_full = pd.read_csv("full_data.csv", parse_dates=["Date"])
```

```
print(df_full.head())

# Preview signal data
df_signals = pd.read_csv("signals_data.csv", parse_dates=["Date"])
print(df_signals.head())
```

	Date	close	high	low	open	volume	\
0	2023-06-16	279.681885	286.499759	278.815971	283.414294	43102500	
1	2023-06-20	282.996216	283.464012	274.924274	277.422509	20701600	
2	2023-06-21	280.318848	282.667762	277.054205	282.199966	20556200	
3	2023-06-22	283.543640	283.921862	276.486903	277.770830	17563100	
4	2023-06-23	287.375610	288.311203	277.641488	280.189477	50988400	

	EMA_50	EMA_100	EMA_200	EMA_500	...	RSI_Sell	MCDX_Buy	\
0	279.681885	279.681885	279.681885	279.681885	...	1	0	
1	279.811859	279.747515	279.714863	279.695116	...	1	0	
2	279.831740	279.758829	279.720873	279.697606	...	1	0	
3	279.977305	279.833775	279.758910	279.712959	...	1	0	
4	280.267435	279.983119	279.834699	279.743548	...	0	0	

	MCDX_Sell	DSS_Buy	DSS_Sell	ZeroLag	MACD_Buy	ZeroLag	MACD_Sell	\
0	1	0	1		0		1	
1	1	0	1		1		0	
2	1	0	1		1		0	
3	1	0	1		1		0	
4	1	0	1		1		0	

	Basic	MACD_Buy	Basic	MACD_Sell	OverallTrade
0		0		1	Sell
1		1		0	Sell
2		1		0	Sell
3		1		0	Sell
4		1		0	Buy

[5 rows x 42 columns]

	Date	Signal	Z-Score
0	2023-06-20	Basic MACD Buy	NaN
1	2023-06-20	ZLMA Buy	NaN
2	2023-06-20	ZeroLag MACD Buy	NaN
3	2023-06-21	RSI Sell	NaN
4	2023-06-23	RSI Buy	NaN

```
[545]: from stable_baselines3.common.callbacks import BaseCallback
```

```
class RewardTrackingCallback(BaseCallback):
    def __init__(self, verbose=0):
        super().__init__(verbose)
```

```

        self.episode_rewards = []
        self.timesteps = []

    def _on_step(self) -> bool:
        if self.locals.get("done") is not None and any(self.locals["done"]):
            ep_rew = self.locals["rewards"]
            self.episode_rewards.append(sum(ep_rew))
            self.timesteps.append(self.num_timesteps)
        return True

```

```

[546]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import gymnasium as gym
from gymnasium import spaces
from sb3_contrib import RecurrentPPO
from sb3_contrib.ppo_recurrent.policies import MlpLstmPolicy
from stable_baselines3.common.vec_env import DummyVecEnv
import torch
import random

# --- Fixed Seed ---
SEED = 73352
np.random.seed(SEED)
random.seed(SEED)
torch.manual_seed(SEED)

# --- Load & Clean Data ---
data_path = "full_data.csv"
data = pd.read_csv(data_path, parse_dates=["Date"])

required_cols = [
    'open', 'high', 'low', 'close',
    'zlma', 'ema_value',
    'DSSb', 'DSSsignal',
    'rsi_ma_base',
    'ZeroLagMACD', 'ZeroLagMACD_signal',
    'basicMACD', 'basicMACD_signal',
    'ZScore',
    'ZLMA_Buy', 'ZLMA_Sell',
    'RSI_Buy', 'RSI_Sell',
    'MCDX_Buy', 'MCDX_Sell',
    'DSS_Buy', 'DSS_Sell',
    'ZeroLag MACD_Buy', 'ZeroLag MACD_Sell',
    'Basic MACD_Buy', 'Basic MACD_Sell'
]

```

```

missing = [col for col in required_cols if col not in data.columns]
assert not missing, f"Missing required columns: {missing}"

data = data.dropna(subset=required_cols).reset_index(drop=True)

SWITCH_COST = 1.0
TRANSACTION_COST = 0.001

ACTIVATION = torch.nn.Tanh
# ACTIVATION = partial(nn.LeakyReLU, negative_slope=0.01)

policy_kwargs = dict(
    activation_fn=ACTIVATION
)

# --- Main Trading Environment ---
class TradingEnvRL(gym.Env):
    metadata = {'render_modes': ['human']}

    def __init__(self, data, initial_balance=10000, hold_cost=0.02,
                 volatility_window=14, exploration_steps=500,
                 switch_cost=1.0, reentry_threshold=0.01,
                 dynamic_threshold=False, loss_penalty=0.75, drawdown_penalty=5.
↪0, large_loss_threshold=-2.0):

        super().__init__()
        self.loss_penalty = loss_penalty
        self.drawdown_penalty = drawdown_penalty
        self.large_loss_threshold = large_loss_threshold
        self.data = data.reset_index(drop=True).copy()
        self.n_steps = len(self.data)
        self.initial_balance = initial_balance
        self.hold_cost = hold_cost
        self.volatility_window = volatility_window
        self.exploration_steps = exploration_steps
        self.switch_cost = switch_cost
        self.reentry_threshold = reentry_threshold
        self.dynamic_threshold = dynamic_threshold
        self.step_rewards = []

        self.feature_cols = [
            'open', 'high', 'low', 'close',
            'basicMACD', 'basicMACD_signal',
            'Basic MACD_Buy', 'Basic MACD_Sell'
        ]

        obs_dim = len(self.feature_cols) + 1

```

```

        self.observation_space = spaces.Box(low=-np.inf, high=np.inf,
↪shape=(obs_dim,), dtype=np.float32)
        self.action_space = spaces.Discrete(2) # 0=Long, 1=Short

        self._compute_volatility_limit()

    def _compute_volatility_limit(self):
        returns = self.data['close'].pct_change()
        self.data['volatility'] = returns.rolling(self.volatility_window).std()
        self.data['adaptive_hold'] = (10 / (self.data['volatility'] * 100)).
↪clip(lower=3, upper=20).fillna(10).astype(int)

    def reset(self, seed=None, options=None):
        if seed is not None:
            np.random.seed(seed)
            random.seed(seed)
            torch.manual_seed(seed)

        self.current_step = 0
        self.position = 0
        self.entry_price = 0.0
        self.entry_date = None
        self.hold_counter = 0
        self.switch_count = 0
        self.balance = self.initial_balance
        self.cumulative_pnl = 0.0
        self.trade_log = []
        self.action_counts = {0: 0, 1: 0}
        self.reward_tracker = {0: [], 1: []}
        self.consecutive_losses = 0
        self.equity_curve = [self.initial_balance]
        self.step_rewards = []
        return self._get_obs(), {}

    def _get_obs(self):
        row = self.data.iloc[self.current_step]
        features = row[self.feature_cols].values.astype(np.float32)
        pos_feature = np.array([self.position], dtype=np.float32)
        return np.concatenate([features, pos_feature])

    def _force_close(self):
        row = self.data.iloc[self.current_step]
        current_price = float(row['close'])
        current_date = row['Date'].strftime("%Y-%m-%d")

        if self.position == 0:
            return 0.0

```

```

# --- Core Return Logic ---
trade_pct = ((current_price / self.entry_price - 1) * 100) if self.
↪position == 1 else ((self.entry_price / current_price - 1) * 100)
pos_str = 'Long' if self.position == 1 else 'Short'
gross_return = trade_pct / 100
transaction_cost = TRANSACTION_COST * current_price

old_balance = self.balance
self.balance -= transaction_cost
self.balance *= (1 + gross_return)
net_profit = self.balance - old_balance

reward = net_profit

# --- Track Trade History ---
self.cumulative_pnl += trade_pct
compounded_pnl = (self.trade_log[-1]['CompoundedFactor'] * (1 + ↪
↪gross_return)) if self.trade_log else (1 + gross_return)
compounded_pnl_pct = (compounded_pnl - 1) * 100

self.trade_log.append({
    'EntryDate': self.entry_date,
    'ExitDate': current_date,
    'Position': pos_str,
    'EntryPrice': round(self.entry_price, 4),
    'ExitPrice': round(current_price, 4),
    'PnL%': round(trade_pct, 2),
    'CumulativePnL%': round(self.cumulative_pnl, 2),
    'CompoundedFactor': compounded_pnl,
    'CompoundedPnL%': round(compounded_pnl_pct, 2)
})

# --- Penalty for Large Loss ---
if trade_pct < -2.0:
    reward += trade_pct * 2 # Stronger penalty for large loss

# --- Track & Penalize Consecutive Losses ---
if trade_pct < 0:
    self.consecutive_losses += 1
    reward -= self.consecutive_losses * 0.75 # Growing penalty
else:
    self.consecutive_losses = 0

# --- Optional: Drawdown Penalty ---
self.equity_curve.append(self.balance)
max_balance = max(self.equity_curve)

```

```

    if max_balance > 0:
        drawdown = (max_balance - self.balance) / max_balance
        reward -= 5.0 * drawdown # Penalize deeper drawdown

    # --- Reset State ---
    self.position = 0
    self.entry_price = 0.0
    self.entry_date = None
    self.hold_counter = 0

    return reward

def step(self, action):
    if isinstance(action, np.ndarray):
        action = int(action.item())

    row = self.data.iloc[self.current_step]
    current_price = float(row['close'])
    current_date = row['Date'].strftime("%Y-%m-%d")

    if self.current_step < self.exploration_steps:
        action = self.action_space.sample()

    self.action_counts[action] += 1
    reward = 0.0
    target_position = 1 if action == 0 else -1

    # === Holding Same Position ===
    if self.position == target_position:
        self.hold_counter += 1
        price_change = (current_price - self.entry_price) / self.entry_price
        step_return = price_change if self.position == 1 else -price_change

        step_reward = np.sign(step_return) * np.sqrt(abs(step_return)) * 10
        step_reward = np.clip(step_reward, -50, 50)
        step_reward -= self.hold_cost
        reward += step_reward
        self.reward_tracker[action].append(step_reward)
        self.step_rewards.append(reward)

    # === New Position from Flat ===
    elif self.position == 0:
        self.position = target_position
        self.entry_price = current_price
        self.entry_date = current_date
        self.hold_counter = 1
        reward -= TRANSACTION_COST * current_price

```



```

# === Switch Position ===
else:
    hold_penalty = max(0, 3 - self.hold_counter) * 5.0
    switch_penalty = self.switch_cost + hold_penalty
    reward += self._force_close()
    reward -= switch_penalty

# === Dynamic Re-entry Threshold ===
if self.current_step > 0:
    prev_close = self.data.iloc[self.current_step - 1]['close']
    recent_return = abs((current_price - prev_close) / prev_close)

    # Volatility-aware reentry threshold
    recent_volatility = row['volatility']
    reentry_threshold = 0.5 * recent_volatility
else:
    recent_return = 0
    reentry_threshold = 0.01 # Fallback

if recent_return > reentry_threshold:
    self.position = target_position
    self.entry_price = current_price
    self.entry_date = current_date
    self.hold_counter = 1
    reward -= TRANSACTION_COST * current_price

self.current_step += 1
terminated = self.current_step >= self.n_steps - 1

if terminated and self.position != 0:
    reward += self._force_close()

obs = self._get_obs() if not terminated else np.zeros(self.
↪observation_space.shape, dtype=np.float32)
return obs, reward, terminated, False, {}

def render(self):
    print(f"Step: {self.current_step}, Position: {self.position}, Balance:␣
↪{self.balance:.2f}")

def save_trade_log(self, filename="trade_log.csv"):
    df = pd.DataFrame(self.trade_log)
    if "CompoundedFactor" in df.columns:
        df = df.drop(columns=["CompoundedFactor"])
    df.to_csv(filename, index=False)

```

```

# --- Training ---
env = TradingEnvRL(data, initial_balance=10000)
vec_env = DummyVecEnv([lambda: env])
vec_env.seed(SEED)

model = RecurrentPPO(
    policy=MlpLstmPolicy,
    env=vec_env,
    verbose=1,
    n_steps=64,
    batch_size=32,
    learning_rate=0.01,
    gamma=0.99,
    ent_coef=0.01,
    seed=SEED,
    policy_kwargs=policy_kwargs
)

callback = RewardTrackingCallback()

model.learn(total_timesteps=3000, callback=callback)
print(f"\nTraining Complete")

```

Using cuda device

```

-----
| time/          |      |
|   fps          | 268  |
| iterations     | 1    |
| time_elapsed   | 0    |
| total_timesteps | 64   |
-----

```

```

-----
| time/          |      |
|   fps          | 107  |
| iterations     | 2    |
| time_elapsed   | 1    |
| total_timesteps | 128  |
| train/         |      |
| approx_kl      | 0.026080996 |
| clip_fraction  | 0.317 |
| clip_range     | 0.2   |
| entropy_loss   | -0.677 |
| explained_variance | 8.03e-05 |
| learning_rate  | 0.01  |
| loss           | 4.37e+04 |
| n_updates      | 10    |
| policy_gradient_loss | -0.00713 |
| value_loss     | 1.18e+05 |
-----

```

time/		
fps	113	
iterations	3	
time_elapsed	1	
total_timesteps	192	
train/		
approx_kl	0.010536695	
clip_fraction	0.0906	
clip_range	0.2	
entropy_loss	-0.69	
explained_variance	-3.81e-06	
learning_rate	0.01	
loss	2.89e+05	
n_updates	20	
policy_gradient_loss	-0.000933	
value_loss	5.64e+05	

time/		
fps	117	
iterations	4	
time_elapsed	2	
total_timesteps	256	
train/		
approx_kl	0.011035234	
clip_fraction	0.15	
clip_range	0.2	
entropy_loss	-0.686	
explained_variance	0	
learning_rate	0.01	
loss	2.61e+05	
n_updates	30	
policy_gradient_loss	-0.00259	
value_loss	3.43e+05	

time/		
fps	118	
iterations	5	
time_elapsed	2	
total_timesteps	320	
train/		
approx_kl	0.010167208	
clip_fraction	0.0953	
clip_range	0.2	
entropy_loss	-0.674	

explained_variance	0	
learning_rate	0.01	
loss	1.12e+05	
n_updates	40	
policy_gradient_loss	0.00344	
value_loss	1.66e+05	

time/		
fps	113	
iterations	6	
time_elapsed	3	
total_timesteps	384	
train/		
approx_kl	0.010145061	
clip_fraction	0.0469	
clip_range	0.2	
entropy_loss	-0.671	
explained_variance	1.19e-07	
learning_rate	0.01	
loss	4.09e+04	
n_updates	50	
policy_gradient_loss	0.00399	
value_loss	1.79e+05	

time/		
fps	111	
iterations	7	
time_elapsed	4	
total_timesteps	448	
train/		
approx_kl	0.011051932	
clip_fraction	0.109	
clip_range	0.2	
entropy_loss	-0.646	
explained_variance	1.19e-07	
learning_rate	0.01	
loss	3.69e+05	
n_updates	60	
policy_gradient_loss	-0.00431	
value_loss	5.29e+05	

time/		
fps	100	
iterations	8	
time_elapsed	5	

	total_timesteps		512	
	train/			
	approx_kl		0.0011400054	
	clip_fraction		0	
	clip_range		0.2	
	entropy_loss		-0.648	
	explained_variance		0	
	learning_rate		0.01	
	loss		2.14e+06	
	n_updates		70	
	policy_gradient_loss		0.00464	
	value_loss		2.89e+06	

	time/			
	fps		102	
	iterations		9	
	time_elapsed		5	
	total_timesteps		576	
	train/			
	approx_kl		0.013453225	
	clip_fraction		0	
	clip_range		0.2	
	entropy_loss		-0.67	
	explained_variance		1.19e-07	
	learning_rate		0.01	
	loss		8.85e+04	
	n_updates		80	
	policy_gradient_loss		-0.000388	
	value_loss		3.56e+05	

	time/			
	fps		104	
	iterations		10	
	time_elapsed		6	
	total_timesteps		640	
	train/			
	approx_kl		0.0029394957	
	clip_fraction		0.148	
	clip_range		0.2	
	entropy_loss		-0.553	
	explained_variance		-1.19e-07	
	learning_rate		0.01	
	loss		1.41e+05	
	n_updates		90	
	policy_gradient_loss		-0.000916	
	value_loss		3.26e+05	

time/		
fps	106	
iterations	11	
time_elapsed	6	
total_timesteps	704	
train/		
approx_kl	0.0072773485	
clip_fraction	0.0672	
clip_range	0.2	
entropy_loss	-0.549	
explained_variance	0	
learning_rate	0.01	
loss	7.68e+04	
n_updates	100	
policy_gradient_loss	-0.00392	
value_loss	1.76e+05	

time/		
fps	108	
iterations	12	
time_elapsed	7	
total_timesteps	768	
train/		
approx_kl	0.041851282	
clip_fraction	0.173	
clip_range	0.2	
entropy_loss	-0.484	
explained_variance	1.19e-07	
learning_rate	0.01	
loss	2.93e+04	
n_updates	110	
policy_gradient_loss	0.00426	
value_loss	4.1e+04	

time/		
fps	109	
iterations	13	
time_elapsed	7	
total_timesteps	832	
train/		
approx_kl	0.37481505	
clip_fraction	0.0906	
clip_range	0.2	
entropy_loss	-0.152	

	explained_variance	0	
	learning_rate	0.01	
	loss	1.46e+04	
	n_updates	120	
	policy_gradient_loss	0.0113	
	value_loss	5.84e+04	

	time/		
	fps	110	
	iterations	14	
	time_elapsed	8	
	total_timesteps	896	
	train/		
	approx_kl	2.3841858e-07	
	clip_fraction	0	
	clip_range	0.2	
	entropy_loss	-0.00488	
	explained_variance	0	
	learning_rate	0.01	
	loss	2.86e+04	
	n_updates	130	
	policy_gradient_loss	-1.12e-08	
	value_loss	5.58e+04	

	time/		
	fps	105	
	iterations	15	
	time_elapsed	9	
	total_timesteps	960	
	train/		
	approx_kl	0.0	
	clip_fraction	0	
	clip_range	0.2	
	entropy_loss	-0.00304	
	explained_variance	0	
	learning_rate	0.01	
	loss	9.99e+04	
	n_updates	140	
	policy_gradient_loss	8.72e-08	
	value_loss	2.33e+05	

	time/		
	fps	106	
	iterations	16	
	time_elapsed	9	

	total_timesteps		1024	
	train/			
	approx_kl		0.0	
	clip_fraction		0	
	clip_range		0.2	
	entropy_loss		-0.00291	
	explained_variance		0	
	learning_rate		0.01	
	loss		4.15e+04	
	n_updates		150	
	policy_gradient_loss		1.04e-08	
	value_loss		5.15e+04	

	time/			
	fps		107	
	iterations		17	
	time_elapsed		10	
	total_timesteps		1088	
	train/			
	approx_kl		8.381903e-09	
	clip_fraction		0	
	clip_range		0.2	
	entropy_loss		-0.00297	
	explained_variance		-1.19e-07	
	learning_rate		0.01	
	loss		3.91e+05	
	n_updates		160	
	policy_gradient_loss		-5.22e-09	
	value_loss		5.28e+05	

	time/			
	fps		108	
	iterations		18	
	time_elapsed		10	
	total_timesteps		1152	
	train/			
	approx_kl		0.0	
	clip_fraction		0	
	clip_range		0.2	
	entropy_loss		-0.00299	
	explained_variance		0	
	learning_rate		0.01	
	loss		1.13e+05	
	n_updates		170	
	policy_gradient_loss		-5.59e-09	
	value_loss		2.2e+05	

time/		
fps	109	
iterations	19	
time_elapsed	11	
total_timesteps	1216	
train/		
approx_kl	0.0	
clip_fraction	0	
clip_range	0.2	
entropy_loss	-0.003	
explained_variance	0	
learning_rate	0.01	
loss	2.84e+04	
n_updates	180	
policy_gradient_loss	-2.79e-09	
value_loss	3.91e+05	

time/		
fps	110	
iterations	20	
time_elapsed	11	
total_timesteps	1280	
train/		
approx_kl	2.9802322e-08	
clip_fraction	0	
clip_range	0.2	
entropy_loss	-0.00303	
explained_variance	0	
learning_rate	0.01	
loss	1.7e+04	
n_updates	190	
policy_gradient_loss	-4.28e-09	
value_loss	5.17e+04	

time/		
fps	110	
iterations	21	
time_elapsed	12	
total_timesteps	1344	
train/		
approx_kl	2.9802322e-08	
clip_fraction	0	
clip_range	0.2	
entropy_loss	-0.00309	

	explained_variance	0	
	learning_rate	0.01	
	loss	4.57e+04	
	n_updates	200	
	policy_gradient_loss	8.94e-09	
	value_loss	8.93e+04	

	time/		
	fps	108	
	iterations	22	
	time_elapsed	13	
	total_timesteps	1408	
	train/		
	approx_kl	0.0	
	clip_fraction	0	
	clip_range	0.2	
	entropy_loss	-0.00325	
	explained_variance	-1.19e-07	
	learning_rate	0.01	
	loss	2.93e+04	
	n_updates	210	
	policy_gradient_loss	-4.58e-08	
	value_loss	1.03e+05	

	time/		
	fps	108	
	iterations	23	
	time_elapsed	13	
	total_timesteps	1472	
	train/		
	approx_kl	0.0	
	clip_fraction	0	
	clip_range	0.2	
	entropy_loss	-0.00338	
	explained_variance	0	
	learning_rate	0.01	
	loss	6.76e+04	
	n_updates	220	
	policy_gradient_loss	2.11e-08	
	value_loss	1.84e+05	

	time/		
	fps	109	
	iterations	24	
	time_elapsed	14	

total_timesteps	1536
train/	
approx_kl	-5.9604645e-08
clip_fraction	0
clip_range	0.2
entropy_loss	-0.00341
explained_variance	0
learning_rate	0.01
loss	2.74e+05
n_updates	230
policy_gradient_loss	-2.37e-08
value_loss	5.65e+05

time/	
fps	110
iterations	25
time_elapsed	14
total_timesteps	1600
train/	
approx_kl	0.0
clip_fraction	0
clip_range	0.2
entropy_loss	-0.00342
explained_variance	-1.19e-07
learning_rate	0.01
loss	3.91e+05
n_updates	240
policy_gradient_loss	-3.35e-08
value_loss	1.11e+06

time/	
fps	110
iterations	26
time_elapsed	15
total_timesteps	1664
train/	
approx_kl	-5.9604645e-08
clip_fraction	0
clip_range	0.2
entropy_loss	-0.00343
explained_variance	0
learning_rate	0.01
loss	9.65e+04
n_updates	250
policy_gradient_loss	1.66e-08
value_loss	1.71e+05

```

-----
| time/          |          |
|   fps          |   111    |
|   iterations    |   27     |
|   time_elapsed  |   15     |
|   total_timesteps | 1728    |
| train/         |          |
|   approx_kl     |   0.0    |
|   clip_fraction |   0      |
|   clip_range    |   0.2    |
|   entropy_loss  | -0.00348 |
|   explained_variance | 0        |
|   learning_rate |   0.01   |
|   loss          | 9.12e+04 |
|   n_updates     | 260      |
|   policy_gradient_loss | 3.07e-09 |
|   value_loss    | 1.37e+05 |
-----

```

```

-----
| time/          |          |
|   fps          |   108    |
|   iterations    |   28     |
|   time_elapsed  |   16     |
|   total_timesteps | 1792    |
| train/         |          |
|   approx_kl     |   0.0    |
|   clip_fraction |   0      |
|   clip_range    |   0.2    |
|   entropy_loss  | -0.0036  |
|   explained_variance | 1.19e-07 |
|   learning_rate |   0.01   |
|   loss          | 8.36e+05 |
|   n_updates     | 270      |
|   policy_gradient_loss | -8.57e-09 |
|   value_loss    | 2.06e+06 |
-----

```

```

-----
| time/          |          |
|   fps          |   109    |
|   iterations    |   29     |
|   time_elapsed  |   16     |
|   total_timesteps | 1856    |
| train/         |          |
|   approx_kl     | 4.1909516e-08 |
|   clip_fraction |   0      |
|   clip_range    |   0.2    |
|   entropy_loss  | -0.00369 |

```

	explained_variance		0	
	learning_rate		0.01	
	loss		2.51e+04	
	n_updates		280	
	policy_gradient_loss		3.17e-08	
	value_loss		8.6e+04	

	time/			
	fps		109	
	iterations		30	
	time_elapsed		17	
	total_timesteps		1920	
	train/			
	approx_kl		2.7008355e-08	
	clip_fraction		0	
	clip_range		0.2	
	entropy_loss		-0.00416	
	explained_variance		0	
	learning_rate		0.01	
	loss		1.34e+05	
	n_updates		290	
	policy_gradient_loss		1.97e-08	
	value_loss		1.29e+06	

	time/			
	fps		110	
	iterations		31	
	time_elapsed		17	
	total_timesteps		1984	
	train/			
	approx_kl		1.0244548e-08	
	clip_fraction		0	
	clip_range		0.2	
	entropy_loss		-0.00432	
	explained_variance		0	
	learning_rate		0.01	
	loss		4.57e+05	
	n_updates		300	
	policy_gradient_loss		-4e-08	
	value_loss		5.3e+05	

	time/			
	fps		110	
	iterations		32	
	time_elapsed		18	

	total_timesteps		2048	
	train/			
	approx_kl		2.142042e-08	
	clip_fraction		0	
	clip_range		0.2	
	entropy_loss		-0.00438	
	explained_variance		0	
	learning_rate		0.01	
	loss		1.78e+05	
	n_updates		310	
	policy_gradient_loss		-1.43e-08	
	value_loss		3.05e+05	

	time/			
	fps		111	
	iterations		33	
	time_elapsed		18	
	total_timesteps		2112	
	train/			
	approx_kl		1.9557774e-08	
	clip_fraction		0	
	clip_range		0.2	
	entropy_loss		-0.00443	
	explained_variance		-1.19e-07	
	learning_rate		0.01	
	loss		1.34e+05	
	n_updates		320	
	policy_gradient_loss		-2.57e-08	
	value_loss		4.95e+05	

	time/			
	fps		111	
	iterations		34	
	time_elapsed		19	
	total_timesteps		2176	
	train/			
	approx_kl		1.4901161e-08	
	clip_fraction		0	
	clip_range		0.2	
	entropy_loss		-0.00447	
	explained_variance		0	
	learning_rate		0.01	
	loss		1.61e+05	
	n_updates		330	
	policy_gradient_loss		3.51e-08	
	value_loss		2.45e+05	

time/		
fps	109	
iterations	35	
time_elapsed	20	
total_timesteps	2240	
train/		
approx_kl	1.0244548e-08	
clip_fraction	0	
clip_range	0.2	
entropy_loss	-0.0046	
explained_variance	1.19e-07	
learning_rate	0.01	
loss	1.05e+06	
n_updates	340	
policy_gradient_loss	-1.88e-07	
value_loss	1.08e+06	

time/		
fps	110	
iterations	36	
time_elapsed	20	
total_timesteps	2304	
train/		
approx_kl	7.450581e-09	
clip_fraction	0	
clip_range	0.2	
entropy_loss	-0.00487	
explained_variance	0	
learning_rate	0.01	
loss	4.27e+04	
n_updates	350	
policy_gradient_loss	-8.57e-09	
value_loss	1.94e+05	

time/		
fps	110	
iterations	37	
time_elapsed	21	
total_timesteps	2368	
train/		
approx_kl	0.00044056587	
clip_fraction	0.00313	
clip_range	0.2	
entropy_loss	-0.00543	

	explained_variance	0	
	learning_rate	0.01	
	loss	3.41e+05	
	n_updates	360	
	policy_gradient_loss	3.41e-05	
	value_loss	4.62e+05	

	time/		
	fps	111	
	iterations	38	
	time_elapsed	21	
	total_timesteps	2432	
	train/		
	approx_kl	2.7008355e-08	
	clip_fraction	0	
	clip_range	0.2	
	entropy_loss	-0.00629	
	explained_variance	0	
	learning_rate	0.01	
	loss	1.7e+05	
	n_updates	370	
	policy_gradient_loss	-1.27e-08	
	value_loss	4.33e+05	

	time/		
	fps	111	
	iterations	39	
	time_elapsed	22	
	total_timesteps	2496	
	train/		
	approx_kl	2.9802322e-08	
	clip_fraction	0	
	clip_range	0.2	
	entropy_loss	-0.00647	
	explained_variance	0	
	learning_rate	0.01	
	loss	1.74e+05	
	n_updates	380	
	policy_gradient_loss	7.82e-09	
	value_loss	1.89e+05	

	time/		
	fps	111	
	iterations	40	
	time_elapsed	22	

	total_timesteps	2560	
	train/		
	approx_kl	0.0	
	clip_fraction	0	
	clip_range	0.2	
	entropy_loss	-0.00658	
	explained_variance	0	
	learning_rate	0.01	
	loss	7.42e+04	
	n_updates	390	
	policy_gradient_loss	-1.49e-09	
	value_loss	1.54e+05	

	time/		
	fps	112	
	iterations	41	
	time_elapsed	23	
	total_timesteps	2624	
	train/		
	approx_kl	5.9604645e-08	
	clip_fraction	0	
	clip_range	0.2	
	entropy_loss	-0.00682	
	explained_variance	-1.19e-07	
	learning_rate	0.01	
	loss	3.38e+04	
	n_updates	400	
	policy_gradient_loss	-2.61e-09	
	value_loss	9.8e+04	

	time/		
	fps	110	
	iterations	42	
	time_elapsed	24	
	total_timesteps	2688	
	train/		
	approx_kl	2.8871e-08	
	clip_fraction	0	
	clip_range	0.2	
	entropy_loss	-0.00709	
	explained_variance	0	
	learning_rate	0.01	
	loss	2.76e+05	
	n_updates	410	
	policy_gradient_loss	-2.01e-08	
	value_loss	3.19e+05	

time/		
fps	110	
iterations	43	
time_elapsed	24	
total_timesteps	2752	
train/		
approx_kl	0.0	
clip_fraction	0	
clip_range	0.2	
entropy_loss	-0.0073	
explained_variance	0	
learning_rate	0.01	
loss	1.12e+05	
n_updates	420	
policy_gradient_loss	-5.59e-10	
value_loss	2.83e+05	

time/		
fps	111	
iterations	44	
time_elapsed	25	
total_timesteps	2816	
train/		
approx_kl	2.9802322e-08	
clip_fraction	0	
clip_range	0.2	
entropy_loss	-0.00751	
explained_variance	-1.19e-07	
learning_rate	0.01	
loss	2.62e+05	
n_updates	430	
policy_gradient_loss	-1.53e-08	
value_loss	6.87e+05	

time/		
fps	111	
iterations	45	
time_elapsed	25	
total_timesteps	2880	
train/		
approx_kl	2.9802322e-08	
clip_fraction	0	
clip_range	0.2	
entropy_loss	-0.00772	

	explained_variance	0	
	learning_rate	0.01	
	loss	3.82e+04	
	n_updates	440	
	policy_gradient_loss	1.86e-09	
	value_loss	2.73e+05	

	time/		
	fps	111	
	iterations	46	
	time_elapsed	26	
	total_timesteps	2944	
	train/		
	approx_kl	2.9802322e-08	
	clip_fraction	0	
	clip_range	0.2	
	entropy_loss	-0.008	
	explained_variance	1.19e-07	
	learning_rate	0.01	
	loss	5.89e+04	
	n_updates	450	
	policy_gradient_loss	3.43e-08	
	value_loss	1.12e+05	

	time/		
	fps	112	
	iterations	47	
	time_elapsed	26	
	total_timesteps	3008	
	train/		
	approx_kl	0.0	
	clip_fraction	0	
	clip_range	0.2	
	entropy_loss	-0.00849	
	explained_variance	0	
	learning_rate	0.01	
	loss	7.37e+04	
	n_updates	460	
	policy_gradient_loss	4.84e-09	
	value_loss	1.03e+05	

Training Complete

```

[547]: # --- Evaluation ---
obs, _ = env.reset()
state = None
done = False
total_reward = 0
portfolio_values = []
final_action = None

while not done:
    action, state = model.predict(obs, state=state, deterministic=True)
    final_action = action
    obs, reward, done, _, _ = env.step(action)
    total_reward += reward

    current_index = min(env.current_step, len(env.data) - 1)
    current_price = env.data.loc[current_index, 'close']
    unrealized = (
        (current_price - env.entry_price) if env.position == 1 else
        (env.entry_price - current_price) if env.position == -1 else
        0.0
    )
    mtm_equity = env.balance + unrealized
    portfolio_values.append(mtm_equity)

# --- Save trade log ---
env.save_trade_log("trade_log_recurrent.csv")

# --- Final Model Signal ---
signal_str = "BUY" if final_action == 0 else "SELL"
latest_date = env.data['Date'].iloc[env.current_step - 1].strftime("%Y-%m-%d")
print(f"\nLatest model signal at {latest_date}: {signal_str}")

# --- Summary Output ---
print(f"Total Reward: {total_reward:.2f}")
print(f"Final Balance: {env.balance:.2f}")
print("Trade Log (Recurrent PPO):")
print(pd.DataFrame(env.trade_log))
print("Action counts:", env.action_counts)
print("Average reward per action:")
for k, v in env.reward_tracker.items():
    mean_r = np.mean(v) if v else 0
    print(f"Action {k} ({'Long' if k==0 else 'Short'}): {mean_r:.4f}")

```

```

Latest model signal at 2025-04-09: SELL
Total Reward: 11332.70
Final Balance: 22961.60
Trade Log (Recurrent PPO):

```

	EntryDate	ExitDate	Position	EntryPrice	ExitPrice	PnL%	\
0	2023-07-27	2023-07-28	Short	310.2478	323.9532	-4.23	
1	2023-07-31	2023-08-03	Short	317.1055	311.7209	1.73	
2	2023-08-04	2023-08-07	Short	309.2724	315.0750	-1.84	
3	2023-08-08	2023-08-09	Short	311.1735	303.7783	2.43	
4	2023-08-10	2023-08-16	Long	304.3058	292.9095	-3.75	
..	
172	2025-03-18	2025-03-19	Short	582.3600	584.0600	-0.29	
173	2025-03-20	2025-03-24	Long	586.0000	618.8500	5.61	
174	2025-03-24	2025-03-31	Short	618.8500	576.3600	7.37	
175	2025-04-01	2025-04-07	Short	586.0000	516.2500	13.51	
176	2025-04-07	2025-04-10	Long	516.2500	546.2900	5.82	

	CumulativePnL%	CompoundedFactor	CompoundedPnL%
0	-4.23	0.957693	-4.23
1	-2.50	0.974236	-2.58
2	-4.34	0.956294	-4.37
3	-1.91	0.979574	-2.04
4	-5.66	0.942889	-5.71
..
172	60.51	1.697206	69.72
173	66.12	1.792348	79.23
174	73.49	1.924482	92.45
175	87.00	2.184497	118.45
176	92.82	2.311610	131.16

[177 rows x 9 columns]
Action counts: {0: 214, 1: 214}
Average reward per action:
Action 0 (Long): 0.6296
Action 1 (Short): 0.1092

```
[548]: # --- Get Latest Signal from Model ---
last_index = len(env.data) - 1
env.current_step = last_index
obs = env._get_obs()


# Add batch dimension and run prediction
obs_input = obs[np.newaxis, :]
action, _ = model.predict(obs_input, deterministic=True)

# Convert action to trading signal
if action == 0:
    current_signal = "BUY"
else:
    current_signal = "SELL"
```

```
print(f"\nLatest model signal at {env.data.iloc[last_index]['Date'].date()}:  
↳ {current_signal}")
```


Latest model signal at 2025-04-10: SELL

```
[549]: # --- Plot Performance with Trade Markers ---  
buy_hold_line = [env.initial_balance * (p / data['close'].iloc[0]) for p in  
↳ data['close'].iloc[:len(portfolio_values)]]  
  
date_to_step = {row['Date'].strftime('%Y-%m-%d'): i for i, row in data.iloc[:  
↳ len(portfolio_values)].iterrows()}  
entry_points = []  
exit_points = []  
  
for trade in env.trade_log:  
    entry_step = date_to_step.get(trade['EntryDate'])  
    exit_step = date_to_step.get(trade['ExitDate'])  
    if entry_step is not None and entry_step < len(portfolio_values):  
        entry_points.append((entry_step, portfolio_values[entry_step]))  
    if exit_step is not None and exit_step < len(portfolio_values):  
        exit_points.append((exit_step, portfolio_values[exit_step]))  
  
plt.figure(figsize=(12, 6))  
plt.plot(portfolio_values, label="Recurrent PPO Portfolio", linewidth=2)  
plt.plot(buy_hold_line, label="Buy & Hold Benchmark", linestyle="--",  
↳ color="gray")  
  
if entry_points:  
    entry_steps, entry_vals = zip(*entry_points)  
    plt.scatter(entry_steps, entry_vals, color='green', marker='^', s=80,  
↳ label="Trade Entry")  
if exit_points:  
    exit_steps, exit_vals = zip(*exit_points)  
    plt.scatter(exit_steps, exit_vals, color='red', marker='v', s=80,  
↳ label="Trade Exit")  
  
plt.title("Portfolio Value: Recurrent PPO vs. Buy & Hold")  
plt.xlabel("Time Step")  
plt.ylabel("Portfolio Value")  
plt.legend()  
plt.grid(True)  
plt.tight_layout()  
plt.savefig("recurrent_ppo_performance.png")  
plt.show()
```



```
[550]: import matplotlib.pyplot as plt
```


```
plt.figure(figsize=(10, 5))
plt.plot(callback.timesteps, callback.episode_rewards, marker='o')
plt.title("Episode Reward Over Time")
plt.xlabel("Timesteps")
plt.ylabel("Episode Reward")
plt.grid(True)
plt.tight_layout()
plt.show()
```



```
[551]: import matplotlib.pyplot as plt
import numpy as np
import pandas as pd

rewards = np.array(env.step_rewards)

# Histogram
plt.hist(rewards, bins=50, alpha=0.7, edgecolor='black')
plt.title("Reward Distribution per Step")
plt.xlabel("Reward")
plt.ylabel("Frequency")
plt.grid(True)
plt.show()
```



```
[552]: import numpy as np
import matplotlib.pyplot as plt
import pandas as pd

rewards = np.array(env.step_rewards)
```



```


# --- Raw volatility metrics ---
reward_std = np.std(rewards)
reward_mean = np.mean(rewards)
reward_variance = np.var(rewards)
reward_range = np.max(rewards) - np.min(rewards)

# --- Rolling volatility ---
window = 20
rolling_std = pd.Series(rewards).rolling(window=window).std()

# --- Plot ---
plt.figure(figsize=(10, 4))
plt.plot(rolling_std, label=f"Rolling Std (window={window})")
plt.axhline(reward_std, color='red', linestyle='--', label='Global Std')
plt.title("Reward Volatility (Rolling Std Dev of Step Rewards)")
plt.xlabel("Step")
plt.ylabel("Rolling Std Dev")
plt.legend()
plt.grid(True)
plt.tight_layout()
plt.show()

# --- Output stats ---
print(f"Mean Step Reward: {reward_mean:.4f}")
print(f"Std Dev of Step Rewards: {reward_std:.4f}")
print(f"Variance: {reward_variance:.4f}")
print(f"Range: {reward_range:.4f}")
print(f"Max Rolling Std ({window}): {rolling_std.max():.4f}")
print(f"Min Rolling Std ({window}): {rolling_std.min():.4f}")

```



Mean Step Reward: 0.3679

Std Dev of Step Rewards: 1.4658
Variance: 2.1486
Range: 6.8229
Max Rolling Std (20): 1.9977
Min Rolling Std (20): 0.8796

```
[553]: import pandas as pd
import matplotlib.pyplot as plt

def compute_switch_rate(trade_log: pd.DataFrame, plot: bool = True):
    if 'Position' not in trade_log.columns or trade_log.empty:
        raise ValueError("Trade log must contain a 'Position' column and be
        ↪non-empty.")

    positions = trade_log['Position'].tolist()
    if len(positions) < 2:
        print("Insufficient trades to compute switching behavior.")
        return 0.0

    switches = sum(1 for i in range(1, len(positions)) if positions[i] !=
    ↪positions[i - 1])
    total_transitions = len(positions) - 1
    switch_rate = switches / total_transitions


    print(f"Total Trades: {len(positions)}")
    print(f"Total Transitions: {total_transitions}")
    print(f"Switches: {switches}")
    print(f"Switch Rate: {switch_rate:.4f}")

    if plot:
        plt.bar(['Same', 'Switch'], [total_transitions - switches, switches],
        ↪color=['gray', 'red'])
        plt.title("Trade Direction Transitions")
        plt.ylabel("Count")
        plt.grid(True, axis='y', linestyle='--', alpha=0.6)
        plt.show()

    return switch_rate

df_trades = pd.DataFrame(env.trade_log)
switch_rate = compute_switch_rate(df_trades)
```

Total Trades: 177
Total Transitions: 176
Switches: 135
Switch Rate: 0.7670



```
[554]: import matplotlib.pyplot as plt


total_actions = sum(env.action_counts.values())

actions = list(env.action_counts.keys())
counts = [env.action_counts[a] for a in actions]
labels = ['Long (0)', 'Short (1)']

proportions = [count / total_actions for count in counts]

plt.bar(labels, proportions, color='gray', edgecolor='black')
plt.title("Action Distribution (Policy Bias)")
plt.ylabel("Proportion of Actions")
plt.ylim(0, 1)
plt.grid(True, axis='y')
plt.show()

for label, count, prop in zip(labels, counts, proportions):
    print(f"{label}: Count = {count}, Proportion = {prop:.2%}")
```



Long (0): Count = 214, Proportion = 50.00%

Short (1): Count = 214, Proportion = 50.00%

— Training with seed 49863 — Seed 49863: Reward = -1790.80, Final Balance = 10230.76 Average reward for action Long: 0.4532 Average reward for action Short: -0.2896 Best model updated and saved.

— Training with seed 73352 — Seed 73352: Reward = 11332.71, Final Balance = 22961.61 Average reward for action Long: 0.6296 Average reward for action Short: 0.1092 Best model updated and saved.

— Training with seed 114146 — Seed 114146: Reward = 1341.21, Final Balance = 13066.03 Average reward for action Long: 0.2273 Average reward for action Short: 0.0563

— Training with seed 37234 — Seed 37234: Reward = 4619.64, Final Balance = 16387.20 Average reward for action Long: 0.2038 Average reward for action Short: 0.3025

— Training with seed 52682 — Seed 52682: Reward = 2486.83, Final Balance = 14374.79 Average reward for action Long: -0.3329 Average reward for action Short: -0.1697

— Training with seed 44348 — Seed 44348: Reward = -5730.73, Final Balance = 6355.58 Average reward for action Long: 0.1135 Average reward for action Short: -0.8020

— Training with seed 70193 — Seed 70193: Reward = -8681.13, Final Balance = 3719.05 Average reward for action Long: -0.1616 Average reward for action Short: -0.5804

— Training with seed 88635 — Seed 88635: Reward = 5027.80, Final Balance = 16530.86 Average reward for action Long: 0.3168 Average reward for action Short: -0.1560

— Training with seed 78409 — Seed 78409: Reward = -4271.84, Final Balance = 7909.80 Average reward for action Long: -0.2609 Average reward for action Short: -0.0922

— Training with seed 85452 — Seed 85452: Reward = 10804.68, Final Balance = 22520.33 Average reward for action Long: 0.2810 Average reward for action Short: -0.1286

— Training with seed 116433 — Seed 116433: Reward = -5220.17, Final Balance = 6784.07 Average reward for action Long: 0.3066 Average reward for action Short: -0.4070

— Seed Sweep Complete — Best Seed: 73352 Best Total Reward: 11332.71 Best Final Balance: 22961.61

```
[555]: model.save("META_best_model")
```

1.12 MSFT

```
[556]: from datetime import datetime, timedelta

# Add +1 day to end date
default_end_date = (datetime.now() + timedelta(days=1)).strftime("%Y-%m-%d")
default_start_date = (datetime.now() - timedelta(days=665)).strftime("%Y-%m-%d")
```

```
[557]: # Parameters
ticker = "MSFT"
start_date = default_start_date
end_date = default_end_date
data_filename = "full_data.csv"
signals_filename = "signals_data.csv"

save_historical_data(
    ticker=ticker,
    start_date=start_date,
    end_date=end_date,
    data_filename=data_filename,
    signals_filename=signals_filename
)
```

```
[*****100%*****] 1 of 1 completed
```

```
[557]: 'Saved full_data.csv and signals_data.csv successfully.'
```

```
[558]: import pandas as pd

# Preview full data
df_full = pd.read_csv("full_data.csv", parse_dates=["Date"])
print(df_full.head())
```

```
# Preview signal data
```

```
df_signals = pd.read_csv("signals_data.csv", parse_dates=["Date"])
print(df_signals.head())
```

	Date	close	high	low	open	volume	\
0	2023-06-16	337.707245	346.723835	337.332402	346.575867	46533600	
1	2023-06-20	333.485016	337.460594	331.324587	334.728011	26375400	
2	2023-06-21	329.055695	333.169397	327.585825	331.827747	25117800	
3	2023-06-22	335.122620	335.527087	328.838644	329.608109	23556800	
4	2023-06-23	330.495972	333.396273	328.947195	329.844881	23084700	

	EMA_50	EMA_100	EMA_200	EMA_500	...	RSI_Sell	MCDX_Buy	\
0	337.707245	337.707245	337.707245	337.707245	...	1	0	
1	337.541667	337.623636	337.665233	337.690390	...	1	0	
2	337.208884	337.453974	337.579566	337.655920	...	1	0	
3	337.127070	337.407809	337.555118	337.645807	...	1	0	
4	336.867027	337.270941	337.484878	337.617265	...	1	0	

	MCDX_Sell	DSS_Buy	DSS_Sell	ZeroLag	MACD_Buy	ZeroLag	MACD_Sell	\
0	1	0	1		0		1	
1	1	0	1		0		1	
2	1	0	1		0		1	
3	1	0	1		0		1	
4	1	0	1		0		1	

	Basic	MACD_Buy	Basic	MACD_Sell	OverallTrade
0		0		1	Sell
1		0		1	Sell
2		0		1	Sell
3		0		1	Sell
4		0		1	Sell

[5 rows x 42 columns]

	Date	Signal	Z-Score
0	2023-06-20	Basic MACD Sell	NaN
1	2023-06-20	ZeroLag MACD Sell	NaN
2	2023-06-20	ZLMA Sell	NaN
3	2023-06-26	DSS Sell	NaN
4	2023-06-27	DSS Buy	NaN

```
[559]: from stable_baselines3.common.callbacks import BaseCallback
```

```
class RewardTrackingCallback(BaseCallback):
    def __init__(self, verbose=0):
        super().__init__(verbose)
        self.episode_rewards = []
        self.timesteps = []
```

```

def _on_step(self) -> bool:
    if self.locals.get("done") is not None and any(self.locals["done"]):
        ep_rew = self.locals["rewards"]
        self.episode_rewards.append(sum(ep_rew))
        self.timesteps.append(self.num_timesteps)
    return True

```

```

[ ]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import gymnasium as gym
from gymnasium import spaces
from sb3_contrib import RecurrentPPO
from sb3_contrib.ppo_recurrent.policies import MlpLstmPolicy
from stable_baselines3.common.vec_env import DummyVecEnv
import torch
import random

# --- Fixed Seed ---
SEED = 96147
np.random.seed(SEED)
random.seed(SEED)
torch.manual_seed(SEED)

# --- Load & Clean Data ---
data_path = "full_data.csv"
data = pd.read_csv(data_path, parse_dates=["Date"])

required_cols = [
    'open', 'high', 'low', 'close',
    'zlma', 'ema_value',
    'DSSb', 'DSSsignal',
    'rsi_ma_base',
    'ZeroLagMACD', 'ZeroLagMACD_signal',
    'basicMACD', 'basicMACD_signal',
    'ZScore',
    'ZLMA_Buy', 'ZLMA_Sell',
    'RSI_Buy', 'RSI_Sell',
    'MCDX_Buy', 'MCDX_Sell',
    'DSS_Buy', 'DSS_Sell',
    'ZeroLag MACD_Buy', 'ZeroLag MACD_Sell',
    'Basic MACD_Buy', 'Basic MACD_Sell'
]

missing = [col for col in required_cols if col not in data.columns]
assert not missing, f"Missing required columns: {missing}"

```

```

data = data.dropna(subset=required_cols).reset_index(drop=True)

SWITCH_COST = 1.0
TRANSACTION_COST = 0.001

ACTIVATION = torch.nn.Tanh
# ACTIVATION = partial(nn.LeakyReLU, negative_slope=0.01)

policy_kwargs = dict(
    activation_fn=ACTIVATION
)

# --- Main Trading Environment ---
class TradingEnvRL(gym.Env):
    metadata = {'render_modes': ['human']}

    def __init__(self, data, initial_balance=10000, hold_cost=0.02,
                 volatility_window=14, exploration_steps=500,
                 switch_cost=1.0, reentry_threshold=0.01,
                 dynamic_threshold=False, loss_penalty=0.75, drawdown_penalty=5.
↪0, large_loss_threshold=-2.0):

        super().__init__()
        self.loss_penalty = loss_penalty
        self.drawdown_penalty = drawdown_penalty
        self.large_loss_threshold = large_loss_threshold
        self.data = data.reset_index(drop=True).copy()
        self.n_steps = len(self.data)
        self.initial_balance = initial_balance
        self.hold_cost = hold_cost
        self.volatility_window = volatility_window
        self.exploration_steps = exploration_steps
        self.switch_cost = switch_cost
        self.reentry_threshold = reentry_threshold
        self.dynamic_threshold = dynamic_threshold
        self.step_rewards = []

        self.feature_cols = [
            'open', 'high', 'low', 'close',
            'basicMACD', 'basicMACD_signal',
            'Basic MACD_Buy', 'Basic MACD_Sell'
        ]

        obs_dim = len(self.feature_cols) + 1
        self.observation_space = spaces.Box(low=-np.inf, high=np.inf,
↪shape=(obs_dim,), dtype=np.float32)

```



```

self.action_space = spaces.Discrete(2) # 0=Long, 1=Short

self._compute_volatility_limit()

def _compute_volatility_limit(self):
    returns = self.data['close'].pct_change()
    self.data['volatility'] = returns.rolling(self.volatility_window).std()
    self.data['adaptive_hold'] = (10 / (self.data['volatility'] * 100)).
↳ clip(lower=3, upper=20).fillna(10).astype(int)

def reset(self, seed=None, options=None):
    if seed is not None:
        np.random.seed(seed)
        random.seed(seed)
        torch.manual_seed(seed)

    self.current_step = 0
    self.position = 0
    self.entry_price = 0.0
    self.entry_date = None
    self.hold_counter = 0
    self.switch_count = 0
    self.balance = self.initial_balance
    self.cumulative_pnl = 0.0
    self.trade_log = []
    self.action_counts = {0: 0, 1: 0}
    self.reward_tracker = {0: [], 1: []}
    self.consecutive_losses = 0
    self.equity_curve = [self.initial_balance]
    self.step_rewards = []
    return self._get_obs(), {}

def _get_obs(self):
    row = self.data.iloc[self.current_step]
    features = row[self.feature_cols].values.astype(np.float32)
    pos_feature = np.array([self.position], dtype=np.float32)
    return np.concatenate([features, pos_feature])

def _force_close(self):
    row = self.data.iloc[self.current_step]
    current_price = float(row['close'])
    current_date = row['Date'].strftime("%Y-%m-%d")

    if self.position == 0:
        return 0.0

# --- Core Return Logic ---

```

```

        trade_pct = ((current_price / self.entry_price - 1) * 100) if self.
↪position == 1 else ((self.entry_price / current_price - 1) * 100)
        pos_str = 'Long' if self.position == 1 else 'Short'
        gross_return = trade_pct / 100
        transaction_cost = TRANSACTION_COST * current_price

        old_balance = self.balance
        self.balance -= transaction_cost
        self.balance *= (1 + gross_return)
        net_profit = self.balance - old_balance

        reward = net_profit

        # --- Track Trade History ---
        self.cumulative_pnl += trade_pct
        compounded_pnl = (self.trade_log[-1]['CompoundedFactor'] * (1 +
↪gross_return)) if self.trade_log else (1 + gross_return)
        compounded_pnl_pct = (compounded_pnl - 1) * 100

        self.trade_log.append({
            'EntryDate': self.entry_date,
            'ExitDate': current_date,
            'Position': pos_str,
            'EntryPrice': round(self.entry_price, 4),
            'ExitPrice': round(current_price, 4),
            'PnL%': round(trade_pct, 2),
            'CumulativePnL%': round(self.cumulative_pnl, 2),
            'CompoundedFactor': compounded_pnl,
            'CompoundedPnL%': round(compounded_pnl_pct, 2)
        })

        # --- Penalty for Large Loss ---
        if trade_pct < -2.0:
            reward += trade_pct * 2 # Stronger penalty for large loss

        # --- Track & Penalize Consecutive Losses ---
        if trade_pct < 0:
            self.consecutive_losses += 1
            reward -= self.consecutive_losses * 0.75 # Growing penalty
        else:
            self.consecutive_losses = 0

        # --- Optional: Drawdown Penalty ---
        self.equity_curve.append(self.balance)
        max_balance = max(self.equity_curve)
        if max_balance > 0:
            drawdown = (max_balance - self.balance) / max_balance

```

```

        reward -= 5.0 * drawdown # Penalize deeper drawdown

# --- Reset State ---
self.position = 0
self.entry_price = 0.0
self.entry_date = None
self.hold_counter = 0

return reward

def step(self, action):
    if isinstance(action, np.ndarray):
        action = int(action.item())

    row = self.data.iloc[self.current_step]
    current_price = float(row['close'])
    current_date = row['Date'].strftime("%Y-%m-%d")

    if self.current_step < self.exploration_steps:
        action = self.action_space.sample()

    self.action_counts[action] += 1
    reward = 0.0
    target_position = 1 if action == 0 else -1

    # === Holding Same Position ===
    if self.position == target_position:
        self.hold_counter += 1
        price_change = (current_price - self.entry_price) / self.entry_price
        step_return = price_change if self.position == 1 else -price_change

        step_reward = np.sign(step_return) * np.sqrt(abs(step_return)) * 10
        step_reward = np.clip(step_reward, -50, 50)
        step_reward -= self.hold_cost
        reward += step_reward
        self.reward_tracker[action].append(step_reward)
        self.step_rewards.append(reward)

    # === New Position from Flat ===
    elif self.position == 0:
        self.position = target_position
        self.entry_price = current_price
        self.entry_date = current_date
        self.hold_counter = 1
        reward -= TRANSACTION_COST * current_price

    # === Switch Position ===

```

```

else:
    hold_penalty = max(0, 3 - self.hold_counter) * 5.0
    switch_penalty = self.switch_cost + hold_penalty
    reward += self._force_close()
    reward -= switch_penalty

    # === Dynamic Re-entry Threshold ===
    if self.current_step > 0:
        prev_close = self.data.iloc[self.current_step - 1]['close']
        recent_return = abs((current_price - prev_close) / prev_close)

        # Volatility-aware reentry threshold
        recent_volatility = row['volatility']
        reentry_threshold = 0.5 * recent_volatility
    else:
        recent_return = 0
        reentry_threshold = 0.01 # Fallback

    if recent_return > reentry_threshold:
        self.position = target_position
        self.entry_price = current_price
        self.entry_date = current_date
        self.hold_counter = 1
        reward -= TRANSACTION_COST * current_price

self.current_step += 1
terminated = self.current_step >= self.n_steps - 1

if terminated and self.position != 0:
    reward += self._force_close()

obs = self._get_obs() if not terminated else np.zeros(self.
↪observation_space.shape, dtype=np.float32)
return obs, reward, terminated, False, {}

def render(self):
    print(f"Step: {self.current_step}, Position: {self.position}, Balance: ↵
↪{self.balance:.2f}")

def save_trade_log(self, filename="trade_log.csv"):
    df = pd.DataFrame(self.trade_log)
    if "CompoundedFactor" in df.columns:
        df = df.drop(columns=["CompoundedFactor"])
    df.to_csv(filename, index=False)

# --- Training ---
env = TradingEnvRL(data, initial_balance=10000)

```

```

vec_env = DummyVecEnv([lambda: env])
vec_env.seed(SEED)

model = RecurrentPP0(
    policy=MlpLstmPolicy,
    env=vec_env,
    verbose=1,
    n_steps=64,
    batch_size=32,
    learning_rate=0.01,
    gamma=0.99,
    ent_coef=0.01,
    seed=SEED,
    policy_kwargs=policy_kwargs
)

callback = RewardTrackingCallback()

model.learn(total_timesteps=3000, callback=callback)
print(f"\nTraining Complete")

```

```

[561]: # --- Evaluation ---
obs, _ = env.reset()
state = None
done = False
total_reward = 0
portfolio_values = []
final_action = None

while not done:
    action, state = model.predict(obs, state=state, deterministic=True)
    final_action = action
    obs, reward, done, _, _ = env.step(action)
    total_reward += reward

    current_index = min(env.current_step, len(env.data) - 1)
    current_price = env.data.loc[current_index, 'close']
    unrealized = (
        (current_price - env.entry_price) if env.position == 1 else
        (env.entry_price - current_price) if env.position == -1 else
        0.0
    )
    mtm_equity = env.balance + unrealized
    portfolio_values.append(mtm_equity)

# --- Save trade log ---
env.save_trade_log("trade_log_recurrent.csv")

```

```

# --- Final Model Signal ---
signal_str = "BUY" if final_action == 0 else "SELL"
latest_date = env.data['Date'].iloc[env.current_step - 1].strftime("%Y-%m-%d")
print(f"\nLatest model signal at {latest_date}: {signal_str}")

# --- Summary Output ---
print(f"Total Reward: {total_reward:.2f}")
print(f"Final Balance: {env.balance:.2f}")
print("Trade Log (Recurrent PPO):")
print(pd.DataFrame(env.trade_log))
print("Action counts:", env.action_counts)
print("Average reward per action:")
for k, v in env.reward_tracker.items():
    mean_r = np.mean(v) if v else 0
    print(f"Action {k} ({'Long' if k==0 else 'Short'}): {mean_r:.4f}")

```

Latest model signal at 2025-04-09: SELL

Total Reward: 7669.29

Final Balance: 19214.39

Trade Log (Recurrent PPO):

	EntryDate	ExitDate	Position	EntryPrice	ExitPrice	PnL%	\
0	2023-07-27	2023-08-01	Short	326.2540	331.7982	-1.67	
1	2023-08-02	2023-08-04	Short	323.0775	323.3537	-0.09	
2	2023-08-07	2023-08-08	Long	325.6522	321.6471	-1.23	
3	2023-08-09	2023-08-11	Long	317.8787	316.6751	-0.38	
4	2023-08-14	2023-08-15	Short	319.6642	317.5137	0.68	
..	
173	2025-03-21	2025-03-24	Short	391.2600	393.0800	-0.46	
174	2025-03-25	2025-04-02	Short	395.1600	382.1400	3.41	
175	2025-04-03	2025-04-08	Short	373.1100	354.5600	5.23	
176	2025-04-08	2025-04-09	Long	354.5600	390.4900	10.13	
177	2025-04-09	2025-04-10	Short	390.4900	381.3500	2.40	

	CumulativePnL%	CompoundedFactor	CompoundedPnL%
0	-1.67	0.983291	-1.67
1	-1.76	0.982451	-1.75
2	-2.99	0.970368	-2.96
3	-3.36	0.966694	-3.33
4	-2.69	0.973241	-2.68
..
173	48.57	1.575948	57.59
174	51.97	1.629643	62.96
175	57.20	1.714903	71.49
176	67.34	1.888686	88.87
177	69.74	1.933953	93.40

```
[178 rows x 9 columns]
Action counts: {0: 217, 1: 211}
Average reward per action:
Action 0 (Long): 0.2989
Action 1 (Short): 0.1202
```

```
[562]: # --- Get Latest Signal from Model ---
last_index = len(env.data) - 1
env.current_step = last_index
obs = env._get_obs()

# Add batch dimension and run prediction
obs_input = obs[np.newaxis, :]
action, _ = model.predict(obs_input, deterministic=True)

# Convert action to trading signal
if action == 0:
    current_signal = "BUY"
else:
    current_signal = "SELL"

print(f"\nLatest model signal at {env.data.iloc[last_index]['Date'].date()}:
↳ {current_signal}")
```

Latest model signal at 2025-04-10: SELL

```
[563]: # --- Plot Performance with Trade Markers ---
buy_hold_line = [env.initial_balance * (p / data['close'].iloc[0]) for p in
↳ data['close'].iloc[:len(portfolio_values)]]

date_to_step = {row['Date'].strftime('%Y-%m-%d'): i for i, row in data.iloc[:
↳ len(portfolio_values)].iterrows()}
entry_points = []
exit_points = []

for trade in env.trade_log:
    entry_step = date_to_step.get(trade['EntryDate'])
    exit_step = date_to_step.get(trade['ExitDate'])
    if entry_step is not None and entry_step < len(portfolio_values):
        entry_points.append((entry_step, portfolio_values[entry_step]))
    if exit_step is not None and exit_step < len(portfolio_values):
        exit_points.append((exit_step, portfolio_values[exit_step]))


plt.figure(figsize=(12, 6))
plt.plot(portfolio_values, label="Recurrent PPO Portfolio", linewidth=2)
plt.plot(buy_hold_line, label="Buy & Hold Benchmark", linestyle="--",
↳ color="gray")
```

```

if entry_points:
    entry_steps, entry_vals = zip(*entry_points)
    plt.scatter(entry_steps, entry_vals, color='green', marker='^', s=80,
    ↪label="Trade Entry")
if exit_points:
    exit_steps, exit_vals = zip(*exit_points)
    plt.scatter(exit_steps, exit_vals, color='red', marker='v', s=80,
    ↪label="Trade Exit")

plt.title("Portfolio Value: Recurrent PPO vs. Buy & Hold")
plt.xlabel("Time Step")
plt.ylabel("Portfolio Value")
plt.legend()
plt.grid(True)
plt.tight_layout()
plt.savefig("recurrent_ppo_performance.png")
plt.show()

```



```


[564]: import matplotlib.pyplot as plt

plt.figure(figsize=(10, 5))
plt.plot(callback.timesteps, callback.episode_rewards, marker='o')
plt.title("Episode Reward Over Time")
plt.xlabel("Timesteps")
plt.ylabel("Episode Reward")
plt.grid(True)
plt.tight_layout()

```




```
plt.show()
```



```
[565]: import matplotlib.pyplot as plt
import numpy as np
import pandas as pd

rewards = np.array(env.step_rewards)

# Histogram
plt.hist(rewards, bins=50, alpha=0.7, edgecolor='black')
plt.title("Reward Distribution per Step")
plt.xlabel("Reward")
plt.ylabel("Frequency")
plt.grid(True)
plt.show()
```



```
[566]: import numpy as np
import matplotlib.pyplot as plt
import pandas as pd

rewards = np.array(env.step_rewards)

# --- Raw volatility metrics ---
reward_std = np.std(rewards)
reward_mean = np.mean(rewards)
reward_variance = np.var(rewards)
reward_range = np.max(rewards) - np.min(rewards)

# --- Rolling volatility ---
window = 20
rolling_std = pd.Series(rewards).rolling(window=window).std()

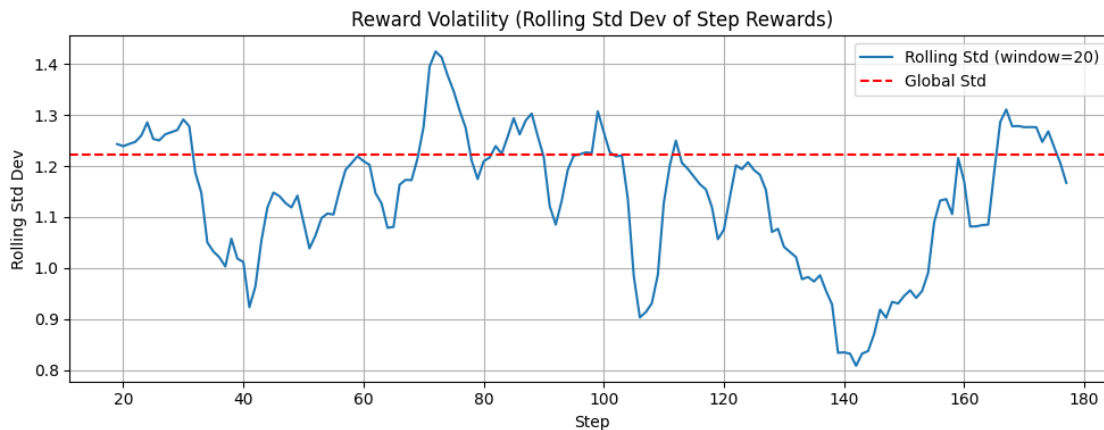
# --- Plot ---
plt.figure(figsize=(10, 4))
plt.plot(rolling_std, label=f"Rolling Std (window={window})")
```

```

plt.axhline(reward_std, color='red', linestyle='--', label='Global Std')
plt.title("Reward Volatility (Rolling Std Dev of Step Rewards)")
plt.xlabel("Step")
plt.ylabel("Rolling Std Dev")
plt.legend()
plt.grid(True)
plt.tight_layout()
plt.show()

# --- Output stats ---
print(f"Mean Step Reward: {reward_mean:.4f}")
print(f"Std Dev of Step Rewards: {reward_std:.4f}")
print(f"Variance: {reward_variance:.4f}")
print(f"Range: {reward_range:.4f}")
print(f"Max Rolling Std ({window}): {rolling_std.max():.4f}")
print(f"Min Rolling Std ({window}): {rolling_std.min():.4f}")

```



```

Mean Step Reward: 0.2126
Std Dev of Step Rewards: 1.2217
Variance: 1.4926
Range: 5.1155
Max Rolling Std (20): 1.4249
Min Rolling Std (20): 0.8083

```

```

[567]: import pandas as pd
import matplotlib.pyplot as plt

def compute_switch_rate(trade_log: pd.DataFrame, plot: bool = True):
    if 'Position' not in trade_log.columns or trade_log.empty:
        raise ValueError("Trade log must contain a 'Position' column and be
        non-empty.")

```

```

positions = trade_log['Position'].tolist()
if len(positions) < 2:
    print("Insufficient trades to compute switching behavior.")
    return 0.0

switches = sum(1 for i in range(1, len(positions)) if positions[i] !=
↳positions[i - 1])
total_transitions = len(positions) - 1
switch_rate = switches / total_transitions

print(f"Total Trades: {len(positions)}")
print(f"Total Transitions: {total_transitions}")
print(f"Switches: {switches}")
print(f"Switch Rate: {switch_rate:.4f}")

if plot:
    plt.bar(['Same', 'Switch'], [total_transitions - switches, switches],
↳color=['gray', 'red'])
    plt.title("Trade Direction Transitions")
    plt.ylabel("Count")
    plt.grid(True, axis='y', linestyle='--', alpha=0.6)
    plt.show()

return switch_rate

df_trades = pd.DataFrame(env.trade_log)
switch_rate = compute_switch_rate(df_trades)

```

```

Total Trades: 178
Total Transitions: 177
Switches: 146
Switch Rate: 0.8249

```



```
[568]: import matplotlib.pyplot as plt


total_actions = sum(env.action_counts.values())

actions = list(env.action_counts.keys())
counts = [env.action_counts[a] for a in actions]
labels = ['Long (0)', 'Short (1)']

proportions = [count / total_actions for count in counts]

plt.bar(labels, proportions, color='gray', edgecolor='black')
plt.title("Action Distribution (Policy Bias)")
plt.ylabel("Proportion of Actions")
plt.ylim(0, 1)
plt.grid(True, axis='y')
plt.show()

for label, count, prop in zip(labels, counts, proportions):
    print(f"{label}: Count = {count}, Proportion = {prop:.2%}")
```



Long (0): Count = 217, Proportion = 50.70%

Short (1): Count = 211, Proportion = 49.30%

— Training with seed 50347 — Seed 50347: Reward = 3326.04, Final Balance = 14929.36 Average reward for action Long: 0.2001 Average reward for action Short: -0.1709 Best model updated and saved.

— Training with seed 115278 — Seed 115278: Reward = 259.69, Final Balance = 11949.41 Average reward for action Long: 0.1923 Average reward for action Short: -0.1650

— Training with seed 131667 — Seed 131667: Reward = 2915.38, Final Balance = 14837.79 Average reward for action Long: 0.1218 Average reward for action Short: -0.4758

— Training with seed 127306 — Seed 127306: Reward = 3358.08, Final Balance = 15112.39 Average reward for action Long: 0.4093 Average reward for action Short: -0.2912 Best model updated and saved.

— Training with seed 76487 — Seed 76487: Reward = 3018.97, Final Balance = 14742.02 Average reward for action Long: 0.2445 Average reward for action Short: -0.6318

— Training with seed 92365 — Seed 92365: Reward = -4420.60, Final Balance = 7503.18 Average reward for action Long: -0.0679 Average reward for action Short: -0.4151

— Training with seed 58262 — Seed 58262: Reward = -1632.51, Final Balance = 10381.70 Average reward for action Long: 0.1275 Average reward for action Short: -0.5116

— Training with seed 76099 — Seed 76099: Reward = 228.53, Final Balance = 12044.49 Average reward for action Long: 0.3265 Average reward for action Short: 0.0292

— Training with seed 96147 — Seed 96147: Reward = 8151.74, Final Balance = 19766.26 Average reward for action Long: 0.5826 Average reward for action Short: 0.1154 Best model updated and saved.

— Training with seed 107905 — Seed 107905: Reward = 3193.40, Final Balance = 14883.34 Average reward for action Long: 0.8358 Average reward for action Short: -0.4058

— Training with seed 113756 — Seed 113756: Reward = -4101.27, Final Balance = 7844.85 Average reward for action Long: 0.0314 Average reward for action Short: -0.2609

— Seed Sweep Complete — Best Seed: 96147 Best Total Reward: 8151.74 Best Final Balance: 19766.26

```
[569]: model.save("MSFT_best_model")
```

1.13 NVDA

```
[486]: from datetime import datetime, timedelta

# Add +1 day to end date
default_end_date = (datetime.now() + timedelta(days=1)).strftime("%Y-%m-%d")
default_start_date = (datetime.now() - timedelta(days=665)).strftime("%Y-%m-%d")
```

```
[487]: # Parameters
ticker = "NVDA"
start_date = default_start_date
end_date = default_end_date
data_filename = "full_data.csv"
signals_filename = "signals_data.csv"

save_historical_data(
    ticker=ticker,
    start_date=start_date,
    end_date=end_date,
    data_filename=data_filename,
    signals_filename=signals_filename
)
```

[*****100%*****] 1 of 1 completed

```
[487]: 'Saved full_data.csv and signals_data.csv successfully.'
```

```
[488]: import pandas as pd

# Preview full data
df_full = pd.read_csv("full_data.csv", parse_dates=["Date"])
print(df_full.head())
```

```
# Preview signal data
df_signals = pd.read_csv("signals_data.csv", parse_dates=["Date"])
print(df_signals.head())
```

	Date	close	high	low	open	volume	\
0	2023-06-16	42.668709	43.697147	42.637724	43.426295	655709000	
1	2023-06-20	43.784096	43.965999	42.650716	42.974540	451153000	
2	2023-06-21	43.021511	43.591204	42.057041	43.477264	551603000	
3	2023-06-22	43.001526	43.402304	42.210957	42.229944	417737000	
4	2023-06-23	42.185974	42.785645	41.992080	42.440836	358140000	

	EMA_50	EMA_100	EMA_200	EMA_500	...	RSI_Sell	MCDX_Buy	\
0	42.668709	42.668709	42.668709	42.668709	...	1	0	
1	42.712449	42.690796	42.679807	42.673161	...	0	0	
2	42.724570	42.697344	42.683207	42.674552	...	0	0	
3	42.735431	42.703368	42.686375	42.675857	...	0	0	
4	42.713883	42.693122	42.681395	42.673902	...	1	0	

	MCDX_Sell	DSS_Buy	DSS_Sell	ZeroLag	MACD_Buy	ZeroLag	MACD_Sell	\
0	1	0	1		0		1	
1	1	0	1		1		0	
2	1	0	1		1		0	
3	1	0	1		1		0	
4	1	0	1		0		1	

	Basic	MACD_Buy	Basic	MACD_Sell	OverallTrade
0		0		1	Sell
1		1		0	Buy
2		1		0	Buy
3		1		0	Buy
4		1		0	Sell

[5 rows x 42 columns]

	Date	Signal	Z-Score
0	2023-06-20	Basic MACD Buy	NaN
1	2023-06-20	RSI Buy	NaN
2	2023-06-20	ZLMA Buy	NaN
3	2023-06-20	ZeroLag MACD Buy	NaN
4	2023-06-23	RSI Sell	NaN

```
[489]: from stable_baselines3.common.callbacks import BaseCallback
```

```
class RewardTrackingCallback(BaseCallback):
    def __init__(self, verbose=0):
        super().__init__(verbose)
        self.episode_rewards = []
```



```

        self.timesteps = []

    def _on_step(self) -> bool:
        if self.locals.get("dones") is not None and any(self.locals["dones"]):
            ep_rew = self.locals["rewards"]
            self.episode_rewards.append(sum(ep_rew))
            self.timesteps.append(self.num_timesteps)
        return True

```

```

[ ]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import gymnasium as gym
from gymnasium import spaces
from sb3_contrib import RecurrentPPO
from sb3_contrib.ppo_recurrent.policies import MlpLstmPolicy
from stable_baselines3.common.vec_env import DummyVecEnv
import torch
import random

# --- Fixed Seed ---
SEED = 34500
np.random.seed(SEED)
random.seed(SEED)
torch.manual_seed(SEED)

# --- Load & Clean Data ---
data_path = "full_data.csv"
data = pd.read_csv(data_path, parse_dates=["Date"])

required_cols = [
    'open', 'high', 'low', 'close',
    'zlma', 'ema_value',
    'DSSb', 'DSSsignal',
    'rsi_ma_base',
    'ZeroLagMACD', 'ZeroLagMACD_signal',
    'basicMACD', 'basicMACD_signal',
    'ZScore',
    'ZLMA_Buy', 'ZLMA_Sell',
    'RSI_Buy', 'RSI_Sell',
    'MCDX_Buy', 'MCDX_Sell',
    'DSS_Buy', 'DSS_Sell',
    'ZeroLag MACD_Buy', 'ZeroLag MACD_Sell',
    'Basic MACD_Buy', 'Basic MACD_Sell'
]

missing = [col for col in required_cols if col not in data.columns]

```

```

assert not missing, f"Missing required columns: {missing}"

data = data.dropna(subset=required_cols).reset_index(drop=True)

SWITCH_COST = 1.0
TRANSACTION_COST = 0.001

ACTIVATION = torch.nn.Tanh
# ACTIVATION = partial(nn.LeakyReLU, negative_slope=0.01)

policy_kwargs = dict(
    activation_fn=ACTIVATION
)

# --- Main Trading Environment ---
class TradingEnvRL(gym.Env):
    metadata = {'render_modes': ['human']}

    def __init__(self, data, initial_balance=10000, hold_cost=0.02,
                  volatility_window=14, exploration_steps=500,
                  switch_cost=1.0, reentry_threshold=0.01,
                  dynamic_threshold=False, loss_penalty=0.75, drawdown_penalty=5.
↪0, large_loss_threshold=-2.0):

        super().__init__()
        self.loss_penalty = loss_penalty
        self.drawdown_penalty = drawdown_penalty
        self.large_loss_threshold = large_loss_threshold
        self.data = data.reset_index(drop=True).copy()
        self.n_steps = len(self.data)
        self.initial_balance = initial_balance
        self.hold_cost = hold_cost
        self.volatility_window = volatility_window
        self.exploration_steps = exploration_steps
        self.switch_cost = switch_cost
        self.reentry_threshold = reentry_threshold
        self.dynamic_threshold = dynamic_threshold
        self.step_rewards = []

        self.feature_cols = [
            'open', 'high', 'low', 'close',
            'basicMACD', 'basicMACD_signal',
            'Basic MACD_Buy', 'Basic MACD_Sell'
        ]

        obs_dim = len(self.feature_cols) + 1

```

```

        self.observation_space = spaces.Box(low=-np.inf, high=np.inf,
↪shape=(obs_dim,), dtype=np.float32)
        self.action_space = spaces.Discrete(2) # 0=Long, 1=Short

        self._compute_volatility_limit()

    def _compute_volatility_limit(self):
        returns = self.data['close'].pct_change()
        self.data['volatility'] = returns.rolling(self.volatility_window).std()
        self.data['adaptive_hold'] = (10 / (self.data['volatility'] * 100)).
↪clip(lower=3, upper=20).fillna(10).astype(int)

    def reset(self, seed=None, options=None):
        if seed is not None:
            np.random.seed(seed)
            random.seed(seed)
            torch.manual_seed(seed)

        self.current_step = 0
        self.position = 0
        self.entry_price = 0.0
        self.entry_date = None
        self.hold_counter = 0
        self.switch_count = 0
        self.balance = self.initial_balance
        self.cumulative_pnl = 0.0
        self.trade_log = []
        self.action_counts = {0: 0, 1: 0}
        self.reward_tracker = {0: [], 1: []}
        self.consecutive_losses = 0
        self.equity_curve = [self.initial_balance]
        self.step_rewards = []
        return self._get_obs(), {}

    def _get_obs(self):
        row = self.data.iloc[self.current_step]
        features = row[self.feature_cols].values.astype(np.float32)
        pos_feature = np.array([self.position], dtype=np.float32)
        return np.concatenate([features, pos_feature])

    def _force_close(self):
        row = self.data.iloc[self.current_step]
        current_price = float(row['close'])
        current_date = row['Date'].strftime("%Y-%m-%d")

        if self.position == 0:
            return 0.0

```

```

# --- Core Return Logic ---
trade_pct = ((current_price / self.entry_price - 1) * 100) if self.
↪position == 1 else ((self.entry_price / current_price - 1) * 100)
pos_str = 'Long' if self.position == 1 else 'Short'
gross_return = trade_pct / 100
transaction_cost = TRANSACTION_COST * current_price

old_balance = self.balance
self.balance -= transaction_cost
self.balance *= (1 + gross_return)
net_profit = self.balance - old_balance

reward = net_profit

# --- Track Trade History ---
self.cumulative_pnl += trade_pct
compounded_pnl = (self.trade_log[-1]['CompoundedFactor'] * (1 + ↪
↪gross_return)) if self.trade_log else (1 + gross_return)
compounded_pnl_pct = (compounded_pnl - 1) * 100

self.trade_log.append({
    'EntryDate': self.entry_date,
    'ExitDate': current_date,
    'Position': pos_str,
    'EntryPrice': round(self.entry_price, 4),
    'ExitPrice': round(current_price, 4),
    'PnL%': round(trade_pct, 2),
    'CumulativePnL%': round(self.cumulative_pnl, 2),
    'CompoundedFactor': compounded_pnl,
    'CompoundedPnL%': round(compounded_pnl_pct, 2)
})

# --- Penalty for Large Loss ---
if trade_pct < -2.0:
    reward += trade_pct * 2 # Stronger penalty for large loss

# --- Track & Penalize Consecutive Losses ---
if trade_pct < 0:
    self.consecutive_losses += 1
    reward -= self.consecutive_losses * 0.75 # Growing penalty
else:
    self.consecutive_losses = 0

# --- Optional: Drawdown Penalty ---
self.equity_curve.append(self.balance)
max_balance = max(self.equity_curve)

```

```

    if max_balance > 0:
        drawdown = (max_balance - self.balance) / max_balance
        reward -= 5.0 * drawdown # Penalize deeper drawdown

    # --- Reset State ---
    self.position = 0
    self.entry_price = 0.0
    self.entry_date = None
    self.hold_counter = 0

    return reward

def step(self, action):
    if isinstance(action, np.ndarray):
        action = int(action.item())

    row = self.data.iloc[self.current_step]
    current_price = float(row['close'])
    current_date = row['Date'].strftime("%Y-%m-%d")

    if self.current_step < self.exploration_steps:
        action = self.action_space.sample()

    self.action_counts[action] += 1
    reward = 0.0
    target_position = 1 if action == 0 else -1

    # === Holding Same Position ===
    if self.position == target_position:
        self.hold_counter += 1
        price_change = (current_price - self.entry_price) / self.entry_price
        step_return = price_change if self.position == 1 else -price_change

        step_reward = np.sign(step_return) * np.sqrt(abs(step_return)) * 10
        step_reward = np.clip(step_reward, -50, 50)
        step_reward -= self.hold_cost
        reward += step_reward
        self.reward_tracker[action].append(step_reward)
        self.step_rewards.append(reward)

    # === New Position from Flat ===
    elif self.position == 0:
        self.position = target_position
        self.entry_price = current_price
        self.entry_date = current_date
        self.hold_counter = 1
        reward -= TRANSACTION_COST * current_price

```

```

# === Switch Position ===
else:
    hold_penalty = max(0, 3 - self.hold_counter) * 5.0
    switch_penalty = self.switch_cost + hold_penalty
    reward += self._force_close()
    reward -= switch_penalty

# === Dynamic Re-entry Threshold ===
if self.current_step > 0:
    prev_close = self.data.iloc[self.current_step - 1]['close']
    recent_return = abs((current_price - prev_close) / prev_close)

    # Volatility-aware reentry threshold
    recent_volatility = row['volatility']
    reentry_threshold = 0.5 * recent_volatility
else:
    recent_return = 0
    reentry_threshold = 0.01 # Fallback

if recent_return > reentry_threshold:
    self.position = target_position
    self.entry_price = current_price
    self.entry_date = current_date
    self.hold_counter = 1
    reward -= TRANSACTION_COST * current_price

self.current_step += 1
terminated = self.current_step >= self.n_steps - 1

if terminated and self.position != 0:
    reward += self._force_close()

obs = self._get_obs() if not terminated else np.zeros(self.
↪observation_space.shape, dtype=np.float32)
return obs, reward, terminated, False, {}

def render(self):
    print(f"Step: {self.current_step}, Position: {self.position}, Balance:␣
↪{self.balance:.2f}")

def save_trade_log(self, filename="trade_log.csv"):
    df = pd.DataFrame(self.trade_log)
    if "CompoundedFactor" in df.columns:
        df = df.drop(columns=["CompoundedFactor"])
    df.to_csv(filename, index=False)

```

```

# --- Training ---
env = TradingEnvRL(data, initial_balance=10000)
vec_env = DummyVecEnv([lambda: env])
vec_env.seed(SEED)

model = RecurrentPPO(
    policy=MlpLstmPolicy,
    env=vec_env,
    verbose=1,
    n_steps=64,
    batch_size=32,
    learning_rate=0.01,
    gamma=0.99,
    ent_coef=0.01,
    seed=SEED,
    policy_kwargs=policy_kwargs
)

callback = RewardTrackingCallback()

model.learn(total_timesteps=3000, callback=callback)
print(f"\nTraining Complete")

```

```

[491]: # --- Evaluation ---
obs, _ = env.reset()
state = None
done = False
total_reward = 0
portfolio_values = []
final_action = None

while not done:
    action, state = model.predict(obs, state=state, deterministic=True)
    final_action = action
    obs, reward, done, _, _ = env.step(action)
    total_reward += reward

    current_index = min(env.current_step, len(env.data) - 1)
    current_price = env.data.loc[current_index, 'close']
    unrealized = (
        (current_price - env.entry_price) if env.position == 1 else
        (env.entry_price - current_price) if env.position == -1 else
        0.0
    )
    mtm_equity = env.balance + unrealized
    portfolio_values.append(mtm_equity)

```

```

# --- Save trade log ---
env.save_trade_log("trade_log_recurrent.csv")

# --- Final Model Signal ---
signal_str = "BUY" if final_action == 0 else "SELL"
latest_date = env.data['Date'].iloc[env.current_step - 1].strftime("%Y-%m-%d")
print(f"\nLatest model signal at {latest_date}: {signal_str}")

# --- Summary Output ---
print(f"Total Reward: {total_reward:.2f}")
print(f"Final Balance: {env.balance:.2f}")
print("Trade Log (Recurrent PP0):")
print(pd.DataFrame(env.trade_log))
print("Action counts:", env.action_counts)
print("Average reward per action:")
for k, v in env.reward_tracker.items():
    mean_r = np.mean(v) if v else 0
    print(f"Action {k} ({'Long' if k==0 else 'Short'}): {mean_r:.4f}")

```

Latest model signal at 2025-04-09: BUY

Total Reward: 26026.52

Final Balance: 37747.27

Trade Log (Recurrent PP0):

	EntryDate	ExitDate	Position	EntryPrice	ExitPrice	PnL%	\
0	2023-07-27	2023-07-28	Long	45.8750	46.7245	1.85	
1	2023-07-31	2023-08-01	Short	46.7035	46.4816	0.48	
2	2023-08-02	2023-08-07	Short	44.2449	45.3922	-2.53	
3	2023-08-08	2023-08-09	Short	44.6396	42.5308	4.96	
4	2023-08-10	2023-08-11	Long	42.3649	40.8327	-3.62	
..	
170	2025-03-24	2025-03-25	Long	121.4100	120.6900	-0.59	
171	2025-03-26	2025-04-04	Short	113.7600	94.3100	20.62	
172	2025-04-04	2025-04-07	Long	94.3100	97.6400	3.53	
173	2025-04-07	2025-04-08	Short	97.6400	96.3000	1.39	
174	2025-04-09	2025-04-10	Long	114.3300	107.5700	-5.91	

	CumulativePnL%	CompoundedFactor	CompoundedPnL%
0	1.85	1.018518	1.85
1	2.33	1.023380	2.34
2	-0.20	0.997512	-0.25
3	4.76	1.046973	4.70
4	1.14	1.009108	0.91
..
170	130.81	3.171008	217.10
171	151.43	3.824980	282.50
172	154.96	3.960037	296.00
173	156.35	4.015140	301.51

174 150.44 3.777736 277.77

[175 rows x 9 columns]
Action counts: {0: 213, 1: 215}
Average reward per action:
Action 0 (Long): 0.4125
Action 1 (Short): -0.1004

```
[492]: # --- Get Latest Signal from Model ---
last_index = len(env.data) - 1
env.current_step = last_index
obs = env._get_obs()

# Add batch dimension and run prediction
obs_input = obs[np.newaxis, :]
action, _ = model.predict(obs_input, deterministic=True)

# Convert action to trading signal
# 0 = Long => BUY, 1 = Short => SELL
if action == 0:
    current_signal = "BUY"
else:
    current_signal = "SELL"

print(f"\nLatest model signal at {env.data.iloc[last_index]['Date'].date()}:
↳{current_signal}")
```

Latest model signal at 2025-04-10: BUY

```
[493]: # --- Plot Performance with Trade Markers ---
buy_hold_line = [env.initial_balance * (p / data['close'].iloc[0]) for p in
↳data['close'].iloc[:len(portfolio_values)]]

date_to_step = {row['Date'].strftime('%Y-%m-%d'): i for i, row in data.iloc[:
↳len(portfolio_values)].iterrows()}
entry_points = []
exit_points = []

for trade in env.trade_log:
    entry_step = date_to_step.get(trade['EntryDate'])
    exit_step = date_to_step.get(trade['ExitDate'])
    if entry_step is not None and entry_step < len(portfolio_values):
        entry_points.append((entry_step, portfolio_values[entry_step]))
    if exit_step is not None and exit_step < len(portfolio_values):
        exit_points.append((exit_step, portfolio_values[exit_step]))

plt.figure(figsize=(12, 6))
```


```

plt.plot(portfolio_values, label="Recurrent PPO Portfolio", linewidth=2)
plt.plot(buy_hold_line, label="Buy & Hold Benchmark", linestyle="--",
        color="gray")

if entry_points:
    entry_steps, entry_vals = zip(*entry_points)
    plt.scatter(entry_steps, entry_vals, color='green', marker='^', s=80,
        label="Trade Entry")
if exit_points:
    exit_steps, exit_vals = zip(*exit_points)
    plt.scatter(exit_steps, exit_vals, color='red', marker='v', s=80,
        label="Trade Exit")

plt.title("Portfolio Value: Recurrent PPO vs. Buy & Hold")
plt.xlabel("Time Step")
plt.ylabel("Portfolio Value")
plt.legend()
plt.grid(True)
plt.tight_layout()
plt.savefig("recurrent_ppo_performance.png")
plt.show()

```




```

[494]: import matplotlib.pyplot as plt

plt.figure(figsize=(10, 5))
plt.plot(callback.timesteps, callback.episode_rewards, marker='o')
plt.title("Episode Reward Over Time")
plt.xlabel("Timesteps")

```

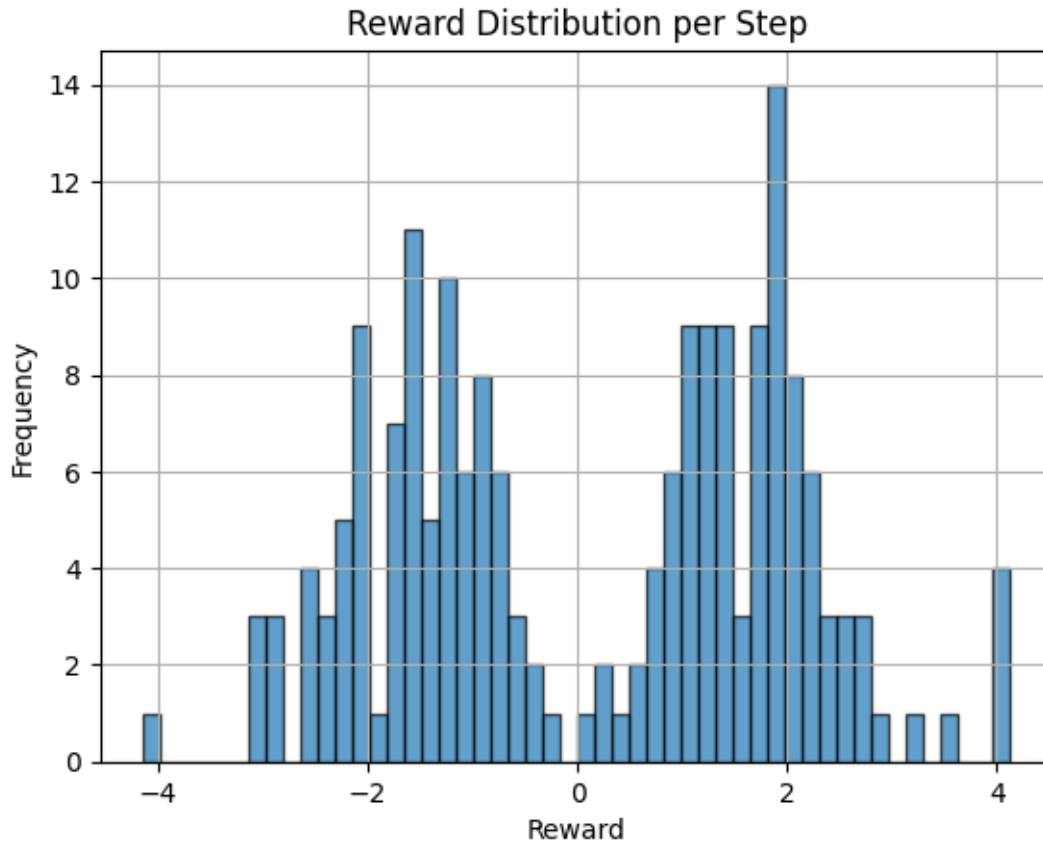
```
plt.ylabel("Episode Reward")
plt.grid(True)
plt.tight_layout()
plt.show()
```



```
[495]: import matplotlib.pyplot as plt
import numpy as np
import pandas as pd

rewards = np.array(env.step_rewards)

# Histogram
plt.hist(rewards, bins=50, alpha=0.7, edgecolor='black')
plt.title("Reward Distribution per Step")
plt.xlabel("Reward")
plt.ylabel("Frequency")
plt.grid(True)
plt.show()
```



```
[496]: import numpy as np
import matplotlib.pyplot as plt
import pandas as pd

rewards = np.array(env.step_rewards)

# --- Raw volatility metrics ---
reward_std = np.std(rewards)
reward_mean = np.mean(rewards)
reward_variance = np.var(rewards)
reward_range = np.max(rewards) - np.min(rewards)

# --- Rolling volatility ---
window = 20
rolling_std = pd.Series(rewards).rolling(window=window).std()


# --- Plot ---
plt.figure(figsize=(10, 4))
plt.plot(rolling_std, label=f"Rolling Std (window={window})")
```

```

plt.axhline(reward_std, color='red', linestyle='--', label='Global Std')
plt.title("Reward Volatility (Rolling Std Dev of Step Rewards)")
plt.xlabel("Step")
plt.ylabel("Rolling Std Dev")
plt.legend()
plt.grid(True)
plt.tight_layout()
plt.show()

# --- Output stats ---
print(f"Mean Step Reward: {reward_mean:.4f}")
print(f"Std Dev of Step Rewards: {reward_std:.4f}")
print(f"Variance: {reward_variance:.4f}")
print(f"Range: {reward_range:.4f}")
print(f"Max Rolling Std ({window}): {rolling_std.max():.4f}")
print(f"Min Rolling Std ({window}): {rolling_std.min():.4f}")

```



```

Mean Step Reward: 0.1547
Std Dev of Step Rewards: 1.8209
Variance: 3.3158
Range: 8.2692
Max Rolling Std (20): 2.3666
Min Rolling Std (20): 1.2189

```

```

[497]: import pandas as pd
import matplotlib.pyplot as plt

def compute_switch_rate(trade_log: pd.DataFrame, plot: bool = True):
    if 'Position' not in trade_log.columns or trade_log.empty:
        raise ValueError("Trade log must contain a 'Position' column and be
        non-empty.")

```

```

positions = trade_log['Position'].tolist()
if len(positions) < 2:
    print("Insufficient trades to compute switching behavior.")
    return 0.0

switches = sum(1 for i in range(1, len(positions)) if positions[i] !=
positions[i - 1])
total_transitions = len(positions) - 1
switch_rate = switches / total_transitions

print(f"Total Trades: {len(positions)}")
print(f"Total Transitions: {total_transitions}")
print(f"Switches: {switches}")
print(f"Switch Rate: {switch_rate:.4f}")

if plot:
    plt.bar(['Same', 'Switch'], [total_transitions - switches, switches],
color=['gray', 'red'])
    plt.title("Trade Direction Transitions")
    plt.ylabel("Count")
    plt.grid(True, axis='y', linestyle='--', alpha=0.6)
    plt.show()

return switch_rate


df_trades = pd.DataFrame(env.trade_log)
switch_rate = compute_switch_rate(df_trades)

```

```

Total Trades: 175
Total Transitions: 174
Switches: 140
Switch Rate: 0.8046

```



```
[498]: import matplotlib.pyplot as plt


total_actions = sum(env.action_counts.values())

actions = list(env.action_counts.keys())
counts = [env.action_counts[a] for a in actions]
labels = ['Long (0)', 'Short (1)']

proportions = [count / total_actions for count in counts]

plt.bar(labels, proportions, color='gray', edgecolor='black')
plt.title("Action Distribution (Policy Bias)")
plt.ylabel("Proportion of Actions")
plt.ylim(0, 1)
plt.grid(True, axis='y')
plt.show()

for label, count, prop in zip(labels, counts, proportions):
    print(f"{label}: Count = {count}, Proportion = {prop:.2%}")
```



Long (0): Count = 213, Proportion = 49.77%

Short (1): Count = 215, Proportion = 50.23%

— Training with seed 71237 — Seed 71237: Reward = -8525.85, Final Balance = 3710.07 Average reward for action Long: 0.2432 Average reward for action Short: 0.0499 Best model updated and saved.

— Training with seed 97600 — Seed 97600: Reward = -6433.13, Final Balance = 5881.59 Average reward for action Long: 0.2968 Average reward for action Short: -0.2705 Best model updated and saved.

— Training with seed 95078 — Seed 95078: Reward = 5224.47, Final Balance = 17003.68 Average reward for action Long: 0.3283 Average reward for action Short: -0.2192 Best model updated and saved.

— Training with seed 57246 — Seed 57246: Reward = -5768.96, Final Balance = 6495.31 Average reward for action Long: 0.3236 Average reward for action Short: -0.0896

— Training with seed 44464 — Seed 44464: Reward = -5472.90, Final Balance = 6903.87 Average reward for action Long: 0.0069 Average reward for action Short: -0.1002

— Training with seed 96800 — Seed 96800: Reward = -7142.38, Final Balance = 5208.52 Average reward for action Long: 0.1932 Average reward for action Short: -0.8241

— Training with seed 52744 — Seed 52744: Reward = 1765.98, Final Balance = 13526.24 Average

reward for action Long: 0.0640 Average reward for action Short: -0.3377

— Training with seed 34500 — Seed 34500: Reward = 26026.52, Final Balance = 37747.27 Average reward for action Long: 0.4125 Average reward for action Short: -0.1004 Best model updated and saved.

— Training with seed 40874 — Seed 40874: Reward = 5781.79, Final Balance = 17401.80 Average reward for action Long: 0.5544 Average reward for action Short: -0.3544

— Training with seed 129147 — Seed 129147: Reward = -9030.47, Final Balance = 3400.98 Average reward for action Long: 0.3973 Average reward for action Short: -0.7070

— Training with seed 71681 — Seed 71681: Reward = -7839.65, Final Balance = 4584.79 Average reward for action Long: 0.1000 Average reward for action Short: -0.1820

— Seed Sweep Complete — Best Seed: 34500 Best Total Reward: 26026.52 Best Final Balance: 37747.27

```
[499]: model.save("NVDA_best_model")
```

1.14 TSLA

```
[570]: from datetime import datetime, timedelta

# Add +1 day to end date
default_end_date = (datetime.now() + timedelta(days=1)).strftime("%Y-%m-%d")
default_start_date = (datetime.now() - timedelta(days=665)).strftime("%Y-%m-%d")
```

```
[571]: # Parameters
ticker = "TSLA"
start_date = default_start_date
end_date = default_end_date
data_filename = "full_data.csv"
signals_filename = "signals_data.csv"

save_historical_data(
    ticker=ticker,
    start_date=start_date,
    end_date=end_date,
    data_filename=data_filename,
    signals_filename=signals_filename
)
```

```
[*****100%*****] 1 of 1 completed
```

```
[571]: 'Saved full_data.csv and signals_data.csv successfully.'
```

```
[572]: import pandas as pd

# Preview full data
df_full = pd.read_csv("full_data.csv", parse_dates=["Date"])
```

```
print(df_full.head())

# Preview signal data
df_signals = pd.read_csv("signals_data.csv", parse_dates=["Date"])
print(df_signals.head())
```

	Date	close	high	low	open	volume	\
0	2023-06-16	260.540009	263.600006	257.209991	258.920013	167563700	
1	2023-06-20	274.450012	274.750000	261.119995	261.500000	165611200	
2	2023-06-21	259.459991	276.989990	257.779999	275.130005	211797100	
3	2023-06-22	264.609985	265.000000	248.250000	250.770004	166875900	
4	2023-06-23	256.600006	262.450012	252.800003	259.290009	176584100	

	EMA_50	EMA_100	EMA_200	EMA_500	...	RSI_Sell	MCDX_Buy	\
0	260.540009	260.540009	260.540009	260.540009	...	1	0	
1	261.085499	260.815454	260.678417	260.595538	...	0	0	
2	261.021753	260.788613	260.666293	260.591004	...	1	0	
3	261.162468	260.864284	260.705534	260.607048	...	1	0	
4	260.983548	260.779843	260.664683	260.591052	...	1	0	

	MCDX_Sell	DSS_Buy	DSS_Sell	ZeroLag	MACD_Buy	ZeroLag	MACD_Sell	\
0	1	0	1		0		1	
1	1	0	1		1		0	
2	1	0	1		1		0	
3	1	0	1		1		0	
4	1	0	1		0		1	

	Basic	MACD_Buy	Basic	MACD_Sell	OverallTrade
0		0		1	Sell
1		1		0	Buy
2		1		0	Sell
3		1		0	Sell
4		1		0	Sell

[5 rows x 42 columns]

	Date	Signal	Z-Score
0	2023-06-20	Basic MACD Buy	NaN
1	2023-06-20	RSI Buy	NaN
2	2023-06-20	ZLMA Buy	NaN
3	2023-06-20	ZeroLag MACD Buy	NaN
4	2023-06-21	RSI Sell	NaN

```
[573]: from stable_baselines3.common.callbacks import BaseCallback
```

```
class RewardTrackingCallback(BaseCallback):
    def __init__(self, verbose=0):
        super().__init__(verbose)
```

```

        self.episode_rewards = []
        self.timesteps = []

    def _on_step(self) -> bool:
        if self.locals.get("done") is not None and any(self.locals["done"]):
            ep_rew = self.locals["rewards"]
            self.episode_rewards.append(sum(ep_rew))
            self.timesteps.append(self.num_timesteps)
        return True

```

```

[ ]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import gymnasium as gym
from gymnasium import spaces
from sb3_contrib import RecurrentPPO
from sb3_contrib.ppo_recurrent.policies import MlpLstmPolicy
from stable_baselines3.common.vec_env import DummyVecEnv
import torch
import random

# --- Fixed Seed ---
SEED = 109961
np.random.seed(SEED)
random.seed(SEED)
torch.manual_seed(SEED)

# --- Load & Clean Data ---
data_path = "full_data.csv"
data = pd.read_csv(data_path, parse_dates=["Date"])

required_cols = [
    'open', 'high', 'low', 'close',
    'zlma', 'ema_value',
    'DSSb', 'DSSsignal',
    'rsi_ma_base',
    'ZeroLagMACD', 'ZeroLagMACD_signal',
    'basicMACD', 'basicMACD_signal',
    'ZScore',
    'ZLMA_Buy', 'ZLMA_Sell',
    'RSI_Buy', 'RSI_Sell',
    'MCDX_Buy', 'MCDX_Sell',
    'DSS_Buy', 'DSS_Sell',
    'ZeroLag MACD_Buy', 'ZeroLag MACD_Sell',
    'Basic MACD_Buy', 'Basic MACD_Sell'
]

```

```

missing = [col for col in required_cols if col not in data.columns]
assert not missing, f"Missing required columns: {missing}"

data = data.dropna(subset=required_cols).reset_index(drop=True)

SWITCH_COST = 1.0
TRANSACTION_COST = 0.001

ACTIVATION = torch.nn.Tanh
# ACTIVATION = partial(nn.LeakyReLU, negative_slope=0.01)

policy_kwargs = dict(
    activation_fn=ACTIVATION
)

# --- Main Trading Environment ---
class TradingEnvRL(gym.Env):
    metadata = {'render_modes': ['human']}

    def __init__(self, data, initial_balance=10000, hold_cost=0.02,
                  volatility_window=14, exploration_steps=500,
                  switch_cost=1.0, reentry_threshold=0.01,
                  dynamic_threshold=False, loss_penalty=0.75, drawdown_penalty=5.
↪0, large_loss_threshold=-2.0):

        super().__init__()
        self.loss_penalty = loss_penalty
        self.drawdown_penalty = drawdown_penalty
        self.large_loss_threshold = large_loss_threshold
        self.data = data.reset_index(drop=True).copy()
        self.n_steps = len(self.data)
        self.initial_balance = initial_balance
        self.hold_cost = hold_cost
        self.volatility_window = volatility_window
        self.exploration_steps = exploration_steps
        self.switch_cost = switch_cost
        self.reentry_threshold = reentry_threshold
        self.dynamic_threshold = dynamic_threshold
        self.step_rewards = []

        self.feature_cols = [
            'open', 'high', 'low', 'close',
            'basicMACD', 'basicMACD_signal',
            'Basic MACD_Buy', 'Basic MACD_Sell'
        ]

        obs_dim = len(self.feature_cols) + 1

```

```

        self.observation_space = spaces.Box(low=-np.inf, high=np.inf,
↪shape=(obs_dim,), dtype=np.float32)
        self.action_space = spaces.Discrete(2) # 0=Long, 1=Short

        self._compute_volatility_limit()

    def _compute_volatility_limit(self):
        returns = self.data['close'].pct_change()
        self.data['volatility'] = returns.rolling(self.volatility_window).std()
        self.data['adaptive_hold'] = (10 / (self.data['volatility'] * 100)).
↪clip(lower=3, upper=20).fillna(10).astype(int)

    def reset(self, seed=None, options=None):
        if seed is not None:
            np.random.seed(seed)
            random.seed(seed)
            torch.manual_seed(seed)

        self.current_step = 0
        self.position = 0
        self.entry_price = 0.0
        self.entry_date = None
        self.hold_counter = 0
        self.switch_count = 0
        self.balance = self.initial_balance
        self.cumulative_pnl = 0.0
        self.trade_log = []
        self.action_counts = {0: 0, 1: 0}
        self.reward_tracker = {0: [], 1: []}
        self.consecutive_losses = 0
        self.equity_curve = [self.initial_balance]
        self.step_rewards = []
        return self._get_obs(), {}

    def _get_obs(self):
        row = self.data.iloc[self.current_step]
        features = row[self.feature_cols].values.astype(np.float32)
        pos_feature = np.array([self.position], dtype=np.float32)
        return np.concatenate([features, pos_feature])

    def _force_close(self):
        row = self.data.iloc[self.current_step]
        current_price = float(row['close'])
        current_date = row['Date'].strftime("%Y-%m-%d")

        if self.position == 0:
            return 0.0

```

```

# --- Core Return Logic ---
trade_pct = ((current_price / self.entry_price - 1) * 100) if self.
↪position == 1 else ((self.entry_price / current_price - 1) * 100)
pos_str = 'Long' if self.position == 1 else 'Short'
gross_return = trade_pct / 100
transaction_cost = TRANSACTION_COST * current_price

old_balance = self.balance
self.balance -= transaction_cost
self.balance *= (1 + gross_return)
net_profit = self.balance - old_balance

reward = net_profit

# --- Track Trade History ---
self.cumulative_pnl += trade_pct
compounded_pnl = (self.trade_log[-1]['CompoundedFactor'] * (1 + ↪
↪gross_return)) if self.trade_log else (1 + gross_return)
compounded_pnl_pct = (compounded_pnl - 1) * 100

self.trade_log.append({
    'EntryDate': self.entry_date,
    'ExitDate': current_date,
    'Position': pos_str,
    'EntryPrice': round(self.entry_price, 4),
    'ExitPrice': round(current_price, 4),
    'PnL%': round(trade_pct, 2),
    'CumulativePnL%': round(self.cumulative_pnl, 2),
    'CompoundedFactor': compounded_pnl,
    'CompoundedPnL%': round(compounded_pnl_pct, 2)
})

# --- Penalty for Large Loss ---
if trade_pct < -2.0:
    reward += trade_pct * 2 # Stronger penalty for large loss

# --- Track & Penalize Consecutive Losses ---
if trade_pct < 0:
    self.consecutive_losses += 1
    reward -= self.consecutive_losses * 0.75 # Growing penalty
else:
    self.consecutive_losses = 0

# --- Optional: Drawdown Penalty ---
self.equity_curve.append(self.balance)
max_balance = max(self.equity_curve)

```

```

    if max_balance > 0:
        drawdown = (max_balance - self.balance) / max_balance
        reward -= 5.0 * drawdown # Penalize deeper drawdown

    # --- Reset State ---
    self.position = 0
    self.entry_price = 0.0
    self.entry_date = None
    self.hold_counter = 0

    return reward

def step(self, action):
    if isinstance(action, np.ndarray):
        action = int(action.item())

    row = self.data.iloc[self.current_step]
    current_price = float(row['close'])
    current_date = row['Date'].strftime("%Y-%m-%d")

    if self.current_step < self.exploration_steps:
        action = self.action_space.sample()

    self.action_counts[action] += 1
    reward = 0.0
    target_position = 1 if action == 0 else -1

    # === Holding Same Position ===
    if self.position == target_position:
        self.hold_counter += 1
        price_change = (current_price - self.entry_price) / self.entry_price
        step_return = price_change if self.position == 1 else -price_change

        step_reward = np.sign(step_return) * np.sqrt(abs(step_return)) * 10
        step_reward = np.clip(step_reward, -50, 50)
        step_reward -= self.hold_cost
        reward += step_reward
        self.reward_tracker[action].append(step_reward)
        self.step_rewards.append(reward)

    # === New Position from Flat ===
    elif self.position == 0:
        self.position = target_position
        self.entry_price = current_price
        self.entry_date = current_date
        self.hold_counter = 1
        reward -= TRANSACTION_COST * current_price

```

```

# === Switch Position ===
else:
    hold_penalty = max(0, 3 - self.hold_counter) * 5.0
    switch_penalty = self.switch_cost + hold_penalty
    reward += self._force_close()
    reward -= switch_penalty

# === Dynamic Re-entry Threshold ===
if self.current_step > 0:
    prev_close = self.data.iloc[self.current_step - 1]['close']
    recent_return = abs((current_price - prev_close) / prev_close)

    # Volatility-aware reentry threshold
    recent_volatility = row['volatility']
    reentry_threshold = 0.5 * recent_volatility
else:
    recent_return = 0
    reentry_threshold = 0.01 # Fallback

if recent_return > reentry_threshold:
    self.position = target_position
    self.entry_price = current_price
    self.entry_date = current_date
    self.hold_counter = 1
    reward -= TRANSACTION_COST * current_price

self.current_step += 1
terminated = self.current_step >= self.n_steps - 1

if terminated and self.position != 0:
    reward += self._force_close()

obs = self._get_obs() if not terminated else np.zeros(self.
↪observation_space.shape, dtype=np.float32)
return obs, reward, terminated, False, {}

def render(self):
    print(f"Step: {self.current_step}, Position: {self.position}, Balance:␣
↪{self.balance:.2f}")

def save_trade_log(self, filename="trade_log.csv"):
    df = pd.DataFrame(self.trade_log)
    if "CompoundedFactor" in df.columns:
        df = df.drop(columns=["CompoundedFactor"])
    df.to_csv(filename, index=False)

```



```

# --- Training ---
env = TradingEnvRL(data, initial_balance=10000)
vec_env = DummyVecEnv([lambda: env])
vec_env.seed(SEED)

model = RecurrentPPO(
    policy=MlpLstmPolicy,
    env=vec_env,
    verbose=1,
    n_steps=64,
    batch_size=32,
    learning_rate=0.01,
    gamma=0.99,
    ent_coef=0.01,
    seed=SEED,
    policy_kwargs=policy_kwargs
)

callback = RewardTrackingCallback()

model.learn(total_timesteps=3000, callback=callback)
print(f"\nTraining Complete")

```

```

[575]: # --- Evaluation ---
obs, _ = env.reset()
state = None
done = False
total_reward = 0
portfolio_values = []
final_action = None

while not done:
    action, state = model.predict(obs, state=state, deterministic=True)
    final_action = action
    obs, reward, done, _, _ = env.step(action)
    total_reward += reward

    current_index = min(env.current_step, len(env.data) - 1)
    current_price = env.data.loc[current_index, 'close']
    unrealized = (
        (current_price - env.entry_price) if env.position == 1 else
        (env.entry_price - current_price) if env.position == -1 else
        0.0
    )
    mtm_equity = env.balance + unrealized
    portfolio_values.append(mtm_equity)

```

```

# --- Save trade log ---
env.save_trade_log("trade_log_recurrent.csv")

# --- Final Model Signal ---
signal_str = "BUY" if final_action == 0 else "SELL"
latest_date = env.data['Date'].iloc[env.current_step - 1].strftime("%Y-%m-%d")
print(f"\nLatest model signal at {latest_date}: {signal_str}")

# --- Summary Output ---
print(f"Total Reward: {total_reward:.2f}")
print(f"Final Balance: {env.balance:.2f}")
print("Trade Log (Recurrent PP0):")
print(pd.DataFrame(env.trade_log))
print("Action counts:", env.action_counts)
print("Average reward per action:")
for k, v in env.reward_tracker.items():
    mean_r = np.mean(v) if v else 0
    print(f"Action {k} ({'Long' if k==0 else 'Short'}): {mean_r:.4f}")

```

Latest model signal at 2025-04-09: BUY

Total Reward: 10069.41

Final Balance: 22325.22

Trade Log (Recurrent PP0):

	EntryDate	ExitDate	Position	EntryPrice	ExitPrice	PnL%	\
0	2023-07-27	2023-07-28	Long	255.71	266.44	4.20	
1	2023-07-31	2023-08-01	Short	267.43	261.07	2.44	
2	2023-08-02	2023-08-04	Long	254.11	253.86	-0.10	
3	2023-08-07	2023-08-08	Long	251.45	249.70	-0.70	
4	2023-08-09	2023-08-10	Short	242.19	245.34	-1.28	
..	
186	2025-04-01	2025-04-03	Long	268.46	267.28	-0.44	
187	2025-04-03	2025-04-04	Short	267.28	239.43	11.63	
188	2025-04-04	2025-04-07	Long	239.43	233.29	-2.56	
189	2025-04-08	2025-04-09	Long	221.86	272.20	22.69	
190	2025-04-09	2025-04-10	Short	272.20	252.40	7.84	

	CumulativePnL%	CompoundedFactor	CompoundedPnL%
0	4.20	1.041962	4.20
1	6.63	1.067345	6.73
2	6.53	1.066295	6.63
3	5.84	1.058874	5.89
4	4.55	1.045279	4.53
..
186	70.37	1.557815	55.78
187	82.01	1.739017	73.90
188	79.44	1.694421	69.44
189	102.13	2.078885	107.89

190 109.98 2.241967 124.20

[191 rows x 9 columns]
Action counts: {0: 205, 1: 223}
Average reward per action:
Action 0 (Long): 0.4063
Action 1 (Short): 0.2677

```
[576]: # --- Get Latest Signal from Model ---
last_index = len(env.data) - 1
env.current_step = last_index
obs = env._get_obs()

# Add batch dimension and run prediction
obs_input = obs[np.newaxis, :]
action, _ = model.predict(obs_input, deterministic=True)

# Convert action to trading signal
# 0 = Long => BUY, 1 = Short => SELL
if action == 0:
    current_signal = "BUY"
else:
    current_signal = "SELL"

print(f"\nLatest model signal at {env.data.iloc[last_index]['Date'].date()}:
↳{current_signal}")
```

Latest model signal at 2025-04-10: BUY

```
[577]: # --- Plot Performance with Trade Markers ---
buy_hold_line = [env.initial_balance * (p / data['close'].iloc[0]) for p in
↳data['close'].iloc[:len(portfolio_values)]]

date_to_step = {row['Date'].strftime('%Y-%m-%d'): i for i, row in data.iloc[:
↳len(portfolio_values)].iterrows()}
entry_points = []
exit_points = []

for trade in env.trade_log:
    entry_step = date_to_step.get(trade['EntryDate'])
    exit_step = date_to_step.get(trade['ExitDate'])
    if entry_step is not None and entry_step < len(portfolio_values):
        entry_points.append((entry_step, portfolio_values[entry_step]))
    if exit_step is not None and exit_step < len(portfolio_values):
        exit_points.append((exit_step, portfolio_values[exit_step]))

plt.figure(figsize=(12, 6))
```


```

plt.plot(portfolio_values, label="Recurrent PPO Portfolio", linewidth=2)
plt.plot(buy_hold_line, label="Buy & Hold Benchmark", linestyle="--",
        color="gray")

if entry_points:
    entry_steps, entry_vals = zip(*entry_points)
    plt.scatter(entry_steps, entry_vals, color='green', marker='^', s=80,
        label="Trade Entry")
if exit_points:
    exit_steps, exit_vals = zip(*exit_points)
    plt.scatter(exit_steps, exit_vals, color='red', marker='v', s=80,
        label="Trade Exit")

plt.title("Portfolio Value: Recurrent PPO vs. Buy & Hold")
plt.xlabel("Time Step")
plt.ylabel("Portfolio Value")
plt.legend()
plt.grid(True)
plt.tight_layout()
plt.savefig("recurrent_ppo_performance.png")
plt.show()

```




```

[578]: import matplotlib.pyplot as plt

plt.figure(figsize=(10, 5))
plt.plot(callback.timesteps, callback.episode_rewards, marker='o')
plt.title("Episode Reward Over Time")
plt.xlabel("Timesteps")

```


```
plt.ylabel("Episode Reward")
plt.grid(True)
plt.tight_layout()
plt.show()
```



```
[579]: import matplotlib.pyplot as plt
import numpy as np
import pandas as pd

rewards = np.array(env.step_rewards)

# Histogram
plt.hist(rewards, bins=50, alpha=0.7, edgecolor='black')
plt.title("Reward Distribution per Step")
plt.xlabel("Reward")
plt.ylabel("Frequency")
plt.grid(True)
plt.show()
```



```
[580]: import numpy as np
import matplotlib.pyplot as plt
import pandas as pd

rewards = np.array(env.step_rewards)

# --- Raw volatility metrics ---
reward_std = np.std(rewards)
reward_mean = np.mean(rewards)
reward_variance = np.var(rewards)
reward_range = np.max(rewards) - np.min(rewards)

# --- Rolling volatility ---
window = 20
rolling_std = pd.Series(rewards).rolling(window=window).std()


# --- Plot ---
plt.figure(figsize=(10, 4))
plt.plot(rolling_std, label=f"Rolling Std (window={window})")
```

```

plt.axhline(reward_std, color='red', linestyle='--', label='Global Std')
plt.title("Reward Volatility (Rolling Std Dev of Step Rewards)")
plt.xlabel("Step")
plt.ylabel("Rolling Std Dev")
plt.legend()
plt.grid(True)
plt.tight_layout()
plt.show()

# --- Output stats ---
print(f"Mean Step Reward: {reward_mean:.4f}")
print(f"Std Dev of Step Rewards: {reward_std:.4f}")
print(f"Variance: {reward_variance:.4f}")
print(f"Range: {reward_range:.4f}")
print(f"Max Rolling Std ({window}): {rolling_std.max():.4f}")
print(f"Min Rolling Std ({window}): {rolling_std.min():.4f}")

```



```

Mean Step Reward: 0.3273
Std Dev of Step Rewards: 1.8234
Variance: 3.3247
Range: 8.1576
Max Rolling Std (20): 2.2853
Min Rolling Std (20): 1.3074

```

```

[581]: import pandas as pd
import matplotlib.pyplot as plt

def compute_switch_rate(trade_log: pd.DataFrame, plot: bool = True):
    if 'Position' not in trade_log.columns or trade_log.empty:
        raise ValueError("Trade log must contain a 'Position' column and be
↪non-empty.")

```

```

positions = trade_log['Position'].tolist()
if len(positions) < 2:
    print("Insufficient trades to compute switching behavior.")
    return 0.0

switches = sum(1 for i in range(1, len(positions)) if positions[i] !=
positions[i - 1])
total_transitions = len(positions) - 1
switch_rate = switches / total_transitions

print(f"Total Trades: {len(positions)}")
print(f"Total Transitions: {total_transitions}")
print(f"Switches: {switches}")
print(f"Switch Rate: {switch_rate:.4f}")

if plot:
    plt.bar(['Same', 'Switch'], [total_transitions - switches, switches],
color=['gray', 'red'])
    plt.title("Trade Direction Transitions")
    plt.ylabel("Count")
    plt.grid(True, axis='y', linestyle='--', alpha=0.6)
    plt.show()

return switch_rate


df_trades = pd.DataFrame(env.trade_log)
switch_rate = compute_switch_rate(df_trades)

```

```

Total Trades: 191
Total Transitions: 190
Switches: 149
Switch Rate: 0.7842

```

```
[582]: import matplotlib.pyplot as plt


total_actions = sum(env.action_counts.values())

actions = list(env.action_counts.keys())
counts = [env.action_counts[a] for a in actions]
labels = ['Long (0)', 'Short (1)']

proportions = [count / total_actions for count in counts]

plt.bar(labels, proportions, color='gray', edgecolor='black')
plt.title("Action Distribution (Policy Bias)")
plt.ylabel("Proportion of Actions")
plt.ylim(0, 1)
plt.grid(True, axis='y')
plt.show()

for label, count, prop in zip(labels, counts, proportions):
    print(f"{label}: Count = {count}, Proportion = {prop:.2%}")
```



Long (0): Count = 205, Proportion = 47.90%

Short (1): Count = 223, Proportion = 52.10%

— Training with seed 84511 — Seed 84511: Reward = -69.15, Final Balance = 12038.12 Average reward for action Long: 0.3914 Average reward for action Short: -0.5380 Best model updated and saved.

— Training with seed 97974 — Seed 97974: Reward = -9070.66, Final Balance = 3522.30 Average reward for action Long: 0.0136 Average reward for action Short: -0.0997

— Training with seed 64982 — Seed 64982: Reward = 7790.18, Final Balance = 19863.73 Average reward for action Long: -0.2441 Average reward for action Short: 0.6258 Best model updated and saved.

— Training with seed 62977 — Seed 62977: Reward = 2570.52, Final Balance = 14840.80 Average reward for action Long: 0.4955 Average reward for action Short: -0.5567

— Training with seed 111581 — Seed 111581: Reward = -6663.71, Final Balance = 5529.87 Average reward for action Long: -0.0897 Average reward for action Short: -0.0450

— Training with seed 43633 — Seed 43633: Reward = -4373.75, Final Balance = 8066.42 Average reward for action Long: -0.0223 Average reward for action Short: 0.1117

— Training with seed 110692 — Seed 110692: Reward = 3013.97, Final Balance = 15078.01 Average reward for action Long: -0.2785 Average reward for action Short: 0.6881

— Training with seed 65785 — Seed 65785: Reward = 2912.61, Final Balance = 15068.03 Average reward for action Long: 0.2839 Average reward for action Short: 0.2509

— Training with seed 94074 — Seed 94074: Reward = -4182.82, Final Balance = 8117.05 Average reward for action Long: -0.0971 Average reward for action Short: -0.4204

— Training with seed 38615 — Seed 38615: Reward = -3628.04, Final Balance = 8626.54 Average reward for action Long: -0.6513 Average reward for action Short: -0.1777

— Training with seed 109961 — Seed 109961: Reward = 10069.41, Final Balance = 22325.22 Average reward for action Long: 0.4063 Average reward for action Short: 0.2677 Best model updated and saved.

— Seed Sweep Complete — Best Seed: 109961 Best Total Reward: 10069.41 Best Final Balance: 22325.22

```
[583]: model.save("TSLA_best_model")
```

1.15 Model Inference Run 1 (NVDA)

Inference determinism test, run this and below to compare.

```
[612]: import numpy as np
import matplotlib.pyplot as plt
from sb3_contrib import RecurrentPPO
from stable_baselines3.common.vec_env import DummyVecEnv

# === Recreate environment ===
inference_env = DummyVecEnv([lambda: TradingEnvRL(data, initial_balance=10000)])

# === Load model ===
model = RecurrentPPO.load("NVDA_best_model", env=inference_env)

# === Inference Setup ===
obs = inference_env.reset()
lstm_states = None
episode_starts = np.ones((inference_env.num_envs,), dtype=bool)
done = [False]
total_reward = 0.0
step_count = 0
portfolio_values = []

env = inference_env.envs[0] # Unwrapped env
data_used = env.data
last_action = None # Track final signal

# === Inference Loop ===
while not done[0]:
    action, lstm_states = model.predict(
        obs,
```

```

        state=lstm_states,
        episode_start=episode_starts,
        deterministic=True
    )
    obs, reward, done, info = inference_env.step(action)
    episode_starts = done
    total_reward += reward[0]
    step_count += 1
    last_action = action[0] # Capture the final action taken

    # Track portfolio value (mark-to-market)
    idx = min(env.current_step, len(data_used) - 1)
    price_now = data_used.loc[idx, 'close']
    unrealized = (
        (price_now - env.entry_price) if env.position == 1 else
        (env.entry_price - price_now) if env.position == -1 else
        0.0
    )
    portfolio_values.append(env.balance + unrealized)

# === Save trades ===
env.save_trade_log("inference_trades.csv")

# === Plot performance ===
buy_hold_line = [env.initial_balance * (p / data_used['close'].iloc[0]) for p
    ↪ in data_used['close'].iloc[:len(portfolio_values)]]

date_to_step = {row['Date'].strftime('%Y-%m-%d'): i for i, row in data_used.
    ↪ iloc[:len(portfolio_values)].iterrows()}
entry_points = []
exit_points = []

for trade in env.trade_log:
    entry_step = date_to_step.get(trade['EntryDate'])
    exit_step = date_to_step.get(trade['ExitDate'])
    if entry_step is not None and entry_step < len(portfolio_values):
        entry_points.append((entry_step, portfolio_values[entry_step]))
    if exit_step is not None and exit_step < len(portfolio_values):
        exit_points.append((exit_step, portfolio_values[exit_step]))

plt.figure(figsize=(12, 6))
plt.plot(portfolio_values, label="Recurrent PPO Portfolio", linewidth=2)
plt.plot(buy_hold_line, label="Buy & Hold Benchmark", linestyle="--",
    ↪ color="gray")

if entry_points:
    entry_steps, entry_vals = zip(*entry_points)

```


```

plt.scatter(entry_steps, entry_vals, color='green', marker='^', s=80,
↳label="Trade Entry")
if exit_points:
    exit_steps, exit_vals = zip(*exit_points)
    plt.scatter(exit_steps, exit_vals, color='red', marker='v', s=80,
↳label="Trade Exit")

plt.title("Portfolio Value: Recurrent PPO vs. Buy & Hold")
plt.xlabel("Time Step")
plt.ylabel("Portfolio Value")
plt.legend()
plt.grid(True)
plt.tight_layout()
plt.savefig("recurrent_ppo_performance.png")
plt.show()

# === Print summary ===
signal_map = {0: "BUY", 1: "SELL"}
latest_signal = signal_map.get(last_action, "UNKNOWN")
print(f"Inference completed in {step_count} steps.")
print(f"Total reward accumulated: {total_reward:.2f}")
print(f"Final model signal: {latest_signal}")

```



Inference completed in 428 steps.
 Total reward accumulated: 12608.24
 Final model signal: BUY

1.16 Model Inference Run 2 (NVDA)

```
[613]: import numpy as np
import matplotlib.pyplot as plt
from sb3_contrib import RecurrentPPO
from stable_baselines3.common.vec_env import DummyVecEnv

# === Recreate environment ===
inference_env = DummyVecEnv([lambda: TradingEnvRL(data, initial_balance=10000)])

# === Load model ===
model = RecurrentPPO.load("NVDA_best_model", env=inference_env)

# === Inference Setup ===
obs = inference_env.reset()
lstm_states = None
episode_starts = np.ones((inference_env.num_envs,), dtype=bool)
done = [False]
total_reward = 0.0
step_count = 0
portfolio_values = []

env = inference_env.envs[0] # Unwrapped env
data_used = env.data
last_action = None # Track final signal

# === Inference Loop ===
while not done[0]:
    action, lstm_states = model.predict(
        obs,
        state=lstm_states,
        episode_start=episode_starts,
        deterministic=True
    )
    obs, reward, done, info = inference_env.step(action)
    episode_starts = done
    total_reward += reward[0]
    step_count += 1
    last_action = action[0] # Capture the final action taken

# Track portfolio value (mark-to-market)
idx = min(env.current_step, len(data_used) - 1)
price_now = data_used.loc[idx, 'close']
unrealized = (
    (price_now - env.entry_price) if env.position == 1 else
    (env.entry_price - price_now) if env.position == -1 else
    0.0
```

```

    )
    portfolio_values.append(env.balance + unrealized)

# === Save trades ===
env.save_trade_log("inference_trades.csv")

# === Plot performance ===
buy_hold_line = [env.initial_balance * (p / data_used['close'].iloc[0]) for p
    ↪in data_used['close'].iloc[:len(portfolio_values)]]

date_to_step = {row['Date'].strftime('%Y-%m-%d'): i for i, row in data_used.
    ↪iloc[:len(portfolio_values)].iterrows()}
entry_points = []
exit_points = []

for trade in env.trade_log:
    entry_step = date_to_step.get(trade['EntryDate'])
    exit_step = date_to_step.get(trade['ExitDate'])
    if entry_step is not None and entry_step < len(portfolio_values):
        entry_points.append((entry_step, portfolio_values[entry_step]))
    if exit_step is not None and exit_step < len(portfolio_values):
        exit_points.append((exit_step, portfolio_values[exit_step]))

plt.figure(figsize=(12, 6))
plt.plot(portfolio_values, label="Recurrent PPO Portfolio", linewidth=2)
plt.plot(buy_hold_line, label="Buy & Hold Benchmark", linestyle="--",
    ↪color="gray")

if entry_points:
    entry_steps, entry_vals = zip(*entry_points)
    plt.scatter(entry_steps, entry_vals, color='green', marker='^', s=80,
    ↪label="Trade Entry")
if exit_points:
    exit_steps, exit_vals = zip(*exit_points)
    plt.scatter(exit_steps, exit_vals, color='red', marker='v', s=80,
    ↪label="Trade Exit")

plt.title("Portfolio Value: Recurrent PPO vs. Buy & Hold")
plt.xlabel("Time Step")
plt.ylabel("Portfolio Value")
plt.legend()
plt.grid(True)
plt.tight_layout()
plt.savefig("recurrent_ppo_performance.png")
plt.show()


# === Print summary ===

```

```

signal_map = {0: "BUY", 1: "SELL"}
latest_signal = signal_map.get(last_action, "UNKNOWN")
print(f"Inference completed in {step_count} steps.")
print(f"Total reward accumulated: {total_reward:.2f}")
print(f"Final model signal: {latest_signal}")

```



Inference completed in 428 steps.
 Total reward accumulated: -4407.63
 Final model signal: BUY

1.17 Changelog

1.17.1 Version 2.8.0

- Added more statistical metrics
- Code cleanup
- Clearer notebook organization
- Added model inference

1.17.2 2.5.0

- Updated model rules and tests

1.17.3 2.2.0

- Added per model saving
- Split entry points from Gradio
- Organization cleanup

1.17.4 2.1.0

- New major features: Recurrent PPO module training, seed search, and portfolio performance testing
- Removed DQN module
- Improved CSV saving function

1.17.5 1.2.3

- Adjusted default weights

1.17.6 1.2.2

- Fixed trade table
- Minor formatting

1.17.7 1.2.1

- Added trade table
- Added user adjustable weights
- Added basic MACD
- Added trade signals